

RealView[®] Compilation Tools

Version 2.2

Developer Guide



RealView Compilation Tools

Developer Guide

Copyright © 2002-2005 ARM Limited. All rights reserved.

Release Information

The following changes have been made to this book.

Change History

Date	Issue	Change
August 2002	A	Release 1.2
January 2003	B	Release 2.0
September 2003	C	Release 2.0.1 for RVDS v2.0
January 2004	D	Release 2.1 for RVDS v2.1
December 2004	E	Release 2.2 for RVDS v2.2
May 2005	F	Release 2.2 for RVDS v2.2 SP1

Proprietary Notice

Words and logos marked with ® or ™ are registered trademarks or trademarks owned by ARM Limited. Other brands and names mentioned herein may be the trademarks of their respective owners.

Neither the whole nor any part of the information contained in, or the product described in, this document may be adapted or reproduced in any material form except with the prior written permission of the copyright holder.

The product described in this document is subject to continuous developments and improvements. All particulars of the product and its use contained in this document are given by ARM in good faith. However, all warranties implied or expressed, including but not limited to implied warranties of merchantability, or fitness for purpose, are excluded.

This document is intended only to assist the reader in the use of the product. ARM Limited shall not be liable for any loss or damage arising from the use of any information in this document, or any error or omission in such information, or any incorrect use of the product.

Confidentiality Status

This document is Non-Confidential. The right to use, copy and disclose this document may be subject to license restrictions in accordance with the terms of the agreement entered into by ARM and the party that ARM delivered this document to.

Product Status

The information in this document is final, that is for a developed product.

Web Address

<http://www.arm.com>

Contents

RealView Compilation Tools Developer Guide

Preface

About this book	viii
Feedback	xi

Chapter 1

Introduction

1.1 About RVCT	1-2
1.2 General programming issues	1-3
1.3 Developing for the ARM processors	1-9
1.4 ARM architecture v6 support	1-13

Chapter 2

Embedded Software Development

2.1 About embedded software development	2-2
2.2 Default compilation tool behavior in the absence of a target system	2-4
2.3 Tailoring the C library to your target hardware	2-11
2.4 Tailoring the image memory map to your target hardware	2-15
2.5 Reset and initialization	2-25
2.6 Further memory map considerations	2-35

Chapter 3

Writing Position Independent Code and Data

3.1 Position independence	3-2
3.2 Read-only position independence	3-3
3.3 Read-write position independence	3-6

Chapter 4	Interworking ARM and Thumb	
4.1	About interworking	4-2
4.2	Assembly language interworking	4-7
4.3	C and C++ interworking and veneers	4-13
4.4	Assembly language interworking using veneers	4-18
Chapter 5	Mixing C, C++, and Assembly Language	
5.1	Using the inline and embedded assemblers	5-2
5.2	Accessing C global variables from assembly code	5-4
5.3	Using C header files from C++	5-5
5.4	Calling between C, C++, and ARM assembly language	5-7
Chapter 6	Handling Processor Exceptions	
6.1	About processor exceptions	6-2
6.2	Determining the processor state	6-6
6.3	Entering and leaving an exception	6-8
6.4	Handling an exception	6-13
6.5	Installing an exception handler	6-14
6.6	SWI handlers	6-19
6.7	Interrupt handlers	6-29
6.8	Reset handlers	6-39
6.9	Undefined Instruction handlers	6-40
6.10	Prefetch Abort handler	6-41
6.11	Data Abort handler	6-42
6.12	Chaining exception handlers	6-44
6.13	System mode	6-46
Chapter 7	Debug Communications Channel	
7.1	About the Debug Communications Channel	7-2
7.2	Target transfer of data	7-3
7.3	Polled debug communications	7-4
7.4	Interrupt-driven debug communications	7-8
7.5	Access from Thumb state	7-9
Appendix A	Using the Procedure Call Standard	
A.1	About the AAPCS	A-2
A.2	Using AAPCS options	A-4
	Glossary	

Preface

This preface introduces the *RealView Compilation Tools Developer Guide*. It contains the following sections:

- *About this book* on page viii
- *Feedback* on page xi.

About this book

This book contains information that helps you with specific issues when developing code for the ARM® family of *Reduced Instruction Set Computing* (RISC) processors. The chapters in this book, and the examples used, assume that you are using the latest release of *RealView® Compilation Tools* (RVCT) to develop your code.

Intended audience

This book is written for all developers writing code for ARM architecture-based processors. It assumes that you are an experienced software developer, and that you are familiar with the ARM development tools described in *RealView Compilation Tools v2.2 Essentials Guide*.

Using this book

This book is organized into the following chapters and appendixes:

Chapter 1 *Introduction*

Read this chapter for an introduction to RVCT.

Chapter 2 *Embedded Software Development*

Read this chapter for details of how to develop embedded applications with RVCT. It describes the default RVCT behavior in the absence of a target system, and how to tailor the C library and image memory map to your target system.

Chapter 3 *Writing Position Independent Code and Data*

Read this chapter for details of how to write position independent code and data that makes use of the *Procedure Call Standard for the ARM Architecture* (AAPCS).

Chapter 4 *Interworking ARM and Thumb*

Read this chapter for details of how to change between ARM state and Thumb state when writing code for processors that implement the Thumb instruction set.

Chapter 5 *Mixing C, C++, and Assembly Language*

Read this chapter for details of how to write mixed C, C++, and ARM assembly language code. It also describes how to use the ARM inline and embedded assembler from C and C++.

Chapter 6 Handling Processor Exceptions

Read this chapter for details of how to handle the various types of exception supported by ARM processors.

Chapter 7 Debug Communications Channel

Read this chapter for a description of how to use the *Debug Communications Channel* (DCC).

Appendix A Using the Procedure Call Standard

Read this appendix for details about command-line options to control compliance with the AAPCS.

Glossary An alphabetically arranged glossary defines the special terms used in this book.

This book assumes that you have installed your ARM software in the default location for example, on Windows this might be *volume:\Program Files\ARM*. This is assumed to be the location of *install_directory* when referring to path names, for example *install_directory\Documentation\...* You might have to change this if you have installed your ARM software in a different location.

Typographical conventions

The following typographical conventions are used in this book:

<code>monospace</code>	Denotes text that can be entered at the keyboard, such as commands, file and program names, and source code.
<u><code>monospace</code></u>	Denotes a permitted abbreviation for a command or option. The underlined text can be entered instead of the full command or option name.
<i>monospace italic</i>	Denotes arguments to commands and functions where the argument is to be replaced by a specific value.
<code>monospace bold</code>	Denotes language keywords when used outside example code.
<i>italic</i>	Highlights important notes, introduces special terminology, denotes internal cross-references, and citations.
bold	Highlights interface elements, such as menu names. Also used for emphasis in descriptive lists, where appropriate, and for ARM processor signal names.

Further reading

This section lists publications from both ARM Limited and third parties that provide additional information on developing code for the ARM family of processors.

ARM periodically provides updates and corrections to its documentation. See <http://www.arm.com> for current errata sheets and addenda, and the ARM Frequently Asked Questions.

ARM publications

This book contains general information on developing applications for the ARM family of processors. See the following books in the RVCT document suite for information on other components:

- *RealView Compilation Tools v2.2 Essentials Guide* (ARM DUI 0202)
- *RealView Compilation Tools v2.2 Assembler Guide* (ARM DUI 0204)
- *RealView Compilation Tools v2.2 Compiler and Libraries Guide* (ARM DUI 0205)
- *RealView Compilation Tools v2.2 Linker and Utilities Guide* (ARM DUI 0206).

For full information about the base standard, software interfaces, and standards supported by ARM, see *install_directory\Documentation\Specifications\....*

In addition, see the following documentation for specific information relating to ARM products:

- *ARM Architecture Reference Manual* (ARM DDI 0100)
- *ARM Reference Peripheral Specification* (ARM DDI 0062)
- the ARM datasheet or technical reference manual for your hardware device.

Other publications

For a comprehensive introduction to ARM architecture, see Steve Furber, *ARM system-on-chip architecture* (2nd edition, 2000). Addison Wesley, ISBN 0-201-67519-6.

Feedback

ARM Limited welcomes feedback on both RealView Compilation Tools, and its documentation.

Feedback on RealView Compilation Tools

If you have any problems with RVCT, contact your supplier. To help them provide a rapid and useful response, give:

- your name and company
- the serial number of the product
- details of the release you are using
- details of the platform you are running on, such as the hardware platform, operating system type and version
- a small standalone sample of code that reproduces the problem
- a clear explanation of what you expected to happen, and what actually happened
- the commands you used, including any command-line options
- sample output illustrating the problem
- the version string of the tool, including the version number and date.

Feedback on this book

If you notice any errors or omissions in this book, send email to errata@arm.com giving:

- the document title
- the document number
- the page number(s) to which your comments apply
- a concise explanation of the problem.

General suggestions for additions and improvements are also welcome.

Chapter 1

Introduction

This chapter introduces this book and begins to describe how the *RealView® Compilation Tools* (RVCT) can be used to develop code. It contains the following sections:

- *About RVCT* on page 1-2
- *General programming issues* on page 1-3
- *Developing for the ARM processors* on page 1-9
- *ARM architecture v6 support* on page 1-13.

1.1 About RVCT

RVCT consists of a suite of applications, together with supporting documentation and examples, that enable you to write applications for the ARM® family of RISC processors. You can use RVCT to build C, C++, and ARM assembly language programs.

The RVCT toolkit consists of the following major components:

- command-line development tools
- utilities
- supporting software.

This book contains information that helps you with specific issues when developing code for ARM-based systems. The chapters in this book, and the examples used, assume that you are using the latest release of RVCT to develop your code.

If you are upgrading to RVCT from a previous release, ensure that you read *RealView Compilation Tools v2.2 Essentials Guide* for details about new features and enhancements in this release.

If you are new to RVCT, read *RealView Compilation Tools v2.2 Essentials Guide* for an overview of the ARM tools and an introduction to using them as part of your development project.

For information about previous releases of RVCT, see Appendix A in *RealView Compilation Tools v2.2 Essentials Guide*.

See *ARM publications* on page x for a list of the other books in the RVCT documentation suite that give information on the ARM assembler, compiler, and supporting software.

1.1.1 Using the examples

This book references examples provided with RealView Developer Suite in the main examples directory `install_directory\RVDS\Examples`. See *RealView Developer Suite Getting Started Guide* for a summary of the examples provided.

1.2 General programming issues

ARM processors are *Reduced Instruction Set Computing* (RISC) processors and many of the programming strategies that give efficient code are generic to this type of device.

As with many RISC processors, ARM processors are designed to access aligned data, that is, words that lie on addresses that are multiples of four, and halfwords that lie on addresses that are multiples of two. This data is located on its natural size boundary.

ARM compilers normally align global variables to these natural size boundaries so that these items can be accessed efficiently using the LDR and STR instructions.

This contrasts with most *Complex Instruction Set Computing* (CISC) architectures where instructions are available to directly access unaligned data. Therefore, you must take care when porting legacy code from CISC architectures to the ARM processors. In particular, accesses to unaligned data can be expensive in code size or performance.

Note

ARM11 processors support unaligned accesses in hardware. This section mainly applies to ARM processors earlier than the ARM11 processor family.

The following sections discuss these programming issues in more detail:

- *Unaligned pointers*
- *Unaligned fields in structures* on page 1-4
- *Unaligned LDR for accessing halfwords* on page 1-6
- *Porting code and detecting unaligned accesses* on page 1-7.

1.2.1 Unaligned pointers

The C and C++ standards specify that a pointer to a type cannot be less aligned than the natural alignment of the type. This improves code size and performance. Therefore, by default, the ARM compiler expects normal C and C++ pointers to point to an aligned word in memory. A type qualifier `__packed` is provided to enable unaligned pointer access (see the section describing variable declaration keywords in the compiler reference in *RealView Compilation Tools v2.2 Compiler and Libraries Guide*).

For example, if the pointer `int *` is used to read a word, the ARM compiler uses an LDR instruction in the generated code. This works as expected when the address is a multiple of four (that is, on a word boundary). However, if the address is not a multiple of four, then an LDR instruction returns a rotated result rather than performing a true unaligned word load. The rotated result depends on the offset and endianness of the system.

If your code loads data from a pointer that points to the address 0x8006, for example, you might expect to load the contents of bytes from 0x8006, 0x8007, 0x8008, and 0x8009. However, on an ARM processor, this access loads the rotated contents of bytes from 0x8004, 0x8005, 0x8006, and 0x8007.

Therefore, if you want to define a pointer to a word that can be at any address (that is, that can be at a non-natural alignment), you must specify this using the `__packed` qualifier when defining the pointer:

```
__packed int *pi; // pointer to unaligned int
```

The ARM compiler does not then use an LDR, but generates code that correctly accesses the value regardless of the alignment of the pointer. This generated code is a sequence of byte accesses or, depending on the compile options, variable alignment-dependent shifting and masking. This approach, however, incurs a performance and code size penalty.

———— Note ————

You must not access memory-mapped peripheral registers using `__packed` because the ARM compiler can use multiple memory accesses to retrieve the data. Therefore, nearby locations can be accessed that might correspond to other peripheral registers. When bitfields are used, the ARM compiler currently accesses the entire container, not just the field specified.

1.2.2 Unaligned fields in structures

In the same way that global variables are located on their natural size boundary, so are the fields in a structure. This means that the compiler often has to insert padding between fields to ensure that fields are aligned. The compiler generates the following remark when it inserts padding:

```
#1301-D: padding inserted in struct mystruct
```

Use the `--remarks` compiler option to display remarks. Alternatively, specify the option `--diag_warning 1301` to raise the remark to a warning.

Sometimes, you might not want the compiler to insert padding. You can use the `__packed` qualifier to create structures without padding between fields. These structures require unaligned accesses.

If the ARM compiler knows the alignment of a particular structure, it can work out whether or not the fields it is accessing are aligned within a packed structure. In these cases, the compiler carries out the more efficient aligned word or halfword accesses,

where possible. Otherwise, the compiler uses multiple aligned memory accesses (LDR, STR, LDM, and STM) combined with fixed shifting and masking to access the correct bytes in memory.

Whether these accesses to unaligned elements are done inline or by calling a function is controlled by using the compiler options `-Ospace` (default, calls a function) and `-Otime` (do unaligned access inline).

For example:

1. Create a file `foo.c` that contains:

```
__packed struct mystruct {
    int aligned_i;
    short aligned_s;
    int unaligned_i;
};
struct mystruct S1;

int foo (int a, short b)
{
    S1.aligned_i=a;
    S1.aligned_s=b;
    return S1.unaligned_i;
}
```

2. Compile this using `armcc -c -Otime foo.c`. The code produced is:

```
MOV     r2,r0
LDR     r0,[L1.84]
MOV     r12,r2,LSR #8
STRB    r2,[r0,#0]
STRB    r12,[r0,#1]
MOV     r12,r2,LSR #16
STRB    r12,[r0,#2]
MOV     r12,r2,LSR #24
STRB    r12,[r0,#3]
MOV     r12,r1,LSR #8
STRB    r1,[r0,#4]
STRB    r12,[r0,#5]
ADD     r0,r0,#6
BIC     r3,r0,#3
AND     r0,r0,#3
LDMIA   r3,{r3,r12}
MOV     r0,r0,LSL #3
MOV     r3,r3,LSR r0
RSB     r0,r0,#0x20
ORR     r0,r3,r12,LSL r0
BX      lr
```

However, you can give the compiler more information to enable it to know which fields are aligned and which are not. To do this you must declare non-aligned fields as `__packed`, and remove the `__packed` attribute from the **struct** itself.

This is the recommended approach, and the only way of guaranteeing fast access to naturally aligned members within the **struct**.

It is also clearer which fields are non-aligned, but care is needed when adding or deleting fields from the **struct**.

3. Now modify the definition of the structure in `foo.c` to:

```
struct mystruct {
    int aligned_i;
    short aligned_s;
    __packed int unaligned_i;
};
struct mystruct S1;
```

4. Compile `foo.c` and the following, more efficient code, is generated:

```
MOV     r2,r0
LDR     r0,[L1.32|
STR     r2,[r0,#0]
STRH    r1,[r0,#4]
LDMIB   r0,{r3,r12}
MOV     r0,r3,LSR #16
ORR     r0,r0,r12,LSL #16
BX      lr
```

The same principle applies to unions. Use the `__packed` attribute on the components of the union that will be unaligned in memory.

———— **Note** ————

Any `__packed` object accessed through a pointer has unknown alignment, even packed structures.

1.2.3 Unaligned LDR for accessing halfwords

In some circumstances the ARM compiler can intentionally generate unaligned LDR instructions. In particular, the compiler does this to load halfwords from memory. This is because by using an appropriate address the required halfword can be loaded into the top half of a register and then shifted down to the bottom half. This requires only one memory access.

Performing the same operation using LDRB instructions requires two memory accesses, plus instructions to merge the two bytes. On ARMv3 and earlier, this is typically done for any halfword loads. On ARMv4 and later, this is done less often because dedicated

halfword load instructions exist, but unaligned LDR instructions might still be generated. For example, to access an unaligned **short** within a packed structure. Such unaligned LDR instructions are only generated by the compiler if you enable them using the `--memaccess +L41` compiler option.

Note

The ARMv3 architecture is obsolete and is no longer supported by RVCT.

1.2.4 Porting code and detecting unaligned accesses

Legacy C code for other architectures (for example, x86 CISC) might perform accesses to unaligned data using pointers that do not work on ARM processors. This is non-portable code, and such accesses must be identified and corrected to work on RISC architectures, which expect aligned data.

Identifying the unaligned accesses can be difficult, because the use of load or store operations with unaligned addresses gives incorrect behavior. It is difficult to trace which part of the C source is causing the problem.

ARM processors with full *Memory Management Units* (MMUs), for example, the ARM920T™, support optional alignment checking, where the processor checks every access to ensure it is correctly aligned. The MMU raises a Data Abort if an incorrectly aligned access occurs.

For simple cores such as the ARM7TDMI®, it is recommended that alignment-checking be implemented within the *Application Specific Integrated Circuit* (ASIC) or *Application Specific Standard Product* (ASSP). You can do this with an additional hardware block that is external to the ARM core, and that monitors the access size and the least significant bits of the address bus for every data access. You can configure the ASIC/ASSP to raise the **ABORT** signal in the case of an unaligned access. ARM Limited recommends that such logic is included on ASIC/ASSP devices where code is ported from other architectures.

If the system is configured to abort on unaligned accesses, a Data Abort exception handler must be installed. When an unaligned access occurs, the Data Abort handler is entered, and this can identify the erroneous data access instruction, which is located at (r14-8).

When identified, you must fix the data access by changing the C source. These changes can be made conditional using the following:

```
#ifdef __arm
#define PACKED __packed
#else
#define PACKED
```

```
#endif  
...  
    PACKED int *pi;  
...
```

It is best to minimize accesses to unaligned data because of code size and performance overheads.

See the description of the `--pointer_alignment` and `--min_array_alignment` options in the section on controlling code generation of *RealView Compilation Tools v2.2 Compiler and Libraries Guide*.

1.3 Developing for the ARM processors

This book gives information and example code for some of the most common ARM programming tasks, and includes information for developers working on ARM architectures:

- *Embedded software development*
- *Interworking ARM and Thumb code* on page 1-10
- *Mixing C, C++, and assembly language* on page 1-10
- *Handling processor exceptions* on page 1-11
- *Using the AAPCS* on page 1-12
- *Compatibility with legacy objects and libraries* on page 1-12.

1.3.1 Embedded software development

Many applications written for ARM architecture-based systems are embedded applications that are contained in ROM and execute on reset. There are a number of factors that you must consider when writing embedded operating systems, or embedded applications that execute from reset without an operating system, including:

- address remapping, for example initializing with ROM at address 0x0000, then remapping RAM to address 0x0000
- initializing the environment and application
- linking an embedded executable image to place code and data in specific locations in memory.

The ARM core usually begins executing instructions from address 0x0000 at reset. For an embedded system, this means that there must be ROM at address 0x0000 when the system is reset. Typically, however, ROM is slow compared to RAM, and often only 8 or 16 bits wide. This affects the speed of exception handling. Having ROM at address 0x0000 means that the exception vectors cannot be modified. A common strategy is to remap ROM to RAM and copy the exception vectors from ROM to RAM after startup. See *ROM/RAM remapping* on page 2-28 for more information.

After reset, an embedded application or operating system must initialize the system, including:

- initializing the execution environment, such as exception vector, stacks, and I/O peripherals
- initializing the application, for example copying initial values of nonzero writable data to the writable data region and zeroing the ZI data region.

See *Initialization sequence* on page 2-26 for more information.

Embedded systems often implement complex memory configurations. For example, an embedded system might use fast, 32-bit RAM for performance-critical code, such as interrupt handlers and the stack, slower 16-bit RAM for application RW data, and ROM for normal application code. You can use the linker scatter-loading mechanism to construct executable images suitable for complex systems. For example, a scatter-load description file can specify the load address and execution address of individual code and data regions. See Chapter 2 *Embedded Software Development* for a series of worked examples, and for information on other issues that affect embedded applications, such as semihosting.

1.3.2 Interworking ARM and Thumb code

If you are writing code for ARM processors that support the Thumb 16-bit instruction set, you can mix ARM and Thumb code as required. If you are writing C or C++ code you must compile with the `--apcs /interwork` option. The linker detects when an ARM function is called from Thumb state, or a Thumb function is called from ARM state and alters call and return sequences, or inserts interworking veneers to change processor state as necessary.

Note

If you want to use absolute addresses to Thumb functions, see *Pointers to functions in Thumb state* on page 4-16.

If you are writing assembly language code you must ensure that you comply with the interworking variant of the *Procedure Call Standard for the ARM Architecture* (AAPCS). There are several ways to change processor state, depending on the target architecture version. See Chapter 4 *Interworking ARM and Thumb* and Appendix A *Using the Procedure Call Standard* for more information.

1.3.3 Mixing C, C++, and assembly language

You can mix separately compiled and assembled C, C++, and ARM assembly language modules in your program. You can write small assembly language routines within your C or C++ code. These routines are compiled using the inline or embedded assembler of the ARM compiler. However, there are a number of restrictions to the assembly language code you can write if you are using the inline or embedded assembler. These restrictions are described in the chapter on inline and embedded assemblers in *RealView Compilation Tools v2.2 Compiler and Libraries Guide*.

In addition, Chapter 5 *Mixing C, C++, and Assembly Language* gives general guidelines and examples of how to call between C, C++, and assembly language modules.

1.3.4 Handling processor exceptions

The ARM processor recognizes the following exception types:

Reset Occurs when the processor reset pin is asserted. This exception is only expected to occur for signaling power-up, or for resetting as if the processor has powered up. A soft reset can be done by branching to the reset vector, 0x0000.

Undefined Instruction

Occurs if neither the processor, nor any attached coprocessor, recognizes the currently executing instruction.

Software Interrupt (SWI)

This is a user-defined interrupt instruction. It enables a program running in User mode, for example, to request privileged operations that run in Supervisor mode, such as an RTOS function.

Prefetch Abort

Occurs when the processor attempts to execute an instruction that has been prefetched from an illegal address. An illegal address is one at which memory does not exist, or one that the memory management subsystem has determined is inaccessible to the processor in its current mode.

Data Abort Occurs when a data transfer instruction attempts to load or store data at an illegal address.

Interrupt (IRQ)

Occurs when the processor external interrupt request pin is asserted (LOW) and IRQ interrupts are enabled (the I bit in the CPSR is clear).

Fast Interrupt (FIQ)

Occurs when the processor external fast interrupt request pin is asserted (LOW) and FIQ interrupts are enabled (the F bit in the CPSR is clear). This exception is typically used where interrupt latency must be kept to a minimum.

In general, if you are writing an application such as an embedded application that does not rely on an operating system to service exceptions, you must write handlers for each exception type.

In cases where an exception type can have more than one source, for example SWI or IRQ interrupts, you can chain exception handlers for each source. See *Chaining exception handlers* on page 6-44 for more information.

On processors that support Thumb instructions, the processor switches to ARM state when an exception is taken. You can either write your exception handler in ARM code, or use a veneer to switch to Thumb state. See *The return address and return instruction* on page 6-10 for more information.

1.3.5 Using the AAPCS

The *Procedure Call Standard for the ARM Architecture* (AAPCS) defines register usage and stack conventions that must be followed to enable separately compiled and assembled modules to work together. There are a number of variants on the base standard. The ARM compiler always generates code that conforms to the selected AAPCS variant. The linker selects an appropriate standard C or C++ library to link with, if required.

When developing code for ARM processors, you must select an appropriate AAPCS variant, for example:

- if you are writing code that interworks between ARM and Thumb state you must select the `--apcs /interwork` option in the compiler and assembler
- if you are writing code in C or C++, you must ensure that you have selected compatible AAPCS options for each compiled module
- if you are writing your own assembly language routines, you must ensure that you conform to the appropriate AAPCS variant.

For more information, see:

- the *Procedure Call Standard for the ARM Architecture* specification, `aapcs.pdf`, in `install_directory\Documentation\Specifications\...`
- Appendix A *Using the Procedure Call Standard*.

———— Note ————

If you are mixing C and assembly language, ensure that you understand the AAPCS implications.

1.3.6 Compatibility with legacy objects and libraries

If you are upgrading to RVCT from a previous release, ensure that you read Appendix A in *RealView Compilation Tools v2.2 Essentials Guide* for details about compatibility between the new release and previous releases of RVCT.

1.4 ARM architecture v6 support

All components of RVCT support ARMv6. `armasm` accepts all ARMv6 instructions, `armlink` can link-in ARMv6 library objects where required, and `fromelf` disassembles the ARMv6 instructions correctly. The embedded assembler of the compiler supports all ARMv6 instructions. The inline assembler supports the majority of ARMv6 instructions.

To compile code for ARMv6 use:

- `--cpu 6` for generic ARMv6 support

To compile code for a specific ARMv6 processor, use the processor name. For example:

- `--cpu ARM1136J-S` to generate code targeted at the ARM1136J-S
- `--cpu ARM1136JF-S` to generate code targeted at the ARM1136JF-S, that includes *Vector Floating Point* (VFP) hardware.

This section includes:

- *Instruction generation*
- *Alignment support* on page 1-14
- *Endian support* on page 1-14
- *Example 1 - instructions generated* on page 1-16
- *Example 2 - instructions generated for packed structures* on page 1-16.

1.4.1 Instruction generation

When compiling code for ARMv6, the compiler generates sign-extend and zero-extend instructions (for example, `SEXT8`), where appropriate (see *Example 1 - instructions generated* on page 1-16). Code scheduling for the specified processor is performed.

In addition, the C libraries contain some functions that are optimized specifically for ARMv6, such as `memcpy()`, `memmove()`, and `strcmp()`.

In RVCT v2.2, the compiler does not make use of SIMD instructions, because these do not map well onto C expressions. The endian reversal instructions (`REV`, `REV16` and `REVSH`) are generated by the compiler if it can deduce that a C expression performs an endian reversal.

1.4.2 Alignment support

By default the compiler utilizes ARMv6 unaligned access support to speed up access to packed structures, by enabling an LDR (or STR) to load from (or store to) a non-word aligned address (see *Example 2 - instructions generated for packed structures* on page 1-16). Structures remain non-packed unless explicitly qualified with `__packed`.

———— Note ————

Code compiled for ARMv6 only runs correctly if you enable unaligned support on the ARM core. You must do this by setting the U bit (bit 22) of CP15 register 1 in your initialization code, or by tying the **UBITINIT** input to the core HIGH.

Code that uses the pre-ARMv6 unaligned accesses behavior can be generated by using the compiler option:

```
--memaccess -UL41
```

1.4.3 Endian support

The ARM compiler has options for producing either little-endian or big-endian objects. ARMv6 supports two different big-endian modes:

- BE8** Specifies ARMv6 Byte Invariant Addressing mode. This produces little-endian code and big-endian data. This is the default Byte Addressing mode for ARMv6 big-endian images.
Byte Invariant Addressing mode is only available on ARM processors that support ARMv6.
- BE32** This is legacy big-endian mode. It produces big-endian code and data. It is identical to the big-endian mode supported prior to ARMv6. This is the default Byte Addressing mode for all pre-ARMv6 big-endian images.

When compiling for ARMv6 big endian, the ARM compiler generates big-endian objects as BE8 rather than BE32. A flag, set in the object code, labels the code as BE8. Therefore, you must enable BE8 support in the ARM core by setting the E-bit in the CPSR.

You can link legacy objects (for example, ARMv4T) with ARMv6 objects (for running on ARMv6), but in this case the linker switches the byte order of the legacy object code into BE8 mode. The resulting image is BE8.

If you want to use the legacy BE32 mode, then you must set the B bit (bit 7) of CP15 register 1 in your initialization code, or tie the **BIGENDINIT** input into the core HIGH.

You can then generate BE32-compatible code by using the following compiler option:

`--memaccess -UL41`

BE32-compatible code must also be linked using the linker option `--BE32`. Otherwise, the ARMv6 attribute of the objects cause a BE8 image to be produced.

1.4.4 Example 1 - instructions generated

This example shows the different instructions generated when compiling for ARMv6 and earlier architectures.

```
signed char unpack(int i)
{
    return (signed char)i;
}
```

Pre-ARMv6 architecture compilations

Compiling with `--cpu 5` gives:

```
unpack PROC
    MOV     r0,r0,LSL #24
    MOV     r0,r0,ASR #24
    BX      lr
    ENDP
```

ARMv6 architecture compilations

Compiling with `--cpu 6` gives:

```
unpack PROC
    SEXT8   r0,r0
    BX      lr
    ENDP
```

1.4.5 Example 2 - instructions generated for packed structures

This example shows the different instructions generated for a packed structure when compiling for ARMv6 and earlier architectures.

```
__packed struct{
    char ch;
    short sh;
    int i;
} foo;

signed char unpack()
{
    return (signed char)foo.i;
}
```

Pre-ARMv6 architecture compilations

Compiling with `--cpu 5` gives:

```
unpack PROC
    STMFD    sp!, {r3,lr}
    LDR      r0, |L1.24|
    BL       __rt_uread4
    MOV      r0, r0, LSL #24
    MOV      r0, r0, ASR #24
    LDMFD    sp!, {r3,pc}
    |L1.24|
    DCD      ||.bss$2|| + 3
    ENDP
```

ARMv6 architecture compilations

Compiling with `--cpu 6` gives:

```
unpack PROC
    LDR      r0, |L1.16|
    LDR      r0, [r0, #3]
    SEXT8    r0, r0
    BX       lr
    |L1.16|
    DCD      ||.bss$2||
    ENDP
```


Chapter 2

Embedded Software Development

This chapter describes how to develop embedded applications with *RealView® Compilation Tools (RVCT)*, with or without a target system present. It contains the following sections:

- *About embedded software development* on page 2-2
- *Default compilation tool behavior in the absence of a target system* on page 2-4
- *Tailoring the C library to your target hardware* on page 2-11
- *Tailoring the image memory map to your target hardware* on page 2-15
- *Reset and initialization* on page 2-25
- *Further memory map considerations* on page 2-35.

2.1 About embedded software development

Most embedded applications are initially developed in a prototype environment with resources that differ from those available in the final product. Therefore, it is important to consider the processes involved in moving an embedded application from one that relies on the facilities of the development or debugging environment to a system that runs standalone on target hardware.

When developing embedded software using RVCT, you must consider the following:

- How the C library uses hardware.
- Some C library functionality executes by using debug environment resources. If used, you must re-implement this functionality to make use of target hardware.
- RVCT has no inherent knowledge of the memory map of any given target. You must tailor the image memory map to the memory layout of the target hardware.
- An embedded application must perform some initialization before the main application can be run. A complete initialization sequence requires code that you implement as well as RVCT C library initialization routines.

2.1.1 Example code

To illustrate the topics covered in this chapter, associated example projects are provided. The code for the Dhrystone builds described in this chapter is in the main examples directory, in `...\emb_sw_dev`. Each build is in a separate directory, and provides an example of the techniques discussed in successive sections of this chapter. Specific information regarding each build can be found in:

- *Example code for Build 1* on page 2-10
- *Example code for Build 2* on page 2-13
- *Example code for Build 3* on page 2-23
- *Example code for Build 4* on page 2-33
- *Example code for Build 5* on page 2-41.

The Dhrystone benchmarking program provides the code base for the example projects. Dhrystone was chosen because it enables many of the concepts described in this chapter to be illustrated.

The example projects are tailored to run on the ARM® Integrator™ development platform. However, the principles illustrated by the examples apply to any target hardware.

Note

The focus of this chapter is not specifically the Dhrystone program, but the steps that must be taken to enable it to run on a fully standalone system. For further discussion of Dhrystone as a benchmarking tool, see Application Note 93 - *Benchmarking with ARMulator®*. You can find the ARM Application Notes in the Documentation area of the ARM website at <http://www.arm.com>.

Running the Dhrystone builds on an Integrator

To run the Dhrystone builds described in this chapter on an Integrator, you must:

- Perform ROM/RAM remapping. To achieve this, run the Boot Monitor by setting switches 1 and 4 to ON, and then reset the board.
- Set `top_of_memory` to `0x40000`, or fit a DIMM memory module. If this is not done, the stack, with a default setting of `0x80000`, might not be in valid memory.

2.2 Default compilation tool behavior in the absence of a target system

When you start work on software for an embedded application, you might not be aware of the full technical specifications of the target hardware. For example, you might not know the details of target peripheral devices, the memory map, or even the processor itself.

To enable you to proceed with software development before such details are known, the compilation tools have a default behavior that enables you to start building and debugging application code immediately. It is useful to be aware of this default behavior, so that you appreciate the steps necessary to move from a default build to a fully standalone application.

This section includes:

- *Semihosting* on page 2-5
- *C library structure* on page 2-6
- *Default memory map* on page 2-7
- *Linker placement rules* on page 2-8
- *Application startup* on page 2-9
- *Example code for Build 1* on page 2-10.

2.2.1 Semihosting

In the ARM C Library, support for some ISO C functionality is provided by the host debugging environment. The mechanism that provides this functionality is known as *semihosting*.

Semihosting is implemented by a set of defined software interrupt (SWI) operations. When a semihosting SWI is executed, the debug agent identifies it and briefly suspends program execution. The semihosting operation is then serviced by the debug agent before code execution is resumed. Therefore, the task performed by the host itself is transparent to the program.

Figure 2-1 shows an example of semihosting operation that prints a string to the debugger console.

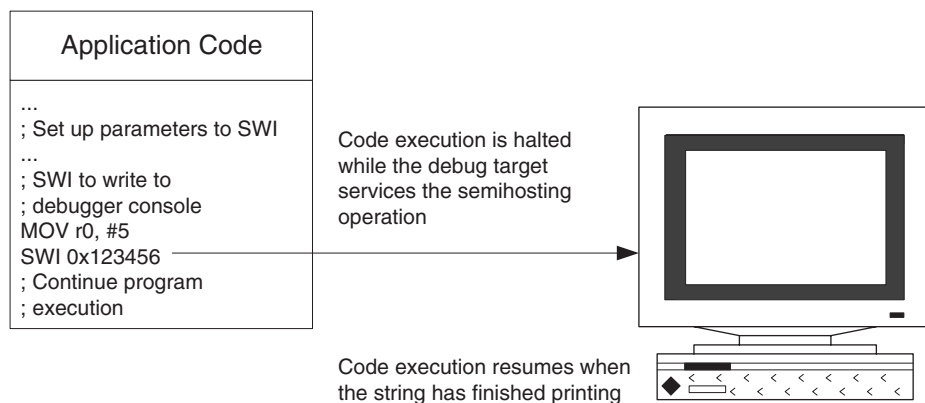


Figure 2-1 Example semihosting operation

Note

For more information, see the chapter describing semihosting in *RealView Compilation Tools v2.2 Compiler and Libraries Guide*.

2.2.2 C library structure

Conceptually, the C library can be divided into functions that are part of the ISO C Language specification and functions that provide support to the ISO C language specification. This is shown in Figure 2-2.

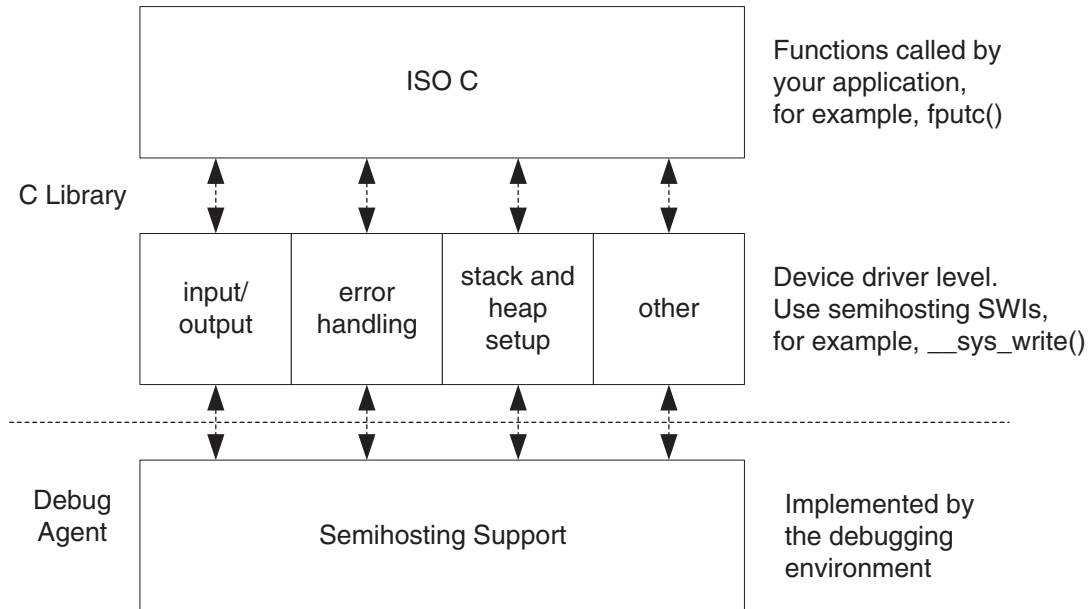


Figure 2-2 C library structure

Support for some ISO C functionality is provided by the host debugging environment at the device driver level.

For example, the RVCT C library implements the ISO C `printf()` family of functions by writing to the debugger console window. This functionality is provided by calling `__sys_write()`. This is a support function that executes a semihosting SWI, resulting in a string being written to the console.

2.2.3 Default memory map

In an image where you have not described the memory map, RVCT places code and data according to a default memory map, as shown in Figure 2-3.




Figure 2-3 Default memory map

The default memory map can be described as follows:

- The image is linked to load and run at address **0x8000**. All RO (Read Only) sections are placed first, followed by RW (Read-Write) sections, then ZI (Zero Initialized) sections.
- The heap follows directly on from the top of the ZI section, so the exact location is decided at link time.
- The stack base location is provided by a semihosting operation during application startup. The value returned by this semihosting operation depends on the debug environment:
 - RealView ARMulator ISS (RVISS) returns the value set in the configuration file `peripherals.ami`. The default is `0x08000000`.
 - Multi-ICE® and RealView ICE return the value of the debugger internal variable `top_of_memory`. The default is `0x00080000`.

2.2.4 Linker placement rules

The linker observes a set of rules, shown in Figure 2-4, to decide where in memory code and data is located.

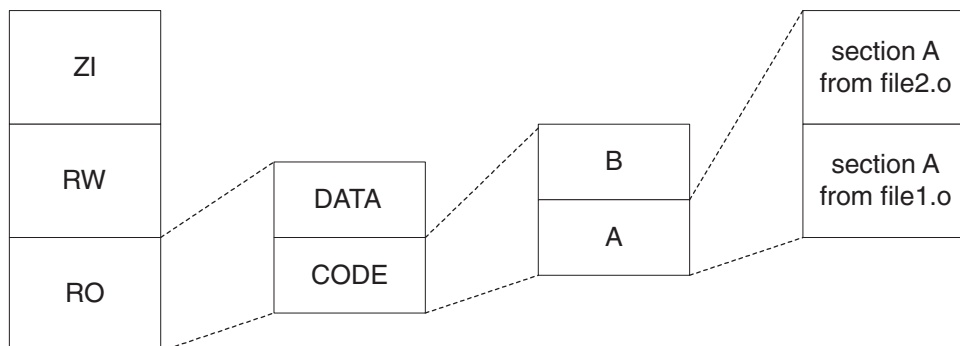


Figure 2-4 Linker placement rules

The image is organized first of all by attribute, with RO at the lowest memory address, then RW, then ZI. Within each attribute code precedes the data.

From there, the linker places input sections alphabetically by name. Input section names correspond with assembler AREA directives.

In input sections, code and data from individual objects are placed according to the order of object files given on the linker command line.

ARM does not recommend relying on these rules for precise placement of code and data. Instead, you must use the scatter-loading mechanism for full control of placement of code and data. See *Tailoring the image memory map to your target hardware* on page 2-15.

————— **Note** —————

See *RealView Compilation Tools v2.2 Linker and Utilities Guide* for more information on placement rules and scatter-loading.

2.2.5 Application startup

In most embedded systems, an initialization sequence executes to set up the system before the main task is executed.

Figure 2-5 shows the default initialization sequence.

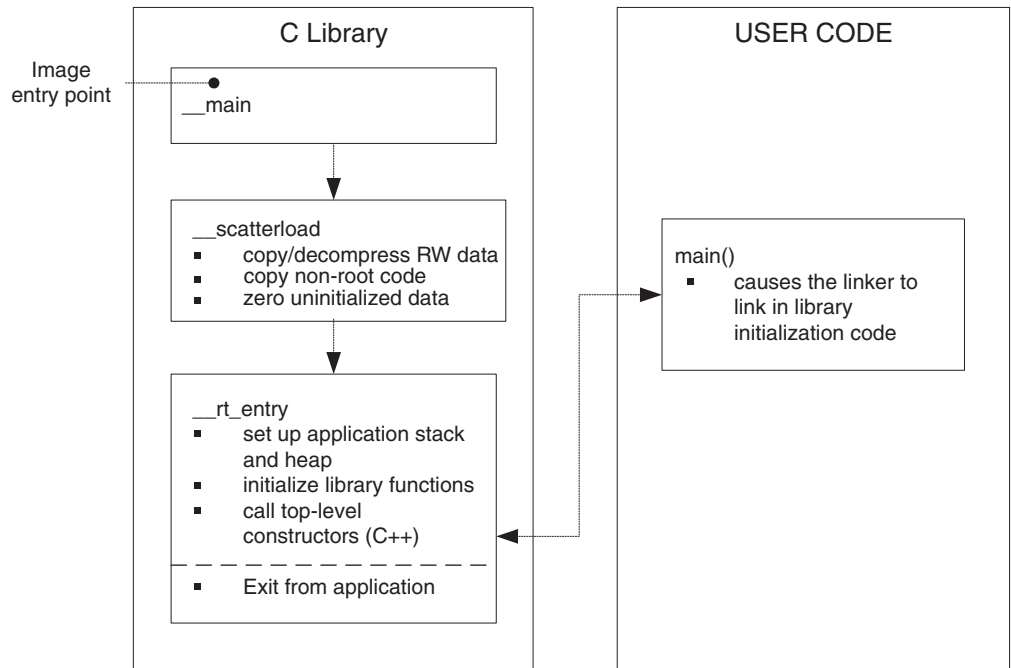


Figure 2-5 Default initialization sequence

At a high level, the initialization sequence can be divided into three functional blocks. **__main** branches directly to **__scatterload**. **__scatterload** is responsible for setting the runtime image memory map, whereas **__rt_entry** (runtime entry) is responsible for initializing the C library.

__scatterload carries out code and data copying, decompression of RW data if necessary, and zeroing of ZI data.

__scatterload branches to **__rt_entry**. This sets up the application stack and heap, initializes library functions and their static data, and calls any constructors of globally declared objects (C++ only).

`__rt_entry` then branches to `main()`, the entry to your application. When the main application has finished executing, `__rt_entry` shuts down the library, then hands control back to the debugger.

The function label `main()` has a special significance. The presence of a `main()` function forces the linker to link in the initialization code in `__main` and `__rt_entry`. Without a function labeled `main()` the initialization sequence is not linked in, and as a result, some standard C library functionality is not supported.

2.2.6 Example code for Build 1

Build 1 is a default build of the Dhrystone benchmark. Therefore, it adheres to the default RVCT behavior described in this section. See *Running the Dhrystone builds on an Integrator* on page 2-3, and the example build files in the main examples directory, in `...\emb_sw_dev\build1`.

2.3 Tailoring the C library to your target hardware

By default the C library makes use of semihosting to provide device driver level functionality, enabling a host computer to act as an input and an output device. This is useful because development hardware often does not have all the input and output facilities of the final system.

This section includes:

- *Retargeting the C library*
- *Avoiding C library semihosting* on page 2-12
- *Example code for Build 2* on page 2-13.

2.3.1 Retargeting the C library

You can provide your own implementation of C Library functions that make use of target hardware, and that are automatically linked in to your image in favor of the C library implementations. This process, known as retargeting the C library, is shown in Figure 2-6.

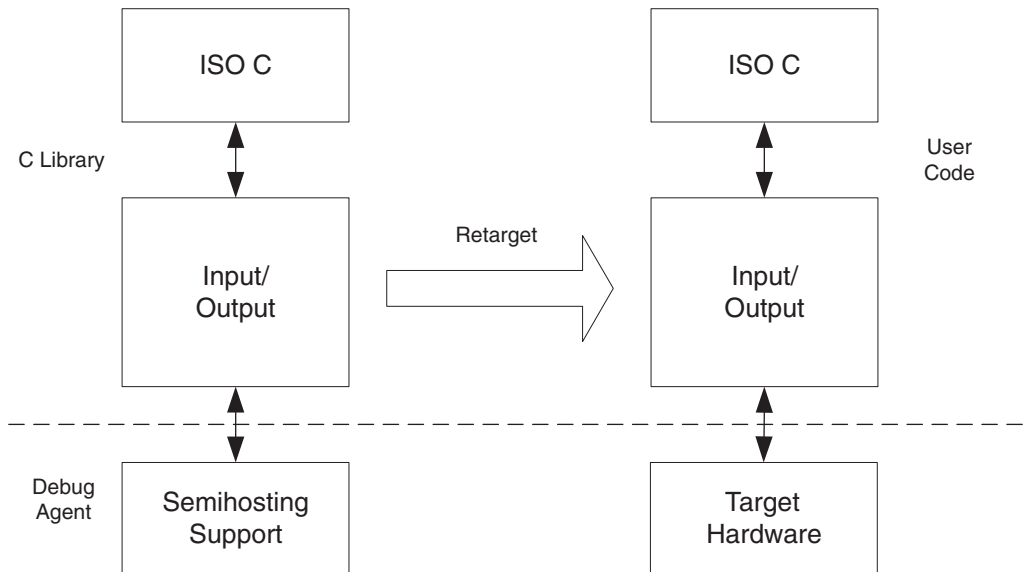


Figure 2-6 Retargeting the C library

For example, you might have a peripheral I/O device such as a UART, and you might want to override the library implementation of `fputc()`, that writes to the debugger console, with one that outputs to the UART. Because this implementation of `fputc()` is linked in to the final image, the entire `printf()` family of functions prints out to the UART.

Example 2-1 shows an example implementation of `fputc()`. The example redirects the input character parameter of `fputc()` to a serial output function `sendchar()` that is assumed to be implemented in a separate source file. In this way, `fputc()` acts as an abstraction layer between target dependent output and the C library standard output functions.

Example 2-1 Implementation of `fputc()`

```
extern void sendchar(char *ch);

int fputc(int ch, FILE *f)
{
    /* e.g. write a character to an UART */
    char tempch = ch;
    sendchar(&tempch);
    return ch;
}
```

2.3.2 Avoiding C library semihosting

In a standalone application, you are unlikely to support semihosting SWI operations. Therefore, you must be certain that no C library semihosting functions are being linked into your application.

To ensure that no functions that use semihosting SWIs are linked in from the C library, you must import the symbol `__use_no_semihosting_swi`. This can be done in any C or assembler source file in your project as follows:

- In a C module, use the `#pragma` directive:
`#pragma import(__use_no_semihosting_swi)`
- In an assembler module, use the `IMPORT` directive:
`IMPORT __use_no_semihosting_swi`

If functions that use semihosting SWIs are still being linked in, the linker reports the following error:

Error: L6200E: Symbol __semlhosting_swi_guard multiply defined (by use_semi.o and use_no_semi.o).

To identify these functions, link using the --verbose option. In the resulting output, C library functions are tagged with __I_use_semlhosting_swi, for example:

```
Loading member sys_exit.o from c_a__un.l.
      definition:  _sys_exit
      reference :  __I_use_semlhosting_swi
```

You must provide your own implementations of these functions (_sys_exit in this example).

Note

The linker does not report any semihosting SWI-using functions in your application code. An error only occurs if this type of function is linked in from the C library.

For a full list of C library functions that use semihosting SWIs, see the chapter describing semihosting in *RealView Compilation Tools v2.2 Compiler and Libraries Guide*.

2.3.3 Example code for Build 2

Build 2 of the Dhrystone benchmark uses the hardware of the Integrator platform for clocking and string I/O. See the example build files in the main examples directory, in ...\\emb_sw_dev\\build2.

The following changes have been made to Build 1 of the example project:

C Library Retargeting

A retargeted layer of ISO C functions has been added. These include standard I/O functions and clock functionality, as well as some additional error signaling and program exit.

Target Dependent Device Driver

A device driver layer has been added that interacts directly with target hardware peripherals.

See *Running the Dhrystone builds on an Integrator* on page 2-3.

The symbol __use_no_semlhosting_swi is not imported into this project. This is because a semihosting SWI is executed during C library initialization to set up the application stack and heap location. Retargeting stack and heap setup is described in detail in *Placing the stack and heap* on page 2-20.

Note

To see the output, a terminal or terminal emulator must be connected to serial port A. The serial port settings must be set to 38400 baud, no parity, 1 stop bit and no flow control. The terminal must be configured to append line feeds to incoming line ends, and echo typed characters locally.

2.4 Tailoring the image memory map to your target hardware

In your final embedded system, without semihosting functionality, you are unlikely to use the default memory map provided by RVCT. Your target hardware usually has several memory devices located at different address ranges. To make the best use of these devices, you must have separate views of memory at load and runtime.

This section includes:

- *Scatter-loading*
- *Scatter-loading description file syntax* on page 2-16
- *Scatter-loading description file example* on page 2-17
- *Placing objects in a scatter-loading description file* on page 2-18
- *Root regions* on page 2-19
- *Placing the stack and heap* on page 2-20
- *Runtime memory models* on page 2-21
- *Example code for Build 3* on page 2-23.

2.4.1 Scatter-loading

Scatter-loading enables you to describe the load-time and runtime location of code and data in memory in a textual description file known as a *scatter-loading description file*. The file is passed to the linker on the command line using the `--scatter` option. For example:

```
armlink --scatter scat.txt file1.o file2.o
```

The scatter-loading description file describes to the linker the desired location of code and data at both load-time and runtime, in terms of addressed memory regions.

Scatter-loading regions

Scatter-loading regions fall into two categories:

- Load Regions that contain application code and data at reset and load-time.
- Execution Regions that contain code and data while the application is executing. One or more execution regions are created from each load region during application startup.

All code and data in the image falls into exactly one load region and one execution region.

During startup, C library initialization code in `__main` carries out the copying and zeroing of code and data necessary to move from the image load view to the execute view.

2.4.2 Scatter-loading description file syntax

The scatter-loading description file syntax reflects the functionality provided by scatter-loading itself. Figure 2-7 shows the file syntax.

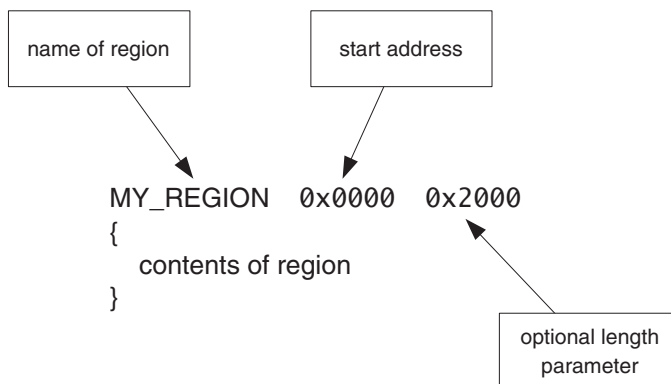


Figure 2-7 Scatter-loading description file syntax

A region is defined by a header tag that contains, as a minimum, a name for the region and a start address. Optionally, a maximum length and various attributes can be added.

The contents of the region depend on the type of region:

- Load regions must contain at least one execution region. In practice, there are usually several execution regions for each load region.
- Execution regions must contain at least one code or data section, unless a region is declared with the `EMPTY` attribute (see *Using the scatter file EMPTY attribute* on page 2-39). Non-EMPTY regions usually contain source or library object files. The wildcard (*) syntax can be used to group all sections of a given attribute not specified elsewhere in the scatter-loading description file.

———— Note ————

For a more detailed description of scatter-loading description file syntax, see *RealView Compilation Tools v2.2 Linker and Utilities Guide*.

2.4.3 Scatter-loading description file example

Figure 2-8 shows a simple example of scatter-loading.

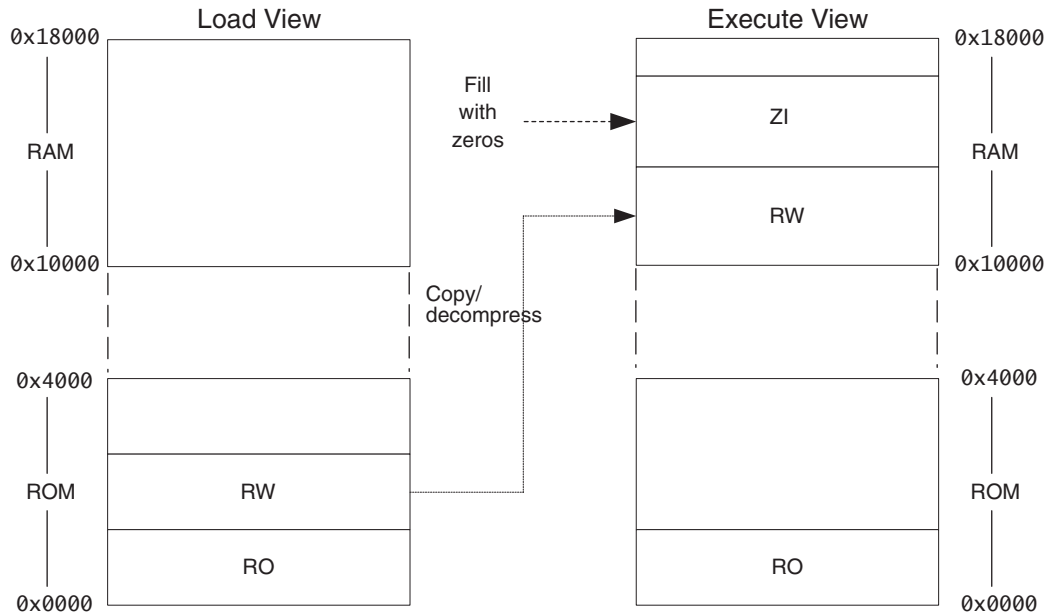


Figure 2-8 Simple scatter-loading example

This example has one load region containing all code and data, starting at address 0x0000. From this load region two execution regions are created. One contains all RO code and data that executes at the same address where it is loaded. The other is at address 0x10000, and contains all RW and ZI data.

Example 2-2 shows the description file that describes the memory map given in Figure 2-8.

Example 2-2 Simple scatter-loading description file

```
ROM_LOAD 0x0000 0x4000
{
    ROM_EXEC 0x0000 0x4000; Root region
    {
        * (+R0); All code and constant data
    }
}
```

```

RAM 0x10000 0x8000
{
    * (+RW, +ZI); All non-constant data
}

```

2.4.4 Placing objects in a scatter-loading description file

For most images, you control the placement of specific code and data sections, rather than grouping all attributes together as in Example 2-2 on page 2-17. You can do this by specifying individual objects directly in the description file, instead of relying only on the wildcard syntax.

————— Note —————

The ordering of objects in a description file execution region does not affect how they are ordered in the output image. The linker placement rules described in *Linker placement rules* on page 2-8 apply to each execution region.

To override the standard linker placement rules, you can use the +FIRST and +LAST scatter-loading directives. Example 2-3 shows a scatter-loading description file that places the vector table at the beginning of an execution region. In this example, the area Vect in vectors.o is placed at address 0x0000.

Example 2-3 Placing a section

```

ROM_LOAD 0x0000 0x4000
{
    ROM_EXEC 0x0000 0x4000
    {
        vectors.o (Vect, +FIRST)
        * (+R0)
    }
    ; more exec regions...
}

```

See *RealView Compilation Tools v2.2 Linker and Utilities Guide* for further information on placing objects in scatter-loading description files.

2.4.5 Root regions

A *root region* is an execution region with a load address that is the same as its execution address. Each scatter-loading description file must have at least one root region.

One restriction placed on scatter-loading is that the code and data responsible for creating execution regions, for example, copying and zeroing code and data, cannot itself be copied to another location. As a result, the following sections must be included in a root region:

- `__main.o` and `__scatter*.o` containing the code that copies code and data
- `__dc*.o` that performs decompression
- `Region$$Table` section containing the addresses of the code and data to be copied or decompressed.

However, these can be described using `InRoot$$Sections`.

Because these sections are defined as read-only, they are grouped by the `*` `(+R0)` wildcard syntax. As a result, if `*` `(+R0)` is specified in a non-root region, these sections must be explicitly declared in a root region. This is shown in Example 2-4.

Example 2-4 Specifying a root region

```

ROM_LOAD 0x0000 0x4000
{
    ROM_EXEC 0x0000 0x4000      ; root region
    {
        vectors.o (Vect, +FIRST) ; Vector table
        * (InRoot$$Sections)      ; All library sections
                                   ; that must be in a root region
                                   ; for example, __main.o, __scatter*.o,
                                   ; __dc*.o and * Region$$Table
    }
    RAM 0x10000 0x8000
    {
        * (+R0, +RW, +ZI)        ; all other sections
    }
}

```

Failing to include `__main.o`, `__scatter.o`, `__dc*.o` and `Region$$Table` in a root region results in the linker generating an error message such as:

Error: L6202E: Section `Region$$Table` cannot be assigned to a non-root region.

2.4.6 Placing the stack and heap

Scatter-loading provides a method for specifying the placement of code and statically allocated data in your image. This section covers how to place the application stack and heap.

The application stack and heap are set up during C library initialization. You can tailor stack and heap placement by re-implementing the routine responsible for stack and heap setup. In the ARM C library, this routine is `__user_initial_stackheap()`.

Figure 2-9 shows the C library initialization process with a re-implemented `__user_initial_stackheap()`.

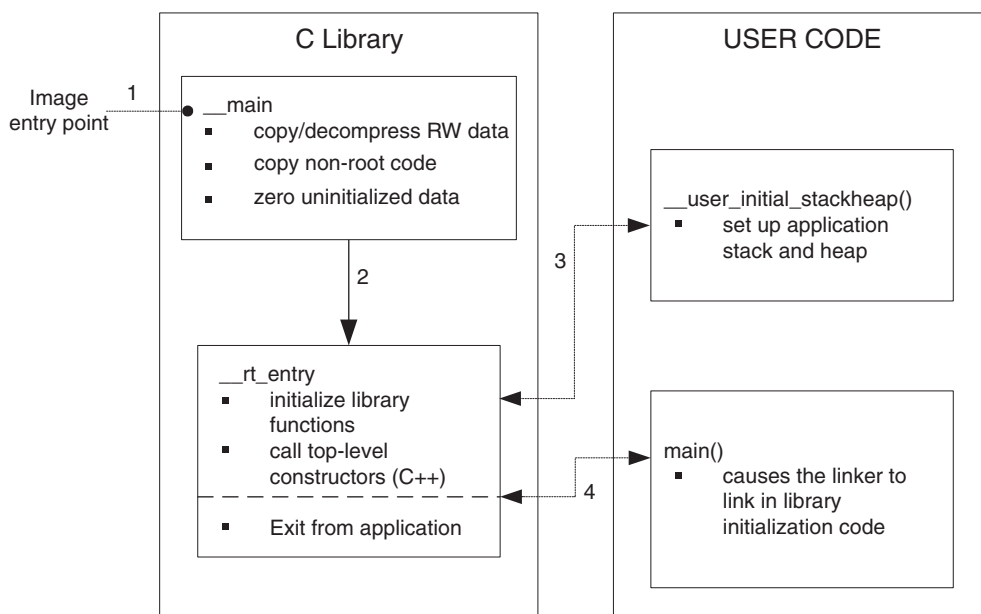


Figure 2-9 Retargeting `__user_initial_stackheap()`

`__user_initial_stackheap()` can be coded in C or ARM assembler. It must return the following parameters:

- heap base in `r0`
- stack base in `r1`
- heap limit in `r2`, if required
- stack limit in `r3`, if required.

You must re-implement `__user_initial_stackheap()` if you are scatter-loading your image. Otherwise, the linker generates the following error:

Error: L6218E: Undefined symbol Image\$\$ZI\$\$Limit (referred from sys_stackheap.o)

2.4.7 Runtime memory models

Two possible runtime memory models are provided:

- *One-region model*, the default
- *Two-region model* on page 2-22.

In both runtime memory models, the stack grows unchecked by default. You can choose to enable software stack checking in your image by compiling all modules with the compiler option `--apcs /swst`. If you are using a two-region model, you must also specify a stack limit in your implementation of `__user_initial_stackheap()`.

———— Note —————

Enabling software stack checking introduces a substantial code size and performance overhead, because the value of the stack pointer must be checked against the stack limit within each function that uses the stack. It also uses register `r10`, so is not generally recommended for embedded systems.

Both these examples are suitable for the Integrator system.

One-region model

In the default, *one-region model*, the application stack and heap grow towards each other in the same region of memory. In this case, the heap is checked against the value of the stack pointer when new heap space is allocated (for example, when `malloc()` is called).

Figure 2-10 on page 2-22 and Example 2-5 on page 2-22 show an example of `__user_initial_stackheap()` implementing a simple one-region model, where the stack grows downwards from address `0x40000`, and the heap grows upwards from `0x20000`.

The routine loads the appropriate values into the registers `r0` and `r1`, and then returns. Registers `r2` and `r3` remain unchanged, because a heap limit and stack limit are not used in a one-region model.




Figure 2-10 One-region model

Example 2-5 One-region model routine

```

EXPORT __user_initial_stackheap

__user_initial_stackheap
    LDR r0, =0x20000 ;HB
    LDR r1, =0x40000 ;SB
    ; r2 not used (HL)
    ; r3 not used (SL)
    MOV pc, lr

```

Two-region model

Your system design might require the stack and heap to be placed in separate regions of memory.

For example, you might have a small block of fast RAM that you want to reserve for stack use only. To inform RVCT that you want to use a *two-region model*, you must import the symbol `__use_two_region_memory` using the assembler `IMPORT` directive. The heap is then checked against a dedicated heap limit, that is set up by `__user_initial_stackheap()`.

Figure 2-11 on page 2-23 and Example 2-6 on page 2-23 show an example of implementing a two-region model.

In this example, the stack grows downwards from `0x40000` towards a limit of `0x20000`. To make use of this stack limit, all modules using this implementation must be compiled for software stack checking. The heap grows upwards from `0x28000000` to `0x28080000`.




Figure 2-11 Two-region model

Example 2-6 Two-region model routine

```

IMPORT __use_two_region_memory
EXPORT __user_initial_stackheap

__user_initial_stackheap
    LDR r0, =0x28000000 ;HB
    LDR r1, =0x40000 ;SB
    LDR r2, =0x28080000 ;HL
    LDR r3, =0x20000 ;SL
    MOV pc, lr

```

2.4.8 Example code for Build 3

Build 3 of the example implements scatter-loading and contains a re-implemented `__user_initial_stackheap()`. See the example build files in the main examples directory, in `...\emb_sw_dev\build3`.

The following modifications have been made to Build 2 of the example project:

Scatter-loading

A simple scatter-loading description file is passed to the linker.

Retargeted `__user_initial_stackheap()`

You have the option of selecting either a one-region or a two-region implementation. The default build uses one region. You can select the two-region implementation by defining `TWO_REGION_MODEL` when assembling.

Avoiding C library Semihosting

The symbol `__use_no_semihosting_swi` is imported into Build 3, because there are no longer any C library semihosting functions present in the image.

Note

To avoid using semihosting for `clock()`, this is retargeted to read the *Real Time Clock* (RTC) on the Integrator AP. This has a resolution of one second, so the results from Dhrystone are not precise. This mechanism is improved in Build 4 (see *Example code for Build 4* on page 2-33).

To run this build on an Integrator AP, you must perform ROM/RAM remapping. To do this, set switches 1 and 4 to ON to run the Boot Monitor.

See *Running the Dhrystone builds on an Integrator* on page 2-3.

Note

You must disable all Vector Catch and semihosting if you are using an ARM7 core-based target. Otherwise the debugger interprets the execution of instructions between address `0x0` and `0x1C` as exceptions, and reports this in a dialog box. See your debugger documentation for details of how to disable Vector Catch and semihosting.

2.5 Reset and initialization

This chapter has so far assumed that execution begins at `__main`, the entry point to the C library initialization routine. In fact, any embedded application on your target hardware performs some system-level initialization at startup. This section describes this in more detail, and includes:

- *Initialization sequence* on page 2-26
- *The vector table* on page 2-27
- *ROM/RAM remapping* on page 2-28
- *Local memory setup considerations* on page 2-30
- *Scatter-loading and memory setup* on page 2-30
- *Stack pointer initialization* on page 2-31
- *Hardware initialization* on page 2-32
- *Execution mode considerations* on page 2-33
- *Example code for Build 4* on page 2-33.

2.5.1 Initialization sequence

Figure 2-12 shows a possible initialization sequence for an embedded system based on an ARM architecture.

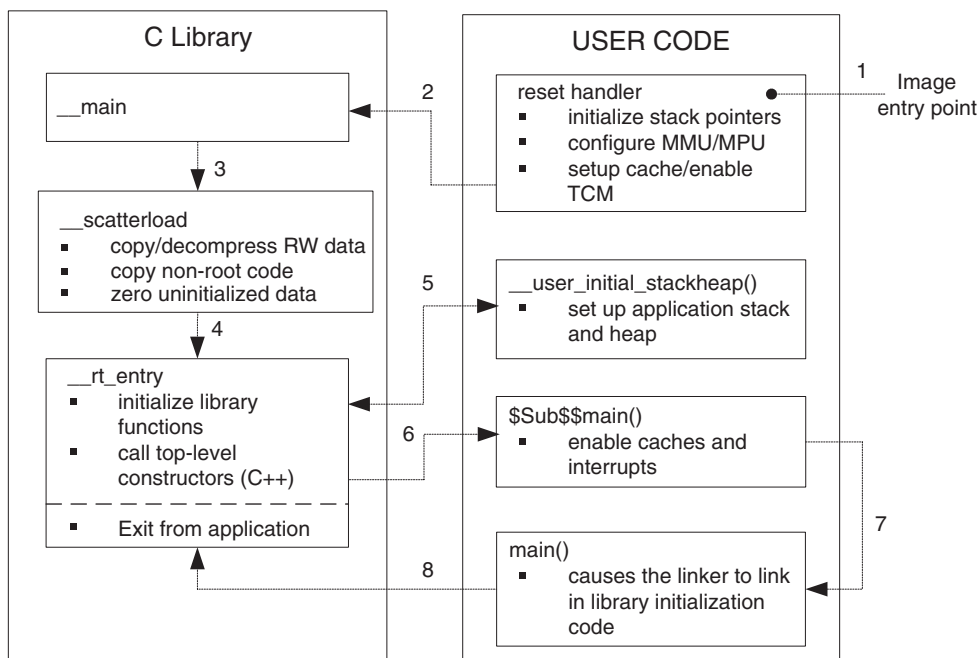


Figure 2-12 Initialization sequence

The reset handler executes immediately on system startup. The block of code labeled `$Sub$$main()` executes immediately before entering the main application.

The reset handler is a short module coded in assembler that is executed on system reset. As a minimum, your reset handler initializes stack pointers for the modes that your application is running in. For cores with local memory systems, such as caches, *Tightly Coupled Memories* (TCMs), *Memory Management Units* (MMUs), and *Memory Protection Units* (MPUs), some configuration must be done at this stage in the initialization process. After executing, the reset handler typically branches to `__main` to begin the C library initialization sequence.

There are some components of system initialization, for example, the enabling of interrupts, that are generally performed after the C library initialization code has finished executing. The block of code labeled `$Sub$$main()` performs these tasks immediately before the main application begins executing.

See *The vector table* for a more detailed description of the various components of the initialization sequence.

2.5.2 The vector table

All ARM systems have a *vector table*. The vector table does not form part of the initialization sequence, but it must be present for any exception to be serviced.

The code in Example 2-7 imports the various exception handlers that might be coded in other modules. The vector table is a list of branch instructions to the exception handlers.

The FIQ handler is placed at address 0x1C directly. This avoids having to execute a branch to the FIQ handler, so optimizing FIQ response time.

Example 2-7 The vector table code

```

PRESERVE8

AREA Vectors, CODE, READONLY
IMPORT Reset_Handler
; import other exception handlers
; ...
ENTRY
B Reset_Handler
B Undefined_Handler
B SWI_Handler
B Prefetch_Handler
B Abort_Handler
NOP ; Reserved vector
B IRQ_Handler
B FIQ_Handler
END

```

———— Note ————

The vector table is marked with the label ENTRY. This label informs the linker that this code is a possible entry point, and so cannot be removed from the image at link time. You must select one of the possible image entry points as the true entry point to your application using the `--entry` linker option. See *RealView Compilation Tools v2.2 Linker and Utilities Guide* for more information.

2.5.3 ROM/RAM remapping

You must consider what sort of memory your system has at address `0x0000`, the address of the first instruction executed.

———— Note ————

This section assumes that the ARM core begins fetching instructions at `0x0000`. This is the norm for systems based on ARM cores. However, some ARM cores can be configured to begin fetching instructions from `0xFFFF0000`.

There has to be a valid instruction at `0x0000` at startup, so you must have non-volatile memory located at `0x0000` at the moment of reset.

One way to achieve this is to have ROM located at `0x0000`. However, there are some drawbacks to this configuration. Access speeds to ROM are generally slower than to RAM, and your system might suffer if there is too great a performance penalty when branching to exception handlers. Also, locating the vector table in ROM does not enable you to modify it at runtime.

Another solution is shown in Figure 2-13. ROM is located at address `0x10000`, but this memory is aliased to zero by the memory controller at reset. Following reset, code in the reset handler branches to the real address of ROM. The memory controller then removes the aliased ROM, so that RAM is shown at address `0x0000`. In `__main`, the vector table is copied into RAM at `0x0000`, so that exceptions can be serviced.

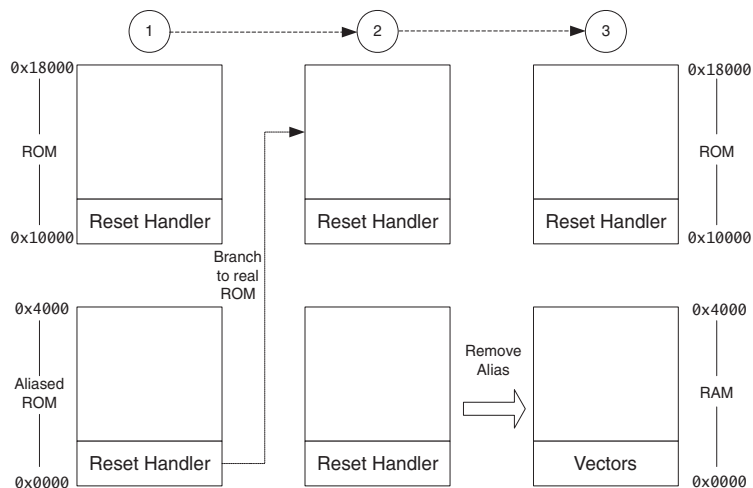


Figure 2-13 ROM/RAM remapping

Example 2-8 shows how you might implement ROM/RAM remapping in an ARM assembler module. The constants shown here are specific to the Integrator platform, but the same method is applicable to any platform that implements ROM/RAM remapping in a similar way.

Example 2-8 ROM/RAM remapping

```

; --- Integrator CM control reg
CM_ctl_reg    EQU    0x1000000C    ; Address of CM Control Register
Remap_bit     EQU    0x04          ; Bit 2 is remap bit of CM_ctl

ENTRY

; Code execution starts here on reset
; On reset, an alias of ROM is at 0x0, so jump to 'real' ROM.
    LDR        pc, =Instruct_2

Instruct_2
; Remap by setting Remap bit of the CM_ctl register
    LDR        r1, =CM_ctl_reg
    LDR        r0, [r1]
    ORR        r0, r0, #Remap_bit
    STR        r0, [r1]

; RAM is now at 0x0.
; The exception vectors must be copied from ROM to RAM (in __main)

; Reset_Handler follows on from here

```

The first instruction is a jump from aliased ROM to real ROM. This can be done because the label `Instruct_2` is located at the real ROM address.

After this step, the alias of ROM is removed by inverting the remap bit of the Integrator Core Module control register.

This code is normally executed immediately after system reset. Remapping must be completed before C library initialization code can be executed.

————— Note —————

In systems with MMUs, remapping can be implemented through MMU configuration at system startup.

2.5.4 Local memory setup considerations

Many ARM cores have on-chip memory systems, such as MMUs or MPUs. These devices are normally set up and enabled during system startup. Therefore, the initialization sequence of cores with local memory systems requires special consideration.

As described in this chapter, C library initialization code in `__main` is responsible for setting up the execution time memory map of the image. Therefore, the runtime memory view of the processor core must be set up before branching to `__main`. This means that any MMU or MPU must be set up and enabled in the reset handler.

TCMs must also be enabled before branching to `__main` (normally before MMU/MPU setup), because you generally want to scatter-load code and data into TCMs. You must be careful that you do not have to access memory that is masked by the TCMs when they are enabled.

You also risk problems with cache coherency if caches are enabled before branching to `__main`. Code in `__main` copies code regions from their load address to their execution address, essentially treating instructions as data. As a result, some instructions can be cached in the data cache, in which case they are not visible to the instruction path.

To avoid these coherency problems, enable caches after the C library initialization sequence finishes executing.

2.5.5 Scatter-loading and memory setup

In a system where the reset-time memory view of the core is altered, either through ROM/RAM remapping or MMU configuration, the scatter-loading description file must describe the image memory map after remapping has taken place.

The description file in Example 2-9 relates to the example in *ROM/RAM remapping* on page 2-28 after remapping.

Example 2-9

```
ROM_LOAD 0x10000 0x8000
{
    ROM_EXEC 0x10000 0x8000
    {
        reset_handler.o (+R0, +FIRST)    ; executed on hard reset
        ...
    }

    RAM 0x0000 0x4000
    {
```

```

        vectors.o (+R0, +FIRST)        ; vector table copied
                                        ; from ROM to RAM at zero
        ...
    }
}

```

The load region ROM_LOAD is placed at 0x10000, because this indicates the load address of code and data after remapping has occurred.

2.5.6 Stack pointer initialization

As a minimum, your reset handler must assign initial values to the stack pointers of any execution modes that are used by your application.

In Example 2-10, the stacks are located at `stack_base`. This symbol can be a hard-coded address, or it can be defined in a separate assembler source file and located by a scatter-loading description file. Details of how this is done are given in *Placing the stack and heap in the scatter-loading description file* on page 2-36.

Example 2-10 Initializing stack pointers

```

; --- Amount of memory (in bytes) allocated for stacks
Len_FIQ_Stack    EQU    256
Len_IRQ_Stack    EQU    256
...
Offset_FIQ_Stack    EQU    0
Offset_IRQ_Stack    EQU    Offset_FIQ_Stack + Len_FIQ_Stack
...
Reset_Handler

; stack_base could be defined above, or located in a description file
    LDR    r0, stack_base ;

; Enter each mode in turn and set up the stack pointer
    MSR    CPSR_c, #Mode_FIQ:OR:I_Bit:OR:F_Bit
    SUB    sp, r0, #Offset_FIQ_Stack

    MSR    CPSR_c, #Mode_IRQ:OR:I_Bit:OR:F_Bit
    SUB    sp, r0, #Offset_IRQ_Stack
    ...
; Set up stack limit if needed
    LDR    r10, stack_limit

```

Example 2-10 on page 2-31 allocates 256 bytes of stack for FIQ and IRQ mode, but you can do the same for any other execution mode. To set up the stack pointers, enter each mode (interrupts disabled) and assign the appropriate value to the stack pointer. To make use of software stack checking, you must also set up a stack limit here.

Stack pointer and stack limit values set up in the reset handler are automatically passed as parameters to `__user_initial_stackheap()` by C library initialization code. Therefore, these values must not be modified by `__user_initial_stackheap()`.

Example 2-11 shows an implementation of `__user_initial_stackheap()` that you can use with the stack pointer setup shown in Example 2-10 on page 2-31.

Example 2-11

```

IMPORT heap_base
EXPORT __user_initial_stackheap()

__user_initial_stackheap()

; heap base could be hard-coded, or placed by description file
LDR    r0,=heap_base
; r1 contains SB value
MOV    pc,r1

```

2.5.7 Hardware initialization

In general, it is beneficial to separate all system initialization code from the main application. However, some components of system initialization, for example, enabling of caches and interrupts, must occur after executing C library initialization code.

You can make use of the `$Sub` and `$Super` function wrapper symbols to (effectively) insert a routine that is executed immediately before entering the main application. This mechanism enables you to extend functions without altering the source code.

Example 2-12 on page 2-33 shows how `$Sub` and `$Super` can be used in this way. The linker replaces the function call to `main()` with a call to `$Sub$$main()`. From there you can call a routine that enables caches and another to enable interrupts.

The code branches to the real `main()` by calling `$Super$$main()`.

———— Note ————

For more information on `$Sub` and `$Super`, see *RealView Compilation Tools v2.2 Linker and Utilities Guide*.

Example 2-12 Use of \$Sub and \$Super

```
extern void $Super$$main(void);

void $Sub$$main(void)
{
    cache_enable();    // enables caches
    int_enable();      // enables interrupts
    $Super$$main();    // calls original main()
}
```

2.5.8 Execution mode considerations

You must consider in what mode the main application is to run. Your choice affects how you implement system initialization.

Much of the functionality that you are likely to implement at startup, both in the reset handler and \$Sub\$\$main, can only be done while executing in privileged modes, for example, on-chip memory manipulation, and enabling interrupts.

If you want to run your application in a privileged mode (for example, Supervisor), this is not an issue. Ensure that you change to the appropriate mode before exiting your reset handler.

If you want to run your application in User mode, however, you can only change to User mode *after* completing the necessary tasks in a privileged mode. The most likely place to do this is in \$Sub\$\$main().

Note

__user_initial_stackheap() must set up the application mode stack. Because of this, you must exit your reset handler in system mode, which uses the User mode registers. __user_initial_stackheap() then executes in system mode, and so the application stack and heap are still set up when User mode is entered.

2.5.9 Example code for Build 4

Build 4 of the example can be run standalone on the Integrator platform. See the example build files in the main examples directory, in `...\emb_sw_dev\build4`.

The following modifications have been made to Build 3 of the example project:

Vector table

A vector table has been added to the project, and placed by the scatter-loading description file.

Reset handler

The reset handler is added in `init.s`. Two separate modules, responsible for TCM and MMU setup respectively, are included in the ARM926EJ-S™ build. These are excluded from the ARM7TDMI® build, which runs on Integrator systems with any core. ROM/RAM remapping occurs immediately after reset.

`Submain()`

For the ARM926EJ-S build, Caches are enabled in `Submain()` before entering the main application.

Embedded description file

An embedded description file is used, that reflects the memory view after remapping.

The build files for both of these builds produce a binary file suitable for downloading into the Integrator AP application Flash at address `0x24000000`.

A precise timer is implemented using a timer on the Integrator AP motherboard. This generates an IRQ, and a handler is installed that increments a counter every one-hundredth of a second (0.01 sec).

2.6 Further memory map considerations

The previous sections in this chapter describe the placement of code and data in a scatter-loading description file. However, the location of target hardware peripherals and the stack and heap limits are assumed to be hard-coded in source or header files. It would be beneficial to locate all information pertaining to the memory map of a target in your description file, so removing all references to absolute addresses from your source code.

This section includes:

- *Locating target peripherals in the scatter-loading description file*
- *Placing the stack and heap in the scatter-loading description file* on page 2-36
- *Example code for Build 5* on page 2-41.

2.6.1 Locating target peripherals in the scatter-loading description file

Conventionally, addresses of peripheral registers are hard-coded in project source or header files. You can also declare structures that map on to peripheral registers, and place these structures in the description file.

For example, a target might have a timer peripheral with two memory mapped 32-bit registers. Example 2-13 shows a C structure that maps to these registers.

Example 2-13 Mapping to a peripheral register

```
struct {
    volatile unsigned ctrl; /* timer control */
    volatile unsigned tmr; /* timer value */
} timer_regs;
```

To place this structure at a specific address in the memory map, create a new execution region to hold the structure.

The description file shown in Example 2-14 on page 2-36 locates the `timer_regs` structure at `0x40000000`.

It is important that the contents of these registers are not initialized to zero during application startup, because this is likely to change the state of your system. Marking an execution region with the `UNINIT` attribute prevents ZI data in that region from being zero initialized.

Example 2-14 Placing the mapped structure

```

ROM_LOAD 0x24000000 0x04000000
{
    ; ...
    TIMER 0x40000000 UNINIT
    {
        timer_regs.o (+ZI)
    }
    ; ...
}

```

2.6.2 Placing the stack and heap in the scatter-loading description file

In many cases, it is preferable to specify the location of the stack and heap in the description file. This has two main advantages:

- all information about the memory map is kept in one file
- changes to the stack and heap only require relinking, not recompiling.

This section describes methods for implementing this:

- *Placing symbols explicitly* (the simplest method)
- *Utilizing linker generated symbols* on page 2-38
- *Using the scatter file EMPTY attribute* on page 2-39.

Placing symbols explicitly

Stack pointer initialization on page 2-31 refers to the symbols `stack_base` and `heap_base` as reference symbols that can be placed in a description file. To do this, create symbols labeled `stack_base` and `heap_base` in an assembler module called `stackheap.s`. The same can be done for the stack and heap limits in a two-region memory model.

You can locate each of the symbols within their own execution region in the description file, as shown in Example 2-15.

Example 2-15 Placing symbols explicitly in `stackheap.s`

```

        AREA    stacks, DATA, NOINIT
        EXPORT  stack_base

stack_base      SPACE    1

        AREA    heap, DATA, NOINIT
        EXPORT  heap_base

```

```
heap_base      SPACE 1
               END
```

Figure 2-14 and Example 2-16 show how you can place the heap base at 0x20000 and the stack base at 0x40000. The stack and heap base locations can be altered by editing the addresses of the respective execution regions.

The disadvantage of this approach is that one word of SPACE (stack_base) is occupied above the stack region.




Figure 2-14 Placing symbols explicitly

Example 2-16 Placing symbols explicitly in a scatter file

```
LOAD_FLASH 0x24000000 0x04000000
{
    ; ...
    HEAP 0x20000 UNINIT
    {
        stackheap.o (heap)
    }

    STACKS 0x40000 UNINIT
    {
        stackheap.o (stacks)
    }
    ; ...
}
```

Utilizing linker generated symbols

This method requires that the stack and heap sizes are specified in an object file.

First, define areas of an appropriate size for the stack and heap in an assembler source file, for example, `stackheap.s`, as shown in Example 2-17.

Use the `SPACE` directive to reserve a zeroed block of memory. Set the `NOINIT` area attribute to prevent this zeroing.

During development, you might choose to zero-initialize the stack so that the maximum stack usage can be seen. Labels are not required in this source file.

Example 2-17 Placing sections for stack and heap

```

        AREA stack, DATA, NOINIT
SPACE   0x3000 ; Reserve stack space

        AREA heap, DATA, NOINIT
SPACE   0x3000 ; Reserve heap space

        END

```

You can then place these sections in their own execution region in the scatter-loading description file, as shown in Example 2-18.

Example 2-18 Placing sections for stack and heap

```

LOAD_FLASH 0x24000000 0x04000000
{
    :
    STACKS 0x1000 UNINIT      ; length = 0x3000
    {
        stackheap.o (stack) ; stack = 0x4000 to 0x1000
    }

    HEAP 0x15000 UNINIT      ; length = 0x3000
    {
        stackheap.o (heap) ; heap = 0x15000 to 0x18000
    }
}

```

The linker generates symbols that point to the base and limit of each execution region, that can be imported into the retargeting code to be used by `__user_initial_stackheap()`:

```
Image$$STACKS$$ZI$$Limit = 0x4000
Image$$STACKS$$ZI$$Base  = 0x1000
Image$$HEAP$$ZI$$Base    = 0x15000
Image$$HEAP$$ZI$$Limit   = 0x18000
```

You can make this code more readable by using the DCD directive to give these values more meaningful names, as shown in Example 2-19.

Example 2-19 Using the DCD directive

IMPORT		Image\$\$STACKS\$\$ZI\$\$Base	
IMPORT		Image\$\$STACKS\$\$ZI\$\$Limit	
IMPORT		Image\$\$HEAP\$\$ZI\$\$Base	
IMPORT		Image\$\$HEAP\$\$ZI\$\$Limit	
stack_base	DCD	Image\$\$STACKS\$\$ZI\$\$Limit	; = 0x4000
stack_limit	DCD	Image\$\$STACKS\$\$ZI\$\$Base	; = 0x1000
heap_base	DCD	Image\$\$HEAP\$\$ZI\$\$Base	; = 0x15000
heap_limit	DCD	Image\$\$HEAP\$\$ZI\$\$Limit	; = 0x18000

You can use these examples to place the heap base at 0x15000 and the stack base at 0x1000. You can then change the stack and heap base locations easily by editing the addresses of the respective execution regions.

Using the scatter file EMPTY attribute

This method uses the scatter file EMPTY attribute of the linker. This enables regions to be defined that contain no object code or data. This is a convenient method of defining a stack or heap. The length of the region is specified after the EMPTY attribute. In the case of a heap, that grows upwards in memory, the region length is positive. In the case of a stack, the region length is marked as negative, to indicate that it grows downwards in memory. Example 2-20 on page 2-40 shows how to use the EMPTY attribute.

The benefit of this approach is that the size and position of the stack and heap is defined in one place, that is, in the scatter-loading description file. You do not have to create a `stackheap.s` file.

Example 2-20 Placing stack and heap regions using EMPTY

```
ROM_LOAD 0x24000000 0x04000000
{
    ...
    HEAP 0x30000 EMPTY 0x3000
    {
    }

    STACKS 0x40000 EMPTY -0x3000
    {
    }
    ...
}
```

At link time, the linker generates symbols to represent these EMPTY regions:

```
Image$$HEAP$$ZI$$Base      = 0x30000
Image$$HEAP$$ZI$$Limit     = 0x33000
Image$$STACKS$$ZI$$Base    = 0x3D000
Image$$STACKS$$ZI$$Limit   = 0x40000
```

Your application code can then process these symbols as shown in Example 2-21.

Example 2-21 Linker generated symbols representing EMPTY regions

	IMPORT	Image\$\$HEAP\$\$ZI\$\$Base
	IMPORT	Image\$\$HEAP\$\$ZI\$\$Limit
heap_base	DCD	Image\$\$HEAP\$\$ZI\$\$Base
heap_limit	DCD	Image\$\$HEAP\$\$ZI\$\$Limit
	IMPORT	Image\$\$STACKS\$\$ZI\$\$Base
	IMPORT	Image\$\$STACKS\$\$ZI\$\$Limit
stack_base	DCD	Image\$\$STACKS\$\$ZI\$\$Limit
stack_limit	DCD	Image\$\$STACKS\$\$ZI\$\$Base

2.6.3 Example code for Build 5

Build 5 of the example is equivalent to Build 4, but with all target memory map information located in the scatter-loading description file as described in *Placing symbols explicitly* on page 2-36:

Scatter-loading description file symbols

Symbols to locate the stack, heap, and peripherals are declared in assembler modules.

Updated Scatter-loading description file

The embedded description file from Build 4 is updated to locate the stack, heap, data TCM, and peripherals.

See the example build files in the main examples directory, in `...\emb_sw_dev\build5`.

The stack and heap are located using linker symbols, see *Utilizing linker generated symbols* on page 2-38.

Chapter 3

Writing Position Independent Code and Data

This chapter describes how to write position independent code and data that makes use of the *Procedure Call Standard for the ARM Architecture* (AAPCS). It contains the following sections:

- *Position independence* on page 3-2
- *Read-only position independence* on page 3-3
- *Read-write position independence* on page 3-6.

3.1 Position independence

Both the ARM® and Thumb® instruction sets support position-independent, or *relocatable*, code through the use of PC-relative instructions (for example BL).

Note

This is not the same as Relocatable ELF (an image type created by the linker).

You can write assembler code that is relocatable but it must not contain any address constants. Any literal addresses used to refer to code must be PC-relative offsets. The PC is added, using an ADD instruction, before the address is accessed.

Both code and data can be position-independent:

- To enable code to execute at different addresses, it must be position-independent or relocatable. However, it can only access a single set of static data at a fixed address.
- Position-independent data requires all data accesses to occur relative to the static base register sb. This is used to implement a shared library mechanism.

RVCT supports position-independent code and data for C and assembler (but not C++), and enables you to write code that is relocatable or reentrant. The rest of this chapter contains information about how to do this.

For more information, see the following in *RealView Compilation Tools v2.2 Compiler and Libraries Guide*:

- the section about position independence qualifiers in the chapter describing how to use the compiler
- the section about writing reentrant and thread-safe code in the chapter describing the C and C++ libraries.

3.1.1 Using the AAPCS

The *Procedure Call Standard for the ARM Architecture* (AAPCS) forms part of the *Application Binary Interface (ABI) for the ARM Architecture (base standard)* [BSABI] specification. By writing code that adheres to the AAPCS, you can ensure that separately compiled and assembled modules can work together.

For more information, see Appendix A *Using the Procedure Call Standard* and the AAPCS specification in `install_directory\Documentation\...`

3.2 Read-only position independence

A program is *Read-Only Position-Independent* (ROPI) if all its read-only segments are position independent.

An ROPI segment is often *Position-Independent Code* (PIC), but could be read-only data, or a combination of PIC and read-only data.

Note

ROPI does not form part of the AAPCS, because it is not supported for C++. However, you can compile your C code or assembler code for ROPI by using the compiler or assembler option `--apcs /ropi`.

Select the ROPI option to avoid committing yourself to loading your code in a particular location in memory. This is particularly useful for routines that are:

- loaded in response to runtime events
- loaded into memory with different combinations of other routines in different circumstances
- mapped at different addresses during their execution.

This section includes:

- *Register usage with ROPI*
- *Writing C and assembler code for ROPI* on page 3-4
- *Linking your code* on page 3-4
- *FPIC addressing* on page 3-4
- *Code example* on page 3-5.

3.2.1 Register usage with ROPI

As defined by the AAPCS, register use is the same with or without ROPI.

For more information, see:

- the *Procedure Call Standard for the ARM Architecture* specification, `aapcs.pdf`, in `install_directory\Documentation\Specifications\...`
- Appendix A *Using the Procedure Call Standard*.

3.2.2 Writing C and assembler code for ROPI

When you are writing C and assembler code for ROPI:

- Every reference from code in an ROPI segment to a symbol in the same ROPI segment must be PC-relative. AAPCS does not define any other base register for a read-only segment. An address of an item in an ROPI segment cannot be assigned to an item in a different ROPI segment.
- Every reference from code in an ROPI segment to a symbol in a different ROPI segment must be PC-relative. The two segments must be fixed relative to each other.
- Every other reference from an ROPI segment must be to either:
 - an absolute address
 - an sb-relative reference to writable data (see *Read-write position independence* on page 3-6).
- A read-write word that addresses a symbol in an ROPI segment must be adjusted whenever the ROPI segment is moved.

3.2.3 Linking your code

Use the linker command-line option `--ropi` to make the load and execution region containing the read-only output section position-independent. If this option is not used the region is marked as absolute. Usually each read-only input section must be read-only position-independent. See *RealView Compilation Tools v2.2 Linker and Utilities Guide* for details.

3.2.4 FPIC addressing

Use the `/fpic` qualifier to generate read-only position-independent code where relative address references are independent of the location where your program is loaded. Relative addressing is only implemented when your code makes use of System V shared libraries. If your code uses shared objects, you do not have to compile with `/fpic`.

For more information see Appendix A *Using the Procedure Call Standard*.

For information on System V shared library support in RVCT, see *RealView Compilation Tools v2.2 Linker and Utilities Guide*.

3.2.5 Code example

For details of writing position-independent code, see the PIC-PID example provided with RealView Developer Suite in the main examples directory, that is in *install_directory\RVDS\Examples\picpid*.

This example consists of a kernel at a fixed address in ROM, together with a collection of application modules that extend kernel functionality. Application modules are loaded into memory following the kernel. However, the address where a module might be loaded is unknown when the module is linked. Therefore, modules must be position-independent (ROPI, PIC).

The example includes source code, a make file, batch files, and a detailed description of how to compile and link the different modules (see *readme.txt*).

3.3 Read-write position independence

A program is *Read-Write Position-Independent* (RWPI) if all its read-write segments are position independent.

An RWPI segment is usually *Position-Independent Data* (PID).

RWPI is an AAPCS variant. Use the compiler or assembler option `--apcs /rwpi` to avoid committing yourself to a particular location of data in memory. This is particularly useful for data that must be multiply instantiated for reentrant routines.

For more information, see:

- the *Procedure Call Standard for the ARM Architecture* specification, `aapcs.pdf`, in `install_directory\Documentation\Specifications\...`
- Appendix A *Using the Procedure Call Standard*.

This section includes:

- *Reentrant routines*
- *Register usage with RWPI*
- *Position-independent data addressing* on page 3-7
- *Writing assembly language for RWPI* on page 3-7
- *Linking your code* on page 3-7
- *Code example* on page 3-7.

3.3.1 Reentrant routines

A reentrant routine can be *threaded* by several processes at the same time. Each process has its own copy of the read-write segments of the routine. Each copy is addressed by a different value of the static base register (sb).

3.3.2 Register usage with RWPI

Register r9 is the static base register, sb. It must point to the base address of the appropriate static data segments whenever you call any externally visible routine.

You can use r9 for other purposes in a routine that does not use sb. If you do this you must save the contents of sb on entry to your routine and restore these before exit. You must also restore the contents before any call to an external routine.

In all other respects the usage of registers is the same with or without RWPI.

3.3.3 Position-independent data addressing

An RWPI segment can be repositioned until it is first used. The address of a symbol in an RWPI segment is calculated as follows:

1. The linker calculates a read-only offset from a fixed location in the segment. By convention, the fixed location is the first byte of the lowest addressed RWPI segment of the program.
2. At runtime, this is used as an offset added to the contents of the static base register, sb.

3.3.4 Writing assembly language for RWPI

Construct references from a read-only segment to the RWPI segment by adding a fixed (read-only) offset to the value of sb (see DCD0 in the *Directives Reference* chapter in *RealView Compilation Tools v2.2 Assembler Guide*).

3.3.5 Linking your code

Use the linker command-line option `--rwp` to make the load and execution region containing the RW and ZI output sections position-independent. If this option is not used the region is marked as absolute. This option requires a value for `--rw-base`. If `--rw-base` is not specified, `--rw-base 0` is assumed. Usually each writable input section must be RWPI. See *RealView Compilation Tools v2.2 Linker and Utilities Guide* for details of these options.

3.3.6 Code example

For details of writing position-independent code, see the PIC-PID example provided with RealView Developer Suite in the main examples directory, that is in `install_directory\RVDS\Examples\picpid`.

This example consists of a kernel at a fixed address in ROM, together with a collection of application modules that extend kernel functionality. A module implements a set of named services that can be multiply instantiated, and that can call one another through the kernel. When a service is called, the kernel creates an instance of its static data and then passes control to the service. However, the service might then call back to the kernel. Therefore, modules must have position-independent data (RWPI, PID).

The example includes source code, a make file, batch files, and a detailed description of how to compile and link the different modules (see `readme.txt`).

Chapter 4

Interworking ARM and Thumb

This chapter explains how to change between ARM® state and Thumb® state when writing code for processors that implement the Thumb instruction set. It contains the following sections:

- *About interworking* on page 4-2
- *Assembly language interworking* on page 4-7
- *C and C++ interworking and veneers* on page 4-13
- *Assembly language interworking using veneers* on page 4-18.

4.1 About interworking

Interworking enables you to mix ARM and Thumb code so that:

- ARM routines return to a Thumb state caller
- Thumb routines return to an ARM state caller.

This means that, if you compile or assemble code for interworking, your code can call a routine in a different module without considering which instruction set it uses.

The ARM linker detects when an ARM function is being called from Thumb state, or a Thumb function is being called from ARM state. The ARM linker changes call and return instructions, or inserts small code segments called *veneers*, to change processor state as necessary.

The ARMv5T architecture provides methods to change processor state without using any extra instructions. There is normally no cost associated with interworking on ARMv5T processors.

———— Note ————

Compiling for ARMv5TE automatically assumes interworking, and always produces code that interworks. However, assembly code built for ARMv5TE does not imply interworking, so you must build assembly code with the `--apcs /interwork` assembler option.

4.1.1 Using the AAPCS

You can mix ARM and Thumb code as you require, provided that the code conforms to the requirements of the AAPCS. For more information, see:

- the *Procedure Call Standard for the ARM Architecture* specification, `aapcs.pdf`, in `install_directory\Documentation\Specifications\...`
- *Appendix A Using the Procedure Call Standard*.

If you are writing ARM assembly language modules you must ensure that your code conforms to the AAPCS. If you are linking several source files together, all your files must use compatible AAPCS options. If incompatible options are detected, the linker produces an error message.

This section describes:

- *When to use interworking* on page 4-3
- *Using the /interwork option* on page 4-4
- *Detecting interworking calls* on page 4-4
- *Linker generated veneers* on page 4-5.

4.1.2 When to use interworking

When you write code for an ARM processor that supports Thumb instructions, you probably write most of your application to run in Thumb state. This gives the best code density. With 8-bit or 16-bit wide memory, it also gives the best performance. However, you might want parts of your application to run in ARM state for reasons such as:

Speed Some parts of an application might be speed critical. These sections might be more efficient running in ARM state than in Thumb state. In some circumstances, a single ARM instruction can do more than the equivalent Thumb instruction.

Some systems include a small amount of fast 32-bit memory. ARM code can be run from this without the overhead of fetching each instruction from 8-bit or 16-bit memory.

Functionality

Thumb instructions are less flexible than their equivalent ARM instructions. Some operations are not possible in Thumb state. A state change to ARM is required to carry out the following operations:

- accesses to CPSR to enable or disable interrupts, and to change mode
- accesses to coprocessors
- DSP math instructions that are not supported by C.

Exception handling

The processor automatically enters ARM state when a processor exception occurs. This means that the first part of an exception handler must be coded with ARM instructions, even if it re-enters Thumb state to carry out the main processing of the exception. At the end of such processing, the processor must be returned to ARM state to return from the handler to the main application.

Standalone Thumb programs

An ARM processor that supports Thumb instructions always starts in ARM state. To run simple Thumb assembly language programs under the debugger, add an ARM header that carries out a state change to Thumb state and then calls the main Thumb routine. See *Example ARM header* on page 4-9 for an example.

4.1.3 Using the /interwork option

The option `--apcs /interwork` is available for the ARM compiler and assembler. If you set this option:

- The compiler or assembler records an interworking attribute in the object file.
- The linker provides interworking veneers for subroutine entry.
- In assembly language, you must write function exit code that returns to the instruction set state of the caller, for example `BX lr`.
- In C or C++, the compiler creates function exit code that returns to the instruction set state of the caller.
- In C or C++, the compiler uses `BX` instructions for indirect or virtual calls.

Use the `--apcs /interwork` option if your object file contains:

- Thumb subroutines that might have to return to ARM code
- ARM subroutines that might have to return to Thumb code
- Thumb subroutines that might make indirect or virtual calls to ARM code
- ARM subroutines that might make indirect or virtual calls to Thumb code.

Note

If a module contains functions marked with `#pragma arm` or `#pragma thumb`, the module must be compiled with `--apcs /interwork`. This ensures that the functions can be called successfully from the other (ARM or Thumb) state.

Otherwise, you do not have to use the `/interwork` option. For example, your object file might contain any of the following without requiring `/interwork`:

- Thumb code that can be interrupted by an exception. The exception forces the processor into ARM state so no veneer is required.
- Exception handling code that can handle exceptions from Thumb code. No veneer is required for the return.

4.1.4 Detecting interworking calls

The linker generates an error if it detects a direct ARM/Thumb interworking call where the called routine is not built for interworking. You must rebuild the called routine for interworking.

For example, Example 4-1 shows the error that is produced if the ARM routine in Example 4-3 on page 4-14 is compiled and linked without the `--apcs /interwork` option.

Example 4-1

```
Error: L6239E: Cannot call ARM symbol 'arm_function' in non-interworking object
armsub.o from THUMB code in thumbmain.o(.text)
```

These types of error indicate that an ARM-to-Thumb or Thumb-to-ARM interworking call has been detected from the object module object to the routine symbol, but the called routine has not been compiled for interworking. You must recompile the module that contains the symbol and specify `--apcs /interwork`.

4.1.5 Linker generated veneers

Veneers are small code segments that are automatically inserted by the linker if a branch involves:

- a change of state
- a destination beyond the range of the branching instruction.

The veneer becomes the target of the original branch, that then branches to the target address.

The linker can reuse a veneer generated for a previous call for subsequent calls to the same function, provided they can be reached from both sections.

For more details on interworking with veneers, see:

- *C and C++ interworking and veneers* on page 4-13
- *Assembly language interworking using veneers* on page 4-18.

Types of veneer

Veneers can be:

long	Can have an optional state change.
short	Performs only a state change.
inline	Performs only a state change, but is added to the start of the function that is being veneered.

Veneer\$\$Code sections

The linker creates one input section called Veneer\$\$Code for each veneer. You can place veneer code in a scatter-loading description file using `*(Veneer$$Code)`. However, the linker only places veneer code there if it safe to do so.

It might not be possible for a veneer input section to be assigned to the region because of problems with address range or limitations on the size of execution regions. If the veneer cannot be added to the specified region, it is added to the execution region containing the relocated input section that generated the veneer.

See *RealView Compilation Tools v2.2 Linker and Utilities Guide* for more details.

Minimizing the use of veneers

You can minimize the use of veneers by:

- structuring the memory map to keep called functions within range of the caller
- encouraging sharing of veneers by keeping calling functions within range
- minimizing state changes.

4.2 Assembly language interworking

In an assembly language source file, you can have several areas (these correspond to ELF sections). Each area can contain ARM instructions, Thumb instructions, or both.

You can use the linker to fix up calls to, and returns from, routines that use a different instruction set from the caller. To do this, use BL to call the routine (see *Assembly language interworking using veneers* on page 4-18).

If you prefer, you can write your code to make the instruction set changes explicitly. In some circumstances you can write smaller or faster code by doing this.

The following instructions perform the processor state changes:

- BX, see *The branch and exchange instruction*
- BLX, LDR, LDM, and POP (ARMv5 and above only), see *ARM architecture v5T* on page 4-11.

The following directives instruct the assembler to assemble instructions from the appropriate instruction set (see *Changing the assembler mode* on page 4-8):

- CODE32
- CODE16

This section includes:

- *The branch and exchange instruction*
- *Changing the assembler mode* on page 4-8
- *Example ARM header* on page 4-9
- *ARM architecture v5T* on page 4-11
- *Labels in Thumb code* on page 4-12.

4.2.1 The branch and exchange instruction

The BX instruction is available only on cores that support Thumb. The instruction branches to the address contained in a specified register, and has a 4GB address range. The value of bit 0 of the branch address determines whether execution continues in ARM state or Thumb state. See *ARM architecture v5T* on page 4-11 for additional instructions available with ARMv5.

Bit 0 of an address can be used in this way because:

- all ARM instructions are word-aligned, so bits 0 and 1 of the address of any ARM instruction are unused
- all Thumb instructions are halfword-aligned, so bit 0 of the address of any Thumb instruction is unused.

Syntax

The syntax of BX is one of:

Thumb BX *Rn*

ARM BX{*cond*} *Rn*

where:

Rn Is a register in the range *r0* to *r15* that contains the address to branch to. The value of bit 0 in this register determines the processor state:

- if bit 0 is set, the instructions at the branch address are executed in Thumb state
- if bit 0 is clear, the instructions at the branch address are executed in ARM state.

cond Is an optional condition code. Only the ARM version of BX can be executed conditionally.

4.2.2 Changing the assembler mode

The ARM assembler can assemble both Thumb code and ARM code. By default, it assembles ARM code unless it is invoked with the `--thumb` (or `--16`) option.

Because all ARM processors that support Thumb start in ARM state, you must use the BX instruction to branch and exchange to Thumb state, and then use the following assembler directives to instruct the assembler to switch mode:

CODE16 Instructs the assembler to assemble Thumb instructions. This also causes an alignment to a two-byte boundary, even if no instructions follow it.

CODE32 Instructs the assembler to return to assembling ARM instructions. This also causes an alignment to a four-byte boundary, even if no instructions follow it.

See *RealView Compilation Tools v2.2 Assembler Guide* for more information on these directives.

4.2.3 Example ARM header

Example 4-2 contains four sections of code. Each of the code sections is described after the example.

Example 4-2

```

PRESERVE8

AREA    AddReg, CODE, READONLY ; Name this block of code.
ENTRY   ; Mark first instruction to call.

; SECTION 1
main
    ADR r0, ThumbProg + 1      ; Generate branch target address
                                ; and set bit 0, hence arrive
                                ; at target in Thumb state.
    BX  r0                     ; Branch exchange to ThumbProg.

; SECTION 2
CODE16                          ; Subsequent instructions are Thumb code.
ThumbProg
    MOV r2, #2                  ; Load r2 with value 2.
    MOV r3, #3                  ; Load r3 with value 3.
    ADD r2, r2, r3              ; r2 = r2 + r3
    ADR r0, ARMPProg
    BX  r0                      ; Branch exchange to ARMPProg.

; SECTION 3
CODE32                          ; Subsequent instructions are ARM code.
ARMPProg
    MOV r4, #4
    MOV r5, #5
    ADD r4, r4, r5

; SECTION 4
stop MOV r0, #0x18              ; angel_SWIreason_ReportException
    LDR r1, =0x20026            ; ADP_Stopped_ApplicationExit
    SWI 0x123456                ; ARM semihosting SWI
    END                         ; Mark end of this file.

```

SECTION 1 implements a short header section of ARM code that changes the processor to Thumb state. The header code uses:

- An ADR pseudo-instruction to load the branch address and set the least significant bit. The ADR pseudo-instruction generates the address by loading r0 with the value pc+offset+1. That is, the address of ThumbProg plus one.

Note

An ADR instruction is used for symbols within the same section. For larger ranges, use the LDR instruction. See *RealView Compilation Tools v2.2 Assembler Guide* for more information on the ADR and LDR pseudo-instructions.

- A BX instruction to branch to the Thumb code and change processor state.

SECTION 2 of the code, labeled ThumbProg, is prefixed by a CODE16 directive. This instructs the assembler to treat the following code as Thumb code. The Thumb code adds the contents of two registers together.

The code again uses an ADR instruction to get the address of the label ARMProg, but this time the least significant bit is left clear. The BX instruction changes the state back to ARM state.

SECTION 3 of the code, labeled ARMProg, adds together the contents of two registers.

SECTION 4 of the code, labeled stop, uses the semihosting SWI to report normal application exit. See *RealView Compilation Tools v2.2 Compiler and Libraries Guide* for more information on semihosting.

Note

The Thumb semihosting SWI is a different number from the ARM semihosting SWI (0xAB rather than 0x123456).

Exporting symbols

If you export a symbol that references Thumb instructions, the linker automatically adds one to the address of any label in Thumb code.

If you do not export a symbol, you must manually add one to the symbol that references the Thumb instructions. In Example 4-2 on page 4-9 it is ThumbProg+1. This is because all references are resolved by the assembler, and the linker never detects the symbol.

Building the example

To build and execute the example:

1. Enter the code using any text editor and save the file as `addreg.s`.
2. Type `armasm -g addreg.s` at the command prompt to assemble the source file.
3. Type `armlink addreg.o -o addreg` to link the file.

4. Run the image using a compatible debugger, for example RealView® Debugger or AXD, with an appropriate debug target. If you step through the program one instruction at a time, you see the processor enter the Thumb state. See the user documentation for the debugger you are using to find out how this change is indicated.

4.2.4 ARM architecture v5T

In ARMv5T and above:

- The following additional interworking instructions are available:

BLX address

The processor performs a PC-relative branch to *address* with link and changes state. *address* must be within 32MB of the PC in ARM code, or within 4MB of the PC in Thumb code.

BLX register

The processor performs a branch with link to an address contained in the specified register. The value of bit[0] determines the new processor state.

In either case, bit[0] of *lr* is set to the current value of the Thumb bit in the CPSR. This means that the return instruction can automatically return to the correct processor state.

- If LDR, LDM, or POP load to the PC, they set the Thumb bit in the CPSR to bit[0] of the value loaded to the PC. You can use this to change instruction sets. This is particularly useful for returning from subroutines. The same return instruction can return to either an ARM or Thumb caller.

For more information, see *RealView Compilation Tools v2.2 Assembler Guide* and *ARM Architecture Reference Manual*.

4.2.5 Labels in Thumb code

The linker distinguishes between labels referring to:

- ARM instructions
- Thumb instructions
- data.

When the linker relocates a value of a label referring to a Thumb instruction, it sets the least significant bit of the relocated value. This means that a branch to a label can automatically select the appropriate instruction set. This works if any of the following instructions are used for the branch:

- BX in ARMv4T
- BX, BLX, or LDR in ARMv5T and above.

In releases of *ARM Developer Suite*[™] (ADS) earlier than 1.2, it was necessary to mark data in Thumb code with the DATA directive. This is no longer necessary.

4.3 C and C++ interworking and veneers

You can freely mix C and C++ code compiled for ARM and Thumb, but in ARMv4T veneers are required between the ARM and Thumb code to carry out state changes. The ARM linker generates these interworking veneers when it detects interworking calls. See *Linker generated veneers* on page 4-5 for more details on veneers.

This section includes:

- *Compiling code for interworking*
- *Basic rules for C and C++ interworking* on page 4-16
- *Pointers to functions in Thumb state* on page 4-16
- *Using two versions of the same function* on page 4-17.

4.3.1 Compiling code for interworking

The `--apcs /interwork` compiler option enables the ARM compiler to compile C and C++ modules containing routines that can be called by routines compiled for the other processor state:

```
armcc --c90 --thumb --apcs /interwork
armcc --c90 --arm --apcs /interwork
armcc --cpp --thumb --apcs /interwork
armcc --cpp --arm --apcs /interwork
```

Note

`--arm` is the default option. `--c90` is the default for files with the extension `.c`, and `--cpp` is the default for files with the extension `.cpp`.

Modules that are compiled for interworking on ARMv4T generate slightly larger code. There is no difference for ARMv5.

In a leaf function, that is a function whose body contains no function calls, the only change in the code generated by the compiler is to replace `MOV pc,lr` with `BX lr`. The `MOV` instruction does not cause the necessary state change.

In nonleaf functions built for ARMv4T in Thumb mode, the compiler must replace, for example, the single instruction:

```
POP {r4,r5,pc}
```

with the sequence:

```
POP {r4,r5}
POP {r3}
BX r3
```

This has a small impact on performance. Compile all source modules for interworking, unless you are sure they are never going to be used with interworking.

The `--apcs /interwork` option also sets the interwork attribute for the code area the modules are compiled into. The linker detects this attribute and inserts the appropriate veneer.

Note

ARM code compiled for interworking can only be used on ARMv4T and later, because earlier processors do not implement the BX instruction.

Use the linker option `--info veneers` to find the amount of space taken by the veneers.

C interworking example

Example 4-3 shows a Thumb routine that carries out an interworking call to an ARM subroutine. The ARM subroutine call makes an interworking call to `printf()` in the Thumb library. These two modules are provided in the main examples directory, in `...\interwork` as `thumbmain.c` and `armsub.c`.

Example 4-3

```

/*****
 *      thumbmain.c  *
 *****/
#include <stdio.h>
extern void arm_function(void);
int main(void)
{
    printf("Hello from Thumb\n");
    arm_function();
    printf("And goodbye from Thumb\n");
    return (0);
}

/*****
 *      armsub.c     *
 *****/
#include <stdio.h>
void arm_function(void)
{
    printf("Hello and Goodbye from ARM\n");
}

```

To compile and link these modules:

1. To compile the Thumb code for interworking, type:
`armcc --thumb -c -g -O1 --apcs /interwork -o thumbmain.o thumbmain.c`
2. To compile the ARM code for interworking, type:
`armcc -c -g -O1 --apcs /interwork -o armsub.o armsub.c`
3. To link the object files, type:
`armlink thumbmain.o armsub.o -o thumbtoarm.axf`
 Alternatively, to view the size of the interworking veneers (as shown in Example 4-4) type:
`armlink armsub.o thumbmain.o -o thumbtoarm.axf --info veneers`

Example 4-4

```

Adding TA veneer (4 bytes, Inline) for call to 'arm_function' from thumbmain.o(.text).
Adding AT veneer (8 bytes, Inline) for call to '__0printf' from armsub.o(.text).
Adding AT veneer (8 bytes, Inline) for call to '__rt_lib_init' from kernel.o(.text).
Adding AT veneer (12 bytes, Long) for call to '__rt_lib_shutdown' from kernel.o(.text).
Adding TA veneer (4 bytes, Inline) for call to '__aeabi_memclr4' from stdio.o(.text).
Adding TA veneer (4 bytes, Inline) for call to '_mutex_initialize' from stdio.o(.text).
Adding TA veneer (4 bytes, Inline) for call to '__rt_raise' from stdio.o(.text).
Adding AT veneer (8 bytes, Inline) for call to '__raise' from rt_raise.o(.text).
Adding TA veneer (4 bytes, Inline) for call to '__heap_extend' from malloc.o(.text).
Adding TA veneer (4 bytes, Inline) for call to '__user_perproc_libspace' from malloc.o(.text).
Adding TA veneer (8 bytes, Short) for call to '__rt_exit' from exit.o(.text).
Adding TA veneer (4 bytes, Inline) for call to '_fp_init' from lib_init.o(.text).
Adding TA veneer (4 bytes, Inline) for call to '__ARM_argv_veneer' from lib_init.o(.text).
Adding TA veneer (4 bytes, Inline) for call to '_sys_exit' from abort.o(.text).

```

14 Veneer(s) (total 80bytes) added to the image.

4.3.2 Basic rules for C and C++ interworking

The following rules apply to interworking within an application:

- You must use the `--apcs /interwork` command-line option to compile any C or C++ modules that contain functions that might return to the other instruction set.
- You must use the `--apcs /interwork` command-line option to compile any C or C++ modules that contain indirect or virtual function calls that might be to functions in the other instruction set.
- Never make indirect calls, such as calls using function pointers, to non-interworking code from code in the other state.
- If any input object contains Thumb code, the linker selects the Thumb runtime libraries. These are built for interworking.

If you specify one of your own libraries explicitly on the linker command line you must ensure that it is an appropriate interworking library.

Note

If a C or C++ module contains functions marked with `#pragma arm` or `#pragma thumb`, you must compile the module with `--apcs /interwork`. This ensures that the functions can be called successfully from the other (ARM or Thumb) state.

4.3.3 Pointers to functions in Thumb state

If you have a Thumb function, that is a function consisting of Thumb code, and that runs in Thumb state, then any pointer to that function must have the least significant bit set. This ensures that interworking works correctly.

When the linker relocates a value of a label referring to a Thumb instruction, it automatically sets the least significant bit of the relocated value. The linker cannot do this if you use absolute addresses to Thumb functions.

Therefore, if you have to use an absolute address to a Thumb function in your code, you must add one to the address. For example, you might have a table of pointers to Thumb functions, such as that shown in Example 4-5 on page 4-17.

See *Assembly language interworking* on page 4-7 for more details.

Example 4-5 Absolute addresses to Thumb functions

```
typedef int (*FN)();

myfunc() {
    FN fnptrs[] = {
        (FN)(0x8084 + 1), // Valid Thumb address
        (FN)(0x8074)      // Invalid Thumb address
    };
    FN* myfunctions = fnptrs;

    myfunctions[0](); // Call OK
    myfunctions[1](); // Call Fails
}

```

4.3.4 Using two versions of the same function

You can have two functions with the same name, one compiled for ARM and the other for Thumb.

ARM/Thumb synonyms

The linker enables multiple definitions of a symbol to coexist in an image, only if each definition is associated with a different processor state. The linker applies the following rules when a reference is made to a symbol with ARM/Thumb synonyms:

- B, BL, or BLX instructions to a symbol from ARM state resolve to the ARM definition
- B, BL, or BLX instructions to a symbol from Thumb state resolve to the Thumb definition.

Any other reference to the symbol resolves to the first definition encountered by the linker. The linker produces a warning that specifies the chosen symbol.

4.4 Assembly language interworking using veneers

The assembly language ARM/Thumb interworking method described in *Assembly language interworking* on page 4-7 carried out all the necessary intermediate processing. There was no requirement for the linker to insert interworking veneers.

This section describes how you can make use of interworking veneers to:

- interwork between assembly language modules, see *Assembly-only interworking using veneers*
- interwork between assembly language and C or C++ modules, see *C, C++, and assembly language interworking using veneers* on page 4-20.

See *Linker generated veneers* on page 4-5 for more details on veneers.

4.4.1 Assembly-only interworking using veneers

You can write assembly language ARM/Thumb interworking code to make use of interworking veneers generated by the linker. To do this, you write:

- A caller routine like any non-interworking routine, using a BL instruction to make the call. A caller routine can be assembled with either `--apcs /interwork` or `--apcs /nointerwork`.

Note

The range of a BL instruction is 32MB in ARM state, and 4MB in Thumb state. During development, your application might have calls to targets that are beyond reach, or calls to functions in another state. The linker automatically inserts a veneer in these cases. The veneer becomes the intermediate target of the original BL, and the veneer code then sets the PC to the desired destination address.

- A callee routine using a BX instruction to return. A callee routine must be assembled with `--apcs /interwork`. Also, you might have to export the function label of the routine, for example, `EXPORT ThumbSub` (see Example 4-6 on page 4-19). Where appropriate, the assembler code must conform to the AAPCS.

This is generally only necessary in ARMv4T, or if the caller and callee are widely separated or in different areas. In ARMv5T and later, if the caller and callee are sufficiently close together, no veneers are necessary.

Example of assembly language interworking using veneers

Example 4-6 shows the code to set registers `r0` to `r2` to the values 1, 2, and 3 respectively. Registers `r0` and `r2` are set by the ARM code. `r1` is set by the Thumb code. Observe that:

- the code must be assembled with the option `--apcs /interwork`
- a `BX lr` instruction is used to return from the subroutine, instead of the usual `MOV pc,lr`.

Example 4-6

```

; *****
; arm.s
; *****

PRESERVE8

AREA    Arm, CODE, READONLY    ; Name this block of code.
IMPORT  ThumbProg
ENTRY   ; Mark 1st instruction to call.
ARMProg
    MOV    r0, #1                ; Set r0 to show in ARM code.
    BL     ThumbProg             ; Call Thumb subroutine.
    MOV    r2, #3                ; Set r2 to show returned to ARM.
                                ; Terminate execution.
    MOV    r0, #0x18             ; angel_SWIreason_ReportException
    LDR    r1, =0x20026          ; ADP_Stopped_ApplicationExit
    SWI    0x123456              ; ARM semihosting SWI
    END

; *****
; thumb.s
; *****
AREA    Thumb, CODE, READONLY   ; Name this block of code.
CODE16                               ; Subsequent instructions are Thumb.
EXPORT  ThumbProg
ThumbProg
    MOV    r1, #2                ; Set r1 to show reached Thumb code.
    BX     lr                    ; Return to ARM subroutine.
    END                          ; Mark end of this file.

```

Follow these steps to build and link the modules, and examine the interworking veneers:

1. Type `armasm -g arm.s` to assemble the ARM code.

2. Type `armasm --thumb -g --apcs /interwork thumb.s` to assemble the Thumb code.
3. Type `armlink arm.o thumb.o -o count` to link the two object files.
4. Run the image using a compatible debugger (for example, RealView Debugger or AXD) with an appropriate debug target.

You can see the interworking veneer that is inserted by the linker in the disassembled code shown in Example 4-7. The veneer is inserted on the next word boundary, and starts at address `0x0000801C`.

Example 4-7

```

ARMProg:
00008000 E3A00001 MOV     r0,#1
00008004 EB000004 BL      0x801c
00008008 E3A02003 MOV     r2,#3
0000800C E3A00018 MOV     r0,#0x18
00008010 E59F1000 LDR      r1,0x8018
00008014 EF123456 SWI      0x123456
00008018 00020026 <Data> '&' 0x00 0x02 0x00
0000801C E28FC001 ADR      r12,{pc}+9 ; #0x8025
00008020 E12FFF1C BX       r12
ThumbProg:
00008024      2102 MOV     r1,#2
00008026      4770 BX      r14

```

4.4.2 C, C++, and assembly language interworking using veneers

C and C++ code compiled to run in one state can call assembly language code designed to run in the other state, and vice versa. To do this, write the caller routine as any non-interworking routine and, if calling from assembly language, use a BL instruction to make the call (see Example 4-8 on page 4-21). Then:

- if the callee routine is in C, compile it using `--apcs /interwork`
- if the callee routine is in assembly language, assemble with the `--apcs /interwork` option and return using BX lr.

———— Note ————

Any assembly language code or user library code used in this manner must conform to the AAPCS where appropriate.

Example 4-8

```

/*****
 *      thumb.c      *
 *****/
#include <stdio.h>
extern int arm_function(int);
int main(void)
{
    int i = 1;
    printf("i = %d\n", i);
    printf("And now i = %d\n", arm_function(i));
    return (0);
}

; *****
; arm.s
; *****
PRESERVE8
AREA Arm, CODE, READONLY ; Name this block of code.
EXPORT arm_function
arm_function
    ADD    r0, r0, #4        ; Add 4 to first parameter.
    BX     lr                ; Return
    END

```

Follow these steps to build and link the modules:

1. Type `armcc --thumb -g -c --apcs /interwork thumb.c` to compile the Thumb code.
2. Type `armasm -g --apcs /interwork arm.s` to assemble the ARM code.
3. Type `armlink arm.o thumb.o -o add --info veneers` to link the two object files and view the size of the interworking veneers.
4. Run the image using a compatible debugger (for example, RealView Debugger or AXD) with an appropriate debug target.

Chapter 5

Mixing C, C++, and Assembly Language

This chapter describes how to write mixed C, C++, and ARM® assembly language code. It also describes how to use the ARM inline and embedded assemblers from C and C++. It contains the following sections:

- *Using the inline and embedded assemblers* on page 5-2
- *Accessing C global variables from assembly code* on page 5-4
- *Using C header files from C++* on page 5-5
- *Calling between C, C++, and ARM assembly language* on page 5-7.

5.1 Using the inline and embedded assemblers

The inline and embedded assemblers that are built into the ARM compiler enable you to use features of the target processor that cannot be accessed directly from C or C++. For example:

- saturating arithmetic (see *RealView Compilation Tools v2.2 Assembler Guide*)
- custom coprocessors
- the *Program Status Register* (PSR).

This section includes:

- *Features of the inline assembler*
- *Features of the embedded assembler*
- *Differences between inline and embedded assembly code* on page 5-3.

For more details, see the chapter on inline and embedded assemblers in *RealView Compilation Tools v2.2 Compiler and Libraries Guide*.

5.1.1 Features of the inline assembler

The inline assembler supports very flexible interworking with C and C++. Any register operand can be an arbitrary C or C++ expression. The inline assembler also expands complex instructions and optimizes the assembly language code.

———— Note —————

Inline assembly language is subject to optimization by the compiler if you use one of the multi-optimization compiler options -01, -02, or -03.

The inline assembler for ARM code implements most of the ARM instruction set including generic coprocessor instructions, halfword instructions and long multiply.

5.1.2 Features of the embedded assembler

The embedded assembler provides unrestricted, low-level access to the target processor, and enables you to use the C and C++ preprocessor directives, and gives easy access to structure member offsets.

The embedded assembler enables you to use the full ARM assembler instruction set, including assembler directives. Embedded assembly code is assembled separately from the C and C++ code. A compiled object is produced that is then combined with the object from the compilation of the C and C++ source.

The embedded assembler is supported in both ARM and Thumb® code. See *RealView Compilation Tools v2.2 Assembler Guide* for details of the ARM/Thumb instruction set.

5.1.3 Differences between inline and embedded assembly code

Table 5-1 summarizes the main differences between inline assembler and embedded assembler.

Table 5-1 Differences between inline and embedded assembler

Feature	Embedded assembler	Inline assembler
Instruction set	ARM and Thumb.	ARM only.
ARM assembler directives	All supported.	None supported.
C/C++ expressions	Constant expressions only.	Full C/C++ expressions.
Optimization of assembly code	No optimization.	Full optimization.
Inlining	No.	Possible.
Register access	Specified physical registers are used. You can also use PC, LR and SP.	Uses virtual registers. Using sp (r13), lr (r14), and pc (r15) gives an error.
Return instructions	You must add them in your code.	Generated automatically. (The BX, BXJ, and BLX instructions are not supported.)
BKPT instruction	Supported directly.	Not supported.

Note

A list of differences between embedded assembler and C/ C++ is provided in the chapter on inline and embedded assemblers in *RealView Compilation Tools v2.2 Compiler and Libraries Guide*.

5.2 Accessing C global variables from assembly code

Global variables can only be accessed indirectly, through their address. To access a global variable, use the `IMPORT` directive to do the import and then load the address into a register. You can access the global variable with load and store instructions, depending on its type.

For **unsigned** variables, for example, use:

- `LDRB/STRB` for **char**
- `LDRH/STRH` for **short**
- `LDR/STR` for **int**.

For **signed** variables, use the equivalent signed instruction, such as `LDRSB` and `LDRSH`.

Small structures of less than eight words can be accessed as a whole using the `LDM` and `STM` instructions. Individual members of structures can be accessed by a load or store instruction of the appropriate type. You must know the offset of a member from the start of the structure in order to access it.

Example 5-1 loads the address of the integer global variable `globvar` into `r1`, loads the value contained in that address into `r0`, adds 2 to it, then stores the new value back into `globvar`.

Example 5-1 Accessing global variables

```

PRESERVE8

AREA    globals, CODE, READONLY

EXPORT  asmsubroutine
IMPORT  globvar

asmsubroutine
    LDR  r1, =globvar    ; read address of globvar into
                        ; r1 from literal pool
    LDR  r0, [r1]
    ADD  r0, r0, #2
    STR  r0, [r1]
    MOV  pc, lr
    END

```

For full details on the instructions available in ARM or Thumb code, see *RealView Compilation Tools v2.2 Assembler Guide*.

5.3 Using C header files from C++

C header files must be wrapped in extern "C" directives before they are called from C++.

This section describes:

- *Including system C header files*
- *Including your own C header files* on page 5-6.

5.3.1 Including system C header files

You do not have to take any special steps to include standard system C header files, such as `stdio.h`. The standard C header files already contain the appropriate extern "C" directives. For example:

```
#include <stdio.h>
int main()
{
    ...          // C++ code
    return 0;
}
```

If you include headers using this syntax, all library names are placed in the global namespace.

The C++ standard specifies that the functionality of the C header files is available through C++ specific header files. These files are installed in `install_directory\RVCT\Data\2.2\build_num\include\platform`, together with the standard C header files, and can be referenced in the usual way. For example:

```
#include <cstdio>
```

In ARM C++, these headers `#include` the C headers. If you include headers using this syntax, all C++ standard library names are defined in the namespace `std`, including the C library names. This means that you must qualify all the library names by using one of the following methods:

- specify the standard namespace, for example:
`std::printf("example\n");`
- use the C++ keyword **using** to import a name to the global namespace:
`using namespace std;
printf("example\n");`
- use the compiler option `--using_std`.

5.3.2 Including your own C header files

To include your own C header files, you must wrap the `#include` directive in an extern "C" statement. You can do this in the following ways:

- When the file is `#included` (shown in Example 5-2).
- By adding the extern "C" statement to the header file (shown in Example 5-3).

Example 5-2 Directive before include file

```
// C++ code

extern "C" {
#include "my-header1.h"
#include "my-header2.h"
}

int main()
{
    // ...
    return 0;
}
```

Example 5-3 Directive in file header

```
/* C header file */

#ifdef __cplusplus    /* Insert start of extern C construct */
extern "C" {
#endif

/* Body of header file */

#ifdef __cplusplus    /* Insert end of extern C construct */
}                    /* The C header file can now be */
#endif               /* included in either C or C++ code. */
```

5.4 Calling between C, C++, and ARM assembly language

This section provides examples that can help you to call C and assembly language code from C++, and to call C++ code from C and assembly language. It also describes calling conventions and data types, and includes:

- *General rules for calling between languages*
- *Information specific to C++* on page 5-8
- *Examples of calling between languages* on page 5-9.

You can mix calls between C and C++ and assembly language routines provided you comply with the AAPCS. For more information, see:

- the *Procedure Call Standard for the ARM Architecture* specification, `aapcs.pdf`, in `install_directory\Documentation\Specifications\...`
- Appendix A *Using the Procedure Call Standard*.

Note

The information in this section is implementation dependent and might change in future releases.

5.4.1 General rules for calling between languages

The following general rules apply to calling between C, C++, and assembly language. For more details, see *RealView Compilation Tools v2.2 Compiler and Libraries Guide*.

The embedded assembler and compliance with the *Application Binary Interface (ABI) for the ARM Architecture (base standard)* [BSABI] make mixed language programming easier to implement. These assist you with:

- name mangling, using the `__cpp` keyword
- the way the implicit **this** parameter is passed
- the way virtual functions are called
- the representation of references
- the layout of C++ class types that have base classes or virtual member functions
- the passing of class objects that are not plain old data structures.

The following general rules apply to mixed language programming:

- Use C calling conventions.

- In C++, nonmember functions can be declared as extern "C" to specify that they have C linkage. In this release of *RealView® Compilation Tools* (RVCT), having C linkage means that the symbol defining the function is not mangled. C linkage can be used to implement a function in one language and call it from another.

———— **Note** —————

Functions that are declared extern "C" cannot be overloaded.

—————

- Assembly language modules must conform to the appropriate AAPCS standard for the memory model used by the application.

The following rules apply to calling C++ functions from C and assembly language:

- To call a global (nonmember) C++ function, declare it extern "C" to give it C linkage.
- Member functions (both static and non-static) always have mangled names. Using the `__cpp` keyword of the embedded assembler means that you do not have to find the mangled names manually.
- C++ inline functions cannot be called from C unless you ensure that the C++ compiler generates an out-of-line copy of the function. For example, taking the address of the function results in an out-of-line copy.
- Nonstatic member functions receive the implicit **this** parameter as a first argument in `r0`, or as a second argument in `r1` if the function returns a non **int**-like structure. Static member functions do not receive an implicit **this** parameter.

5.4.2 Information specific to C++

The following information applies specifically to C++.

C++ calling conventions

ARM C++ uses the same calling conventions as ARM C with one exception:

- Nonstatic member functions are called with the implicit **this** parameter as the first argument, or as the second argument if the called function returns a non **int**-like **struct**. This might change in future implementations.

C++ data types

ARM C++ uses the same data types as ARM C with the following exceptions and additions:

- C++ objects of type **struct** or **class** have the same layout that is expected from ARM C if they have no base classes or virtual functions. If such a **struct** has neither a user-defined copy assignment operator nor a user-defined destructor, it is a plain old data structure.
- References are represented as pointers.
- No distinction is made between pointers to C functions and pointers to C++ (nonmember) functions.

Symbol name mangling

The linker unmangles symbol names in messages.

C names must be declared as extern "C" in C++ programs. This is done already for the ARM ISO C headers. See *Using C header files from C++* on page 5-5 for more information.

5.4.3 Examples of calling between languages

The following sections contain code examples that demonstrate how to mix language calls:

- *Calling assembly language from C* on page 5-10
- *Calling C from assembly language* on page 5-11
- *Calling C from C++* on page 5-12
- *Calling assembly language from C++* on page 5-13
- *Calling C++ from C* on page 5-14
- *Calling C++ from assembly language* on page 5-15
- *Calling C++ from C or assembly language* on page 5-17
- *Passing a reference between C and C++* on page 5-16.

The examples assume the default non software-stack checking AAPCS variant because they perform stack operations without checking for stack overflow.

Calling assembly language from C

Example 5-4 and Example 5-5 show a C program that uses a call to an assembly language subroutine to copy one string over the top of another string.

Example 5-4 Calling assembly language from C

```
#include <stdio.h>
extern void strcpy(char *d, const char *s);
int main()
{   const char *srcstr = "First string - source ";
    char dststr[] = "Second string - destination ";
    /* dststr is an array since we're going to change it */
    printf("Before copying:\n");
    printf(" %s\n %s\n",srcstr,dststr);
    strcpy(dststr,srcstr);
    printf("After copying:\n");
    printf(" %s\n %s\n",srcstr,dststr);
    return (0);
}
```

Example 5-5 Assembly language string copy subroutine

```
PRESERVE8

AREA    SCopy, CODE, READONLY
EXPORT strcpy

strcpy    ; r0 points to destination string.
          ; r1 points to source string.
    LDRB r2, [r1],#1 ; Load byte and update address.
    STRB r2, [r0],#1 ; Store byte and update address.
    CMP r2, #0      ; Check for zero terminator.
    BNE strcpy      ; Keep going if not.
    MOV pc,lr       ; Return.
END
```

Example 5-4 is located in the main examples directory, in ...\\asm as strtest.c and scopy.s.

Follow these steps to build the example from the command line:

1. Type `armasm -g scopy.s` to build the assembly language source.
2. Type `armcc -c -g strtest.c` to build the C source.

3. Type `armlink strttest.o scopy.o -o strttest` to link the object files.
4. Run the image using a compatible debugger (for example, AXD or RealView Debugger) with an appropriate debug target.

Calling C from assembly language

Example 5-6 and Example 5-7 show how to call C from assembly language.

Example 5-6 Defining the function in C

```
int g(int a, int b, int c, int d, int e)
{
    return a + b + c + d + e;
}
```

Example 5-7 Assembly language call

```
; int f(int i) { return g(i, 2*i, 3*i, 4*i, 5*i); }
```

```
PRESERVE8
```

```
EXPORT f
AREA f, CODE, READONLY
IMPORT g          ; i is in r0
STR lr, [sp, #-4]! ; preserve lr
ADD r1, r0, r0    ; compute 2*i (2nd param)
ADD r2, r1, r0    ; compute 3*i (3rd param)
ADD r3, r1, r2    ; compute 5*i
STR r3, [sp, #-4]! ; 5th param on stack
ADD r3, r1, r1    ; compute 4*i (4th param)
BL g             ; branch to C function
ADD sp, sp, #4    ; remove 5th param
LDR pc, [sp], #4  ; return
END
```

Calling C from C++

Example 5-8 and Example 5-9 show how to call C from C++.

Example 5-8 Calling a C function from C++

```
struct S {           // has no base classes
                    // or virtual functions
    S(int s) : i(s) { }
    int i;
};
extern "C" void cfunc(S *);
// declare the C function to be called from C++
int f(){
    S s(2);          // initialize 's'
    cfunc(&s);        // call 'cfunc' so it can change 's'
    return s.i * 3;
}
```

Example 5-9 Defining the function in C

```
struct S {
    int i;
};
void cfunc(struct S *p) {
    /* the definition of the C function to be called from C++ */
    p->i += 5;
}
```

Calling assembly language from C++

Example 5-10 and Example 5-11 show how to call assembly language from C++.

Example 5-10 Calling assembly language from C++

```

struct S {          // has no base classes
                  // or virtual functions
    S(int s) : i(s) { }
    int i;
};

extern "C" void asmfunc(S *); // declare the Asm function
                              // to be called

int f() {
    S s(2);                // initialize 's'
    asmfunc(&s);           // call 'asmfunc' so it
                          // can change 's'
    return s.i * 3;
}

```

Example 5-11 Defining the assembly language function

```

PRESERVE8

AREA Asm, CODE
EXPORT asmfunc
asmfunc          ; the definition of the Asm
LDR r1, [r0]     ; function to be called from C++
ADD r1, r1, #5
STR r1, [r0]
MOV pc, lr
END

```

Calling C++ from C

Example 5-12 and Example 5-13 show how to call C++ from C.

Example 5-12 Defining the function to be called in C++

```
struct S {          // has no base classes or virtual functions
    S(int s) : i(s) { }
    int i;
};

extern "C" void cppfunc(S *p) {
    // Definition of the C++ function to be called from C.
    // The function is written in C++, only the linkage is C
    p->i += 5;          //
}
```

Example 5-13 Declaring and calling the function in C

```
struct S {
    int i;
};

extern void cppfunc(struct S *p);
/* Declaration of the C++ function to be called from C */

int f(void) {
    struct S s;
    s.i = 2;          /* initialize 's' */
    cppfunc(&s);      /* call 'cppfunc' so it */
                     /* can change 's' */
    return s.i * 3;
}
```

Calling C++ from assembly language

Example 5-14 and Example 5-15 show how to call C++ from assembly language.

Example 5-14 Defining the function to be called in C++

```
struct S {           // has no base classes or virtual functions
    S(int s) : i(s) { }
    int i;
};
extern "C" void cppfunc(S * p) {
    // Definition of the C++ function to be called from ASM.
    // The body is C++, only the linkage is C
    p->i += 5;
}
```

In ARM assembly language, import the name of the C++ function and use a Branch with Link (BL) instruction to call it:

Example 5-15 Defining assembly language function

```
AREA Asm, CODE
IMPORT cppfunc          ; import the name of the C++
                        ; function to be called from Asm

EXPORT f

f
    STMFD sp!,{lr}
    MOV    r0,#2
    STR    r0,[sp,#-4]!  ; initialize struct
    MOV    r0,sp         ; argument is pointer to struct
    BL     cppfunc       ; call 'cppfunc' so it can change
                        ; the struct

    LDR    r0, [sp], #4
    ADD    r0, r0, r0, LSL #1
    LDMFD  sp!,{pc}
    END
```

Passing a reference between C and C++

Example 5-16 and Example 5-17 show how to pass a reference between C and C++.

Example 5-16 Defining the C++ function

```
extern "C" int cfunc(const int&);  
// Declaration of the C function to be called from C++  
  
extern "C" int cppfunc(const int& r) {  
// Definition of the C++ to be called from C.  
    return 7 * r;  
}  
  
int f() {  
    int i = 3;  
    return cfunc(i);    // passes a pointer to 'i'  
}
```

Example 5-17 Defining the C function

```
extern int cppfunc(const int*);  
/* declaration of the C++ to be called from C */  
  
int cfunc(const int *p) {  
/* definition of the C function to be called from C++ */  
    int k = *p + 4;  
    return cppfunc(&k);  
}
```

Calling C++ from C or assembly language

The code in Example 5-18, Example 5-19 and Example 5-20 on page 5-18 demonstrates how to call a non-static, non-virtual C++ member function from C or assembly language. Use the assembler output from the compiler to locate the mangled name of the function.

Example 5-18 Calling a C++ member function

```

struct T {
    T(int i) : t(i) { }
    int t;
    int f(int i);
};

int T::f(int i) { return i + t; }
// Definition of the C++ function to be called from C.

extern "C" int cfunc(T*);
// declaration of the C function to be called from C++

int f() {
    T t(5);                // create an object of type T
    return cfunc(&t);
}

```

Example 5-19 Defining the C function

```

struct T;

extern int _ZN1T1fEi(struct T*, int);
/* the mangled name of the C++ */
/* function to be called */

int cfunc(struct T* t) {
/* Definition of the C function to be called from C++. */
    return 3 * _ZN1T1fEi(t, 2); /* like '3 * t->f(2)' */
}

```

Example 5-20 Implementing the function in assembly language

```

EXPORT cfunc
AREA foo, CODE
IMPORT _ZN1T1fEi

cfunc
    STMFD sp!,{lr}          ; r0 already contains the object pointer
    MOV r1, #2
    BL _ZN1T1fEi
    ADD r0, r0, r0, LSL #1   ; multiply by 3
    LDMFD sp!,{pc}
    END

```

Alternatively, you can implement Example 5-18 on page 5-17 and Example 5-20 using embedded assembly, as shown in Example 5-21. In this example, the `__cpp` keyword is used to reference the function. Therefore, you do not have to know the mangled name of the function.

Example 5-21 Implementing the function in embedded assembly

```

struct T {
    T(int i) : t(i) { }
    int t;
    int f(int i);
};

int T::f(int i) { return i + t; }

// Definition of asm function called from C++
__asm int asm_func(T*) {
    STMFD sp!, {lr}
    MOV r1, #2;
    BL __cpp(T::f);
    ADD r0, r0, r0, LSL #1 ; multiply by 3
    LDMFD sp!, {pc}
}

int f() {
    T t(5); // create an object of type T
    return asm_func(&t);
}

```

Chapter 6

Handling Processor Exceptions

This chapter describes how to handle the different types of exception supported by ARM® processors. It contains the following sections:

- *About processor exceptions* on page 6-2
- *Determining the processor state* on page 6-6
- *Entering and leaving an exception* on page 6-8
- *Handling an exception* on page 6-13
- *Installing an exception handler* on page 6-14
- *SWI handlers* on page 6-19
- *Interrupt handlers* on page 6-29
- *Reset handlers* on page 6-39
- *Undefined Instruction handlers* on page 6-40
- *Prefetch Abort handler* on page 6-41
- *Data Abort handler* on page 6-42
- *Chaining exception handlers* on page 6-44
- *System mode* on page 6-46.

6.1 About processor exceptions

During the normal flow of execution through a program, the program counter (PC) increases sequentially through the address space, with branches to nearby labels or branch and links to subroutines.

Processor exceptions occur when this normal flow of execution is diverted, to enable the processor to handle events generated by internal or external sources. Examples of such events are:

- externally generated interrupts
- an attempt by the processor to execute an undefined instruction
- accessing privileged operating system functions.

It is necessary to preserve the previous processor status when handling such exceptions, so that execution of the program that was running when the exception occurred can resume when the appropriate exception routine has completed.

This section includes:

- *Types of exception*
- *The vector table* on page 6-3
- *Use of modes and registers by exceptions* on page 6-3
- *Exception priorities* on page 6-4.

6.1.1 Types of exception

Table 6-1 shows the different types of exception recognized by ARM processors.

Table 6-1 Exception types

Exception	Description
Reset	Occurs when the processor reset pin is asserted. This exception is only expected to occur for signaling power-up, or for resetting as if the processor has powered up. A soft reset can be done by branching to the reset vector (0x0000).
Undefined Instruction	Occurs if neither the processor, nor any attached coprocessor, recognizes the currently executing instruction.
Software Interrupt (SWI)	This is a user-defined synchronous interrupt instruction. It enables a program running in User mode, for example, to request privileged operations that run in Supervisor mode, such as an RTOS function.
Prefetch Abort	Occurs when the processor attempts to execute an instruction that was not fetched, because the address was illegal (see <i>Illegal addresses</i> on page 6-3).

Table 6-1 Exception types (continued)

Exception	Description
Data Abort	Occurs when a data transfer instruction attempts to load or store data at an illegal address (see <i>Illegal addresses</i>).
IRQ	Occurs when the processor external interrupt request pin is asserted (LOW) and the I bit in the CPSR is clear.
FIQ	Occurs when the processor external fast interrupt request pin is asserted (LOW) and the F bit in the CPSR is clear.

Illegal addresses

An illegal virtual address is one that does not currently correspond to an address in physical memory, or one that the memory management subsystem has determined is inaccessible to the processor in its current mode.

6.1.2 The vector table

The *vector table* controls processor exception handling. The vector table is a reserved area of 32 bytes, usually at the bottom of the memory map. It has one word of space allocated to each exception type, and one word that is currently reserved.

This is not enough space to contain the full code for a handler, so the vector entry for each exception type typically contains a branch instruction or load PC instruction to continue execution with the appropriate handler.

6.1.3 Use of modes and registers by exceptions

Typically, an application runs in User mode, but servicing exceptions requires a privileged mode. An exception changes the processor mode, and this in turn means that each exception handler has access to a certain subset of the banked registers:

- its own r13 or *Stack Pointer* (*sp_mode*)
- its own r14 or *Link Register* (*lr_mode*)
- its own *Saved Program Status Register* (*spsr_mode*).

In the case of an FIQ, each exception handler has access to five more general purpose registers (r8_FIQ to r12_FIQ).

Each exception handler must ensure that other registers are restored to their original contents on exit. You can do this by saving the contents of any registers that the handler has to use onto its stack and restoring them before returning. If you are using Angel or RealView ARMulator® ISS, the required stacks are set up for you. Otherwise, you must set them up yourself.

———— **Note** —————

As supplied, the assembler does *not* predeclare symbolic register names of the form *register_mode*. To use this form, you must declare the appropriate symbolic names with the RN assembler directive, for example, `1r_FIQ RN r14` declares the symbolic register name `1r_FIQ` for `r14`. See the directives chapter in *RealView Compilation Tools v2.2 Assembler Guide* for more information on the RN directive.

6.1.4 Exception priorities

When several exceptions occur simultaneously, they are serviced in a fixed order of priority. Each exception is handled in turn before execution of the user program continues. It is not possible for all exceptions to occur concurrently. For example, the Undefined Instruction and SWI exceptions are mutually exclusive because they are both triggered by executing an instruction.

Table 6-2 shows the exceptions, their corresponding processor modes and their handling priorities.

Table 6-2 Exception priorities

Vector address	Exception type	Exception mode	Priority (1=high, 6=low)
0x0	Reset	Supervisor (SVC)	1
0x4	Undefined Instruction	Undef	6
0x8	Software Interrupt (SWI)	Supervisor (SVC)	6
0xC	Prefetch Abort	Abort	5
0x10	Data Abort	Abort	2
0x14	<i>Reserved</i>	<i>Not applicable</i>	<i>Not applicable</i>
0x18	Interrupt (IRQ)	Interrupt (IRQ)	4
0x1C	Fast Interrupt (FIQ)	Fast Interrupt (FIQ)	3

Because the Data Abort exception has a higher priority than the FIQ exception, the Data Abort is actually registered before the FIQ is handled. The Data Abort handler is entered, but control is then passed immediately to the FIQ handler. When the FIQ has been handled, control returns to the Data Abort handler. This means that the data transfer error does not escape detection as it would if the FIQ were handled first.

6.2 Determining the processor state

An exception handler might have to determine whether the processor was in ARM or Thumb® state when the exception occurred.

SWI handlers, especially, might have to read the processor state. This is done by examining the SPSR T-bit. This bit is set for Thumb state and clear for ARM state.

Both ARM and Thumb instruction sets have the SWI instruction. When calling SWIs from Thumb state, you must consider the following:

- The instruction address is at (lr-2), rather than (lr-4).
- The instruction itself is 16-bit, and so requires a halfword load (see Figure 6-1).
- The SWI number is held in 8 bits instead of the 24 bits in ARM state.




Figure 6-1 Thumb SWI instruction

Example 6-1 on page 6-7 shows ARM code that handles a SWI from both sources.

Consider the following:

- Each of the do_swi_x routines could carry out a switch to Thumb state and back again to improve code density if required.
- You can replace the jump table by a call to a C function containing a switch() statement to implement the SWIs.
- It is possible for a SWI number to be handled differently depending on the state it is called from.
- The range of SWI numbers accessible from Thumb state can be increased by calling SWIs dynamically (as described in *SWI handlers* on page 6-19).

Example 6-1 SWI handler

```

T_bit    EQU    0x20                                ; Thumb bit of CPSR/SPSR, that is,
                                                    ; bit 5.
:
:
SWIHandler
    STMFD    sp!, {r0-r3,r12,lr}                ; Store registers.
    MRS      r0, spsr                            ; Move SPSR into
                                                    ; general purpose register.
    TST      r0, #T_bit                          ; Occurred in Thumb state?
    LDRNEH   r0,[lr,#-2]                        ; Yes: load halfword and...
    BICNE    r0,r0,#0xFF00                      ; ...extract comment field.
    LDREQ    r0,[lr,#-4]                        ; No: load word and...
    BICEQ    r0,r0,#0xFF000000                  ; ...extract comment field.

                                                    ; r0 now contains SWI number

    CMP      r0, #MaxSWI                        ; Rangecheck
    LDRLS    pc, [pc, r0, LSL#2]                ; Jump to the appropriate routine.
    B        SWIOutOfRange

switable
    DCD      do_swi_1
    DCD      do_swi_2
    :
    :
do_swi_1
    ; Handle the SWI.
    LDMFD    sp!, {r0-r3,r12,pc}^                ; Restore the registers and return.
do_swi_2
    :

```

6.3 Entering and leaving an exception

This section describes the processor response to an exception, and how to return to the place where an exception occurred after the exception has been handled. The method for returning is different depending on the exception type (see *Types of exception* on page 6-2).

Processors that support Thumb state use the same basic exception handling mechanism as processors that do not support Thumb state. An exception causes the next instruction to be fetched from the appropriate vector table entry.

The same vector table is used for exceptions in both Thumb state and ARM state. An initial step (to switch to ARM state) is added to the exception handling procedure described in *The processor response to an exception*.

In the following descriptions, it is clearly marked if there are further considerations that you must take into account when writing exception handlers suitable for use on processors that support Thumb state.

This section includes:

- *The processor response to an exception*
- *Returning from an exception handler* on page 6-9
- *The return address and return instruction* on page 6-10.

6.3.1 The processor response to an exception

When an exception is generated, the processor performs the following actions:

1. Copies the *Current Program Status Register* (CPSR) into the *Saved Program Status Register* (SPSR) for the mode in which the exception is to be handled. This saves the current mode, interrupt mask, and condition flags.
2. Switches to ARM state, if it is currently in Thumb state.
3. Changes the appropriate CPSR mode bits in order to:
 - change to the appropriate mode, and map in the appropriate banked registers for that mode
 - disable interrupts. IRQs are disabled when any exception occurs. FIQs are disabled when a FIQ occurs and on reset.
4. Sets *lr_mode* to the return address, as defined in *The return address and return instruction* on page 6-10.
5. Sets the PC to the vector address for the exception.

For *ARM processors that do not support Thumb*, this forces a branch to the appropriate exception handler.

For *processors that support Thumb*, the switch from Thumb state to ARM state in step 2 ensures that the ARM instruction installed at this vector address (either a branch or a PC-relative load) is correctly fetched, decoded, and executed. This forces a branch to a top-level veneer that you must write in ARM code.

6.3.2 Returning from an exception handler

The method used to return from an exception depends on whether the exception handler uses stack operations or not. In both cases, to return execution to the place where the exception occurred an exception handler must:

- restore the CPSR from *spsr_mode*
- restore the PC using the return address stored in *lr_mode*.

For a simple return that does not require the destination mode registers to be restored from the stack, the exception handler carries out these operations by performing a data processing instruction with:

- the S flag set
- the PC as the destination register.

The return instruction required depends on the type of exception. See *The return address and return instruction* on page 6-10 for instructions on how to return from each exception type.

————— Note —————

You do not have to return from the reset handler because the reset handler executes your main code directly.

If the exception handler entry code uses the stack to store registers that must be preserved while it handles the exception, it can return using a load multiple instruction with the ^ qualifier. For example, an exception handler can return in one instruction using:

```
LDMFD sp!,{r0-r12,pc}^
```

To do this, the exception handler must save the following onto the stack:

- all the work registers in use when the handler is invoked
- the link register, modified to produce the same effect as the data processing instructions described in *The return address and return instruction* on page 6-10.

The \wedge qualifier specifies that the CPSR is restored from the SPSR. It must be used only from a privileged mode. See the description of how to implement stacks with LDM and STM in the *RealView Compilation Tools v2.2 Assembler Guide* for more general information on stack operations.

6.3.3 The return address and return instruction

The actual location pointed to by the PC when an exception is taken depends on the exception type. The return address might not necessarily be the next instruction pointed to by the PC.

If an exception occurs in ARM state, the processor stores $(PC-4)$ in *lr_mode*. However, for exceptions that occur in Thumb state, the processor automatically stores a different value for each of the exception types. This adjustment is required because Thumb instructions take up only a halfword, rather than the full word that ARM instructions occupy.

If this correction were not made by the processor, the handler would have to determine the original state of the processor, and use a different instruction to return to Thumb code rather than ARM code. By making this adjustment, however, the processor enables the handler to have a single return instruction that returns correctly, regardless of the processor state (ARM or Thumb) at the time the exception occurred.

The following sections detail the instructions to return correctly from handling code for each type of exception.

Returning from SWI and Undefined Instruction handlers

The SWI and Undefined Instruction exceptions are generated by the instruction itself, so the PC is not updated when the exception is taken. The processor stores $(PC-4)$ in *lr_mode*. This makes *lr_mode* point to the next instruction to be executed. Restoring the PC from the link register with:

```
MOVS      pc, lr
```

returns control from the handler.

The handler entry and exit code to stack the return address and pop it on return is:

```
STMFD     sp!,{reglist,lr}
;...
LDMFD     sp!,{reglist,pc}^
```

For exceptions that occur in Thumb state, the handler return instruction (`MOVS pc, lr`) changes the PC to the address of the next instruction to execute. This is at $(PC-2)$, so the value stored by the processor in *lr_mode* is $(PC-2)$.

Returning from FIQ and IRQ handlers

After executing each instruction, the processor checks to see whether the interrupt pins are LOW and whether the interrupt disable bits in the CPSR are clear. As a result, IRQ or FIQ exceptions are generated only after the PC has been updated. The processor stores (PC-4) in *lr_mode*. This makes *lr_mode* point one instruction beyond the end of the instruction in which the exception occurred. When the handler has finished, execution must continue from the instruction prior to the one pointed to by *lr_mode*. The address to continue from is one word (four bytes) less than that in *lr_mode*, so the return instruction is:

```
SUBS      pc, lr, #4
```

The handler entry and exit code to stack the return address and pop it on return is:

```
SUB      lr,lr,#4
STMFD    sp!,{reglist,lr}
;...
LDMFD    sp!,{reglist,pc}^
```

For exceptions that occur in Thumb state, the handler return instruction (SUBS pc,lr,#4) changes the PC to the address of the next instruction to execute. Because the PC is updated before the exception is taken, the next instruction is at (PC-4). The value stored by the processor in *lr_mode* is therefore PC.

Returning from Prefetch Abort handlers

If the processor attempts to fetch an instruction from an illegal address, the instruction is flagged as invalid. Instructions already in the pipeline continue to execute until the invalid instruction is reached, at which point a Prefetch Abort is generated.

The exception handler loads the unmapped instruction into physical memory and uses the MMU, if there is one, to map the virtual memory location into the physical one. The handler must then return to retry the instruction that caused the exception. The instruction now loads and executes.

Because the PC is not updated at the time the prefetch abort is issued, *lr_ABT* points to the instruction following the one that caused the exception. The handler must return to *lr_ABT-4* with:

```
SUBS      pc,lr, #4
```

The handler entry and exit code to stack the return address and pop it on return is:

```
SUB      lr,lr,#4
STMFD    sp!,{reglist,lr}
;...
LDMFD    sp!,{reglist,pc}^
```

For exceptions that occur in Thumb state, the handler return instruction (`SUBS pc, lr, #4`) changes the PC to the address of the aborted instruction. Because the PC is not updated before the exception is taken, the aborted instruction is at (PC-4). The value stored by the processor in `lr_mode` is therefore PC.

Returning from Data Abort handlers

When a load or store instruction tries to access memory, the PC has been updated. The stored value of (PC-4) in `lr_ABT` points to the second instruction beyond the address where the exception occurred. When the MMU, if present, has mapped the appropriate address into physical memory, the handler must return to the original, aborted instruction so that a second attempt can be made to execute it. The return address is therefore two words (eight bytes) less than that in `lr_ABT`, making the return instruction:

```
SUBS    pc, lr, #8
```

The handler entry and exit code to stack the return address and pop it on return is:

```
SUB      lr, lr, #8
STMFD    sp!, {reglist, lr}
;...
LDMFD    sp!, {reglist, pc}^
```

For exceptions that occur in Thumb state, the handler return instruction (`SUBS pc, lr, #8`) changes the PC to the address of the aborted instruction. Because the PC is updated before the exception is taken, the aborted instruction is at (PC-6). The value stored by the processor in `lr_mode` is therefore (PC+2).

6.4 Handling an exception

Your top-level veneer routine must save the processor status and any required registers on the stack. You then have the following options for writing the exception handler:

- Write the whole exception handler in ARM code.
- Perform a BX (Branch and eXchange) to a Thumb code routine that handles the exception. The routine must return to an ARM code veneer in order to return from the exception, because the Thumb instruction set does not have the instructions required to restore CPSR from SPSR.

Figure 6-2 shows how to implement this strategy.




Figure 6-2 Handling an exception in ARM or Thumb state

See Chapter 4 *Interworking ARM and Thumb* for details of how to combine ARM and Thumb code in this way.

6.5 Installing an exception handler

Any new exception handler must be installed in the vector table. When installation is complete, the new handler executes whenever the corresponding exception occurs.

This section includes:

- *Methods of installing exception handlers*
- *Installing the handlers at reset*
- *Installing the handlers from C* on page 6-16.

6.5.1 Methods of installing exception handlers

Exception handlers can be installed in the following ways:

Branch instruction

This is the simplest way to reach the exception handler. Each entry in the vector table contains a branch to the required handler routine. However, this method does have a limitation. Because the branch instruction only has a range of 32MB relative to the PC, with some memory organizations the branch might be unable to reach the handler.

Load PC instruction

With this method, the PC is forced directly to the handler address by:

1. Storing the absolute address of the handler in a suitable memory location (within 4KB of the vector address).
2. Placing an instruction in the vector that loads the PC with the contents of the chosen memory location.

6.5.2 Installing the handlers at reset

If your application does not rely on the debugger or debug monitor to start program execution, you can load the vector table directly from your assembly language reset (or startup) code.

If your ROM is at location 0x0 in memory, you can have a branch statement for each vector at the start of your code. This could also include the FIQ handler if it is running directly from 0x1C (see *Interrupt handlers* on page 6-29).

Example 6-2 on page 6-15 shows code that sets up the vectors if they are located in ROM at address zero. You can substitute branch statements for the loads.

Example 6-2

```

Vector_Init_Block
    LDR    pc, Reset_Addr
    LDR    pc, Undefined_Addr
    LDR    pc, SWI_Addr
    LDR    pc, Prefetch_Addr
    LDR    pc, Abort_Addr
    NOP                                ;Reserved vector
    LDR    pc, IRQ_Addr
    LDR    pc, FIQ_Addr

Reset_Addr    DCD    Start_Boot
Undefined_Addr DCD    Undefined_Handler
SWI_Addr      DCD    SWI_Handler
Prefetch_Addr DCD    Prefetch_Handler
Abort_Addr    DCD    Abort_Handler
              DCD    0                                ;Reserved vector
IRQ_Addr      DCD    IRQ_Handler
FIQ_Addr      DCD    FIQ_Handler

```

You must have ROM at location 0x0 on reset. Your reset code can remap RAM to location 0x0. Before doing this, it must copy the vectors (and the FIQ handler if required) down from an area in ROM into the RAM.

In this case, you must use an LDR pc instruction to address the reset handler, so that the reset vector code can be position independent.

Example 6-3 copies down the vectors given in Example 6-2 to the vector table in RAM.

Example 6-3

```

MOV    r8, #0
ADR    r9, Vector_Init_Block
LDMIA  r9!, {r0-r7}                ;Copy the vectors (8 words)
STMIA  r8!, {r0-r7}
LDMIA  r9!, {r0-r7}                ;Copy the DCD'ed addresses
STMIA  r8!, {r0-r7}                ;(8 words again)

```

Alternatively, you can use the scatter-loading mechanism to define the load and execution address of the vector table. In that case, the C library copies the vector table for you (see Chapter 2 *Embedded Software Development*).

6.5.3 Installing the handlers from C

Sometimes during development work it is necessary to install exception handlers into the vectors directly from the main application. As a result, the required instruction encoding must be written to the appropriate vector address. This can be done for both the branch and the load PC method of reaching the handler.

Branch method

The required instruction can be constructed as follows:

1. Take the address of the exception handler.
2. Subtract the address of the corresponding vector.
3. Subtract 0x8 to provide for prefetching.
4. Shift the result to the right by two to give a word offset, rather than a byte offset.
5. Test that the top eight bits of this are clear, to ensure that the result is only 24 bits long (because the offset for the branch is limited to this).
6. Logically OR this with 0xEA000000 (the opcode for the Branch instruction) to produce the value to be placed in the vector.

Example 6-4 on page 6-17 shows a C function that implements this algorithm.

It takes the following arguments:

- the address of the handler
- the address of the vector in which the handler is to be installed.

The function can install the handler and return the original contents of the vector. This result can be used to create a chain of handlers for a particular exception. See *Chaining exception handlers* on page 6-44 for more details.

The following code calls this to install an IRQ handler:

```
unsigned *irqvec = (unsigned *)0x18;
Install_Handler ((unsigned)IRQHandler, irqvec);
```

In this case, the returned, original contents of the IRQ vector are discarded.

Example 6-4 Implementing the branch method

```

unsigned Install_Handler (unsigned routine, unsigned *vector)
/* Updates contents of 'vector' to contain branch instruction */
/* to reach 'routine' from 'vector'. Function return value is */
/* original contents of 'vector'.*/
/* NB: 'Routine' must be within range of 32MB from 'vector'.*/

{
    unsigned vec, oldvec;
    vec = ((routine - (unsigned)vector - 0x8)>>2);
    if ((vec & 0xFF000000))
    {
        /* diagnose the fault */
        printf ("Installation of Handler failed");
        exit (1);
    }
    vec = 0xEA000000 | vec;
    oldvec = *vector;
    *vector = vec;
    return (oldvec);
}

```

Load PC method

The required instruction can be constructed as follows:

1. Take the address of the word containing the address of the exception handler.
2. Subtract the address of the corresponding vector.
3. Subtract 0x8 to provide for prefetching.
4. Check that the result can be represented in 12 bits.
5. Logically OR this with 0xe59FF000 (the opcode for LDR pc, [pc,#offset]) to produce the value to be placed in the vector.
6. Put the address of the handler into the storage location.

Example 6-5 on page 6-18 shows a C routine that implements this method.

Again in this example the returned, original contents of the IRQ vector are discarded, but they could be used to create a chain of handlers. See *Chaining exception handlers* on page 6-44 for more information.

Example 6-5 Implementing the load PC method

```

unsigned Install_Handler (unsigned location, unsigned *vector)

/* Updates contents of 'vector' to contain LDR pc, [pc, #offset] */
/* instruction to cause long branch to address in 'location'. */
/* Function return value is original contents of 'vector'. */

{   unsigned vec, oldvec;
    vec = ((unsigned)location - (unsigned)vector - 0x8) | 0xe59ff000;
    oldvec = *vector;
    *vector = vec;
    return (oldvec);
}

```

The following code calls this to install an IRQ handler:

```

unsigned *irqvec = (unsigned *)0x18;
static unsigned pIRQ_Handler = (unsigned)IRQ_handler
Install_Handler (&pIRQHandler, irqvec);

```

Note

If you are using a processor with separate instruction and data caches, such as StrongARM®, or ARM940T, you must ensure that cache coherence problems do not prevent the new contents of the vectors from being used.

The data cache (or at least the entries containing the modified vectors) must be cleaned to ensure the new vector contents are written to main memory. You must then flush the instruction cache to ensure that the new vector contents are read from main memory.

For details of cache clean and flush operations, see the Technical Reference Manual for your target processor.

6.6 SWI handlers

This section describes SWI handlers, and includes:

- *Determining the SWI to be called*
- *SWI handlers in assembly language on page 6-20*
- *SWI handlers in C and assembly language on page 6-21*
- *Using SWIs in Supervisor mode on page 6-22*
- *Calling SWIs from an application on page 6-24*
- *Calling SWIs dynamically from an application on page 6-26.*

6.6.1 Determining the SWI to be called

When the SWI handler is entered, it must establish which SWI is being called. This information can be stored in bits 0-23 of the instruction itself, as shown in Figure 6-3, or passed in an integer register, usually one of `r0-r3`.




Figure 6-3 ARM SWI instruction

The top-level SWI handler can load the SWI instruction relative to the link register. Do this in assembly language, C/C++ inline, or embedded assembler.

The handler must first load the SWI instruction that caused the exception into a register. At this point, `lr_SVC` holds the address of the instruction that follows the SWI instruction, so the SWI is loaded into the register (in this case `r0`) using:

```
LDR r0, [r, #-4]
```

The handler can then examine the comment field bits, to determine the required operation. The SWI number is extracted by clearing the top eight bits of the opcode:

```
BIC r0, r0, #0xFF000000
```

Example 6-6 on page 6-20 shows how you can put these instructions together to form a top-level SWI handler.

See *Determining the processor state* on page 6-6 for an example of a handler that deals with SWI instructions in both ARM state and Thumb state.

Example 6-6 Top-level SWI handler

```

PRESERVE8

AREA TopLevelSwi, CODE, READONLY ; Name this block of code.
EXPORT SWI_Handler
SWI_Handler
    STMFD    sp!,{r0-r12,lr}      ; Store registers.
    LDR      r0,[lr,#-4]          ; Calculate address of SWI instruction
                                      ; and load it into r0.
    BIC      r0,r0,#0xff000000    ; Mask off top 8 bits of instruction
                                      ; to give SWI number.

    ;
    ; Use value in r0 to determine which SWI routine to execute.
    ;
    LDMFD    sp!, {r0-r12,pc}^    ; Restore registers and return.
    END                                ; Mark end of this file.

```

6.6.2 SWI handlers in assembly language

The easiest way to call the handler for the requested SWI number is to use a jump table. If `r0` contains the SWI number, the code in Example 6-7 can be inserted into the top-level handler given in Example 6-6, following on from the `BIC` instruction.

Example 6-7 SWI jump table

```

    CMP      r0,#MaxSWI          ; Range check
    LDRLS    pc, [pc,r0,LSL #2]
    B        SWIOutOfRange
SWIJumpTable
    DCD      SWInum0
    DCD      SWInum1
                                      ; DCD for each of other SWI routines
SWInum0
    B        EndofSWI            ; SWI number 0 code
SWInum1
    B        EndofSWI            ; SWI number 1 code
                                      ; Rest of SWI handling code
    ;
EndofSWI
                                      ; Return execution to top level
                                      ; SWI handler so as to restore
                                      ; registers and return to program.

```

6.6.3 SWI handlers in C and assembly language

Although the top-level handler must always be written in ARM assembly language, the routines that handle each SWI can be written in either assembly language or in C. See *Using SWIs in Supervisor mode* on page 6-22 for a description of restrictions.

The top-level handler uses a BL (Branch with Link) instruction to jump to the appropriate C function. Because the SWI number is loaded into r0 by the assembly routine, this is passed to the C function as the first parameter (in accordance with the AAPCS as described in Appendix A *Using the Procedure Call Standard*). The function can use this value in, for example, a switch() statement.

You can add the following line to the SWI_Handler routine in Example 6-6 on page 6-20:

```
BL    C_SWI_Handler    ; Call C routine to handle the SWI
```

Example 6-8 shows how to implement the C function.

Example 6-8

```
void C_SWI_handler (unsigned number)
{
    switch (number)
    {
        case 0 :                /* SWI number 0 code */
            break;
        case 1 :                /* SWI number 1 code */
            break;
        ...
        default :               /* Unknown SWI - report error */
    }
}
```

The supervisor stack space might be limited, so avoid using functions that require a large amount of stack space.

```
MOV    r1, sp        ; Second parameter to C routine...
                        ; ...is pointer to register values.
BL     C_SWI_Handler ; Call C routine to handle the SWI
```

You can pass values in and out of a SWI handler written in C, provided that the top-level handler passes the stack pointer value into the C function as the second parameter (in r1):

and the C function is updated to access it:

```
void C_SWI_handler(unsigned number, unsigned *reg)
```

The C function can now access the values contained in the registers at the time the SWI instruction was encountered in the main application code (see Figure 6-4). It can read from them:

```
value_in_reg_0 = reg [0];
value_in_reg_1 = reg [1];
value_in_reg_2 = reg [2];
value_in_reg_3 = reg [3];
```

and also write back to them:

```
reg [0] = updated_value_0;
reg [1] = updated_value_1;
reg [2] = updated_value_2;
reg [3] = updated_value_3;
```

This causes the updated value to be written into the appropriate stack position, and then restored into the register by the top-level handler.




Figure 6-4 Accessing the supervisor stack

6.6.4 Using SWIs in Supervisor mode

When a SWI instruction is executed:

1. The processor enters Supervisor mode.
2. The CPSR is stored into spsr_SVC.
3. The return address is stored in lr_SVC (see *The processor response to an exception* on page 6-8).

If the processor is already in Supervisor mode, lr_SVC and spsr_SVC are corrupted.

If you call a SWI while in Supervisor mode you must store `lr_SVC` and `spsr_SVC` to ensure that the original values of the link register and the SPSR are not lost. For example, if the handler routine for a particular SWI number calls another SWI, you must ensure that the handler routine stores both `lr_SVC` and `spsr_SVC` on the stack. This guarantees that each invocation of the handler saves the information required to return to the instruction following the SWI that invoked it. Example 6-9 shows how to do this.

Example 6-9 SWI Handler

```

        AREA SWI_Area, CODE, READONLY

        PRESERVE8

        EXPORT SWI_Handler
        IMPORT C_SWI_Handler

T_bit EQU 0x20

SWI_Handler

        STMFD    sp!,{r0-r3,r12,lr}    ; Store registers.
        MOV     r1, sp                  ; Set pointer to parameters.
        MRS     r0, spsr                ; Get spsr.
        STMFD    sp!, {r0, r3}          ; Store spsr onto stack and another register to maintain
                                          ; 8-byte-aligned stack. This is only really needed in case of
                                          ; nested SWIs.

        ; the next two instructions only work for SWI calls from ARM state.
        ; See Example 6-18 on page 6-36 for a version that works for calls from either ARM or Thumb.

        LDR     r0,[lr,#-4]             ; Calculate address of SWI instruction and load it into r0.
        BIC     r0,r0,#0xFF000000      ; Mask off top 8 bits of instruction to give SWI number.

        ; r0 now contains SWI number
        ; r1 now contains pointer to stacked registers

        BL      C_SWI_Handler          ; Call C routine to handle the SWI.
        LDMFD    sp!, {r0, r3}          ; Get spsr from stack.
        MSR     spsr_cf, r0             ; Restore spsr.
        LDMFD    sp!, {r0-r3,r12,pc}^   ; Restore registers and return.

        END

```

Nested SWIs in C and C++

You can write nested SWIs in C or C++. Code generated by the ARM compiler stores and reloads `1r_SVC` as necessary.

6.6.5 Calling SWIs from an application

You can call a SWI from assembly language or C/C++.

In assembly language, set up any required register values and issue the relevant SWI. For example:

```
MOV    r0, #65    ; load r0 with the value 65
SWI     0x0        ; Call SWI 0x0 with parameter value in r0
```

The SWI instruction can be conditionally executed, as can almost all ARM instructions.

From C/C++, declare the SWI as an `__swi` function, and call it. For example:

```
__swi(0) void my_swi(int);
.
.
.
my_swi(65);
```

This enables a SWI to be compiled inline, without additional calling overhead, provided that:

- any arguments are passed in `r0-r3` only
- any results are returned in `r0-r3` only.

The parameters are passed to the SWI as if the SWI were a real function call. However, if there are between two and four return values, you must tell the compiler that the return values are being returned in a structure, and use the `__value_in_regs` directive. This is because a **struct**-valued function is usually treated as if it were a **void** function whose first argument is the address where the result structure must be placed.

Example 6-10 on page 6-25 and Example 6-11 on page 6-25 show a SWI handler that provides SWI numbers `0x0`, `0x1`, `0x2` and `0x3`. SWIs `0x0` and `0x1` each take two integer parameters and return a single result. SWI `0x2` takes four parameters and returns a single result. SWI `0x3` takes four parameters and returns four results. This example is in the main examples directory, in `...\swi\main.c` and `...\swi\swi.h`.

Example 6-10 main.c

```

#include <stdio.h>
#include "swi.h"

unsigned *swi_vec = (unsigned *)0x08;
extern void SWI_Handler(void);

int main( void )
{
    int result1, result2;
    struct four_results res_3;
    Install_Handler( (unsigned) SWI_Handler, swi_vec );
    printf("result1 = multiply_two(2,4) = %d\n", result1 = multiply_two(2,4));
    printf("result2 = multiply_two(3,6) = %d\n", result2 = multiply_two(3,6));
    printf("add_two( result1, result2 ) = %d\n", add_two( result1, result2 ));
    printf("add_multiply_two(2,4,3,6) = %d\n", add_multiply_two(2,4,3,6));
    res_3 = many_operations( 12, 4, 3, 1 );
    printf("res_3.a = %d\n", res_3.a );
    printf("res_3.b = %d\n", res_3.b );
    printf("res_3.c = %d\n", res_3.c );
    printf("res_3.d = %d\n", res_3.d );
    return 0;
}

```

Example 6-11 swi.h

```

__swi(0) int multiply_two(int, int);
__swi(1) int add_two(int, int);
__swi(2) int add_multiply_two(int, int, int, int);

struct four_results
{
    int a;
    int b;
    int c;
    int d;
};

__swi(3) __value_in_regs struct four_results
many_operations(int, int, int, int);

```

6.6.6 Calling SWIs dynamically from an application

In some circumstances it might be necessary to call a SWI whose number is not known until runtime. This situation might occur, for example, when there are a number of related operations that can be performed on an object, and each operation has its own SWI. In this case, the methods described in the previous sections are not appropriate.

There are several ways of dealing with this, for example:

- Construct the SWI instruction from the SWI number, store it somewhere, then execute it.
- Use a generic SWI that takes, as an extra argument, a code for the actual operation to be performed on its arguments. The generic SWI decodes the operation and performs it.

The second mechanism can be implemented in assembly language by passing the required operation number in a register, typically `r0` or `r12`. You can then rewrite the SWI handler to act on the value in the appropriate register.

Because some value has to be passed to the SWI in the comment field, it is possible for a combination of these two methods to be used.

For example, an operating system might make use of only a single SWI instruction and employ a register to pass the number of the required operation. This leaves the rest of the SWI space available for application-specific SWIs. You can use this method if the overhead of extracting the SWI number from the instruction is too great in a particular application. This is how the ARM (0x123456) and Thumb (0xAB) semihosted SWIs are implemented.

Example 6-12 shows how `__swi` can be used to map a C function call onto a semihosting SWI. It is derived from `retarget.c` in the main examples directory, in `...\emb_sw_dev\source\retarget.c`.

Example 6-12 Mapping a C function onto a semihosting SWI

```
#ifdef __thumb
/* Thumb Semihosting SWI */
#define SemiSWI 0xAB
#else
/* ARM Semihosting SWI */
#define SemiSWI 0x123456
#endif

/* Semihosting SWI to write a character */
__swi(SemiSWI) void Semihosting(unsigned op, char *c);
```

```
#define WriteC(c) Semihosting (0x3,c)

void write_a_character(int ch)
{
    char tempch = ch;
    WriteC( &tempch );
}
```

The compiler includes a mechanism to support the use of r12 to pass the value of the required operation. Under the AAPCS, r12 is the ip register and has a dedicated role only during function calls. At other times, you can use it as a scratch register. The arguments to the generic SWI are passed in registers r0-r3 and values are optionally returned in r0-r3 as described earlier (see *Calling SWIs from an application* on page 6-24). The operation number passed in r12 can be the number of the SWI to be called by the generic SWI. However, this is not required.

Example 6-13 shows a C fragment that uses a generic, or *indirect* SWI.

Example 6-13

```
__swi_indirect(0x80)
    unsigned SWI_ManipulateObject(unsigned operationNumber,
                                   unsigned object,unsigned parameter);

unsigned DoSelectedManipulation(unsigned object,
                                unsigned parameter, unsigned operation)
{ return SWI_ManipulateObject(operation, object, parameter);
}
```

This produces the following code:

```
DoSelectedManipulation PROC
    STMFD    sp!,{r3,lr}
    MOV      r12,r2
    SWI      0x80
    LDMFD    sp!,{r3,pc}
    ENDP
```

It is also possible to pass the SWI number in r0 from C using the __swi mechanism. For example, if SWI 0x0 is used as the generic SWI and operation 0 is a character read and operation 1 a character write, you can set up the following:

```
__swi (0) char __ReadCharacter (unsigned op);
__swi (0) void __WriteCharacter (unsigned op, char c);
```

These can be used in a more reader-friendly way by defining the following:

```
#define ReadCharacter () __ReadCharacter (0);  
#define WriteCharacter (c) __WriteCharacter (1, c);
```

However, if you use `r0` in this way, then only three registers are available for passing parameters to the SWI. Usually, if you have to pass more parameters to a subroutine in addition to `r0-r3`, you can do this using the stack. However, stacked parameters are not easily accessible to a SWI handler, because they typically exist on the User mode stack rather than the supervisor stack employed by the SWI handler.

Alternatively, one of the registers (typically `r1`) can be used to point to a block of memory storing the other parameters.

6.7 Interrupt handlers

This section describes how to write interrupt handlers to service the external interrupts FIQ and IRQ, and includes:

- *Levels of external interrupt*
- *Simple interrupt handlers in C*
- *Reentrant interrupt handlers* on page 6-31
- *Example interrupt handlers in assembly language* on page 6-33.

6.7.1 Levels of external interrupt

The ARM processor has two levels of external interrupt, FIQ and IRQ, both of which are level-sensitive active LOW signals into the core. For an interrupt to be taken, the appropriate disable bit in the CPSR must be clear.

FIQs have higher priority than IRQs in the following ways:

- FIQs are serviced first when multiple interrupts occur.
- Servicing an FIQ causes IRQs to be disabled, preventing them from being serviced until after the FIQ handler has re-enabled them. This is usually done by restoring the CPSR from the SPSR at the end of the handler.

The FIQ vector is the last entry in the vector table (at address 0x1C) so that the FIQ handler can be placed directly at the vector location and run sequentially from that address. This removes the requirement for a branch and its associated delays, and also means that if the system has a cache, the vector table and FIQ handler might all be locked down in one block within it. This is important because FIQs are designed to service interrupts as quickly as possible. The five extra FIQ mode banked registers enable status to be held between calls to the handler, again increasing execution speed.

———— Note ————

An interrupt handler must contain code to clear the source of the interrupt.

6.7.2 Simple interrupt handlers in C

You can write simple C interrupt handlers by using the `__irq` function declaration keyword. You can use the `__irq` keyword both for simple one-level interrupt handlers, and interrupt handlers that call subroutines. However, you cannot use the `__irq` keyword for *reentrant* interrupt handlers, because it does not cause the SPSR to be saved or restored. In this context, reentrant means that the handler re-enables interrupts, and can itself be interrupted. See *Reentrant interrupt handlers* on page 6-31 for more information.

The `__irq` keyword:

- preserves all AAPCS corruptible registers
- preserves all other registers (excluding the floating-point registers) used by the function
- exits the function by setting the PC to (1r-4) and restoring the CPSR to its original value.

If the function calls a subroutine, `__irq` preserves the link register for the interrupt mode in addition to preserving the other corruptible registers. See *Calling subroutines from interrupt handlers* for more information.

Note

C interrupt handlers cannot be produced in this way when compiling Thumb C code. When compiling for Thumb (`--thumb` option or `#pragma thumb`), any functions specified as `__irq` are compiled for ARM.

However, the subroutine called by an `__irq` function can be compiled for Thumb, with interworking enabled. See Chapter 4 *Interworking ARM and Thumb* for more information on interworking.

Calling subroutines from interrupt handlers

If you call subroutines from your top-level interrupt handler, the `__irq` keyword also restores the value of `lr_IRQ` from the stack so that it can be used by a `SUBS` instruction to return to the correct address after the interrupt has been handled.

Example 6-14 shows how this works. The top level interrupt handler reads the value of a memory-mapped interrupt controller base address at `0x80000000`. If the value of the address is 1, the top-level handler branches to a handler written in C.

Example 6-14

```
__irq void IRQHandler (void)
{
    volatile unsigned int *base = (unsigned int *) 0x80000000;

    if (*base == 1)          // which interrupt was it?
    {
        C_int_handler();    // process the interrupt
    }
    *(base+1) = 0;          // clear the interrupt
}
```

Compiled with armcc, Example 6-14 on page 6-30 produces the following code:

```
IRQHandler PROC
    STMFD    sp!,{r0-r4,r12,lr}
    MOV      r4,#0x80000000
    LDR      r0,[r4,#0]
    SUB      sp,sp,#4
    CMP      r0,#1
    BLEQ     C_int_handler
    MOV      r0,#0
    STR      r0,[r4,#4]
    ADD      sp,sp,#4
    LDMFD    sp!,{r0-r4,r12,lr}
    SUBS     pc,lr,#4
ENDP
```

Compare this with the result when the `__irq` keyword is not used:

```
IRQHandler PROC
    STMFD    sp!,{r4,lr}
    MOV      r4,#0x80000000
    LDR      r0,[r4,#0]
    CMP      r0,#1
    BLEQ     C_int_handler
    MOV      r0,#0
    STR      r0,[r4,#4]
    LDMFD    sp!,{r4,pc}
ENDP
```

6.7.3 Reentrant interrupt handlers

If an interrupt handler re-enables interrupts, then calls a subroutine, and another interrupt occurs, the return address of the subroutine (stored in `lr_IRQ`) is corrupted when the second IRQ is taken. Using the `__irq` keyword in C does not cause the SPSR to be saved and restored, as required by reentrant interrupt handlers, so you must write your top level interrupt handler in assembly language.

A reentrant interrupt handler must save the IRQ state, switch processor modes, and save the state for the new processor mode before branching to a nested subroutine or C function.

In ARMv4 or later you can switch to System mode. System mode uses the User mode registers, and enables privileged access that might be required by your exception handler. See *System mode* on page 6-46 for more information. In ARM architectures prior to ARMv4 you must switch to Supervisor mode instead.

Note

This method works for both IRQ and FIQ interrupts. However, because FIQ interrupts are meant to be serviced as quickly as possible there is normally only one interrupt source, so it might not be necessary to provide for reentrancy.

The steps required to safely re-enable interrupts in an IRQ handler are:

1. Construct the return address and save it on the IRQ stack.
2. Save the work registers and `spsr_IRQ`.
3. Clear the source of the interrupt.
4. Switch to System mode and re-enable interrupts.
5. Save User mode link register and non callee-saved registers.
6. Call the C interrupt handler function.
7. When the C interrupt handler returns, restore User mode registers and disable interrupts.
8. Switch to IRQ mode, disabling interrupts.
9. Restore work registers and `spsr_IRQ`.
10. Return from the IRQ.

Example 6-15 shows how this works for System mode. Registers `r12` and `r14` are used as temporary work registers after `lr_IRQ` is pushed on the stack.

Example 6-15

```

PRESERVE8

AREA INTERRUPT, CODE, READONLY
IMPORT C_irq_handler
IRQ
SUB    lr, lr, #4          ; construct the return address
STMFD  sp!, {lr}          ; and push the adjusted lr_IRQ
MRS    r14, SPSR           ; copy spsr_IRQ to r14
STMFD  sp!, {r12, r14}    ; save work regs and spsr_IRQ

; Add instructions to clear the interrupt here
; then re-enable interrupts.
```



```

MSR      CPSR_c, #0x1F      ; switch to SYS mode, FIQ and IRQ
                                ; enabled. USR mode registers
                                ; are now current.
STMFD    sp!, {r0-r3, lr}   ; save lr_USR and non-callee
                                ; saved registers
BL       C_irq_handler      ; branch to C IRQ handler.
LDMFD    sp!, {r0-r3, lr}   ; restore registers
MSR      CPSR_c, #0x92      ; switch to IRQ mode and disable
                                ; IRQs. FIQ is still enabled.

LDMFD    sp!, {r12, r14}    ; restore work regs and spsr_IRQ
MSR      SPSR_cf, r14
LDMFD    sp!, {pc}^         ; return from IRQ.
END

```

This example assumes that FIQ remains permanently enabled.

6.7.4 Example interrupt handlers in assembly language

Interrupt handlers are often written in assembly language to ensure that they execute quickly. The following sections give some examples:

- *Single-channel DMA transfer*
- *Dual-channel DMA transfer* on page 6-34
- *Interrupt prioritization* on page 6-35
- *Context switch* on page 6-37.

Single-channel DMA transfer

Example 6-16 on page 6-34 shows an interrupt handler that performs interrupt driven I/O to memory transfers (soft DMA). The code is an FIQ handler. It uses the banked FIQ registers to maintain state between interrupts. This code is best situated at location 0x1C.

In the example code:

- r8** Points to the base address of the I/O device that data is read from.
- IOData** Is the offset from the base address to the 32-bit data register that is read. Reading this register clears the interrupt.
- r9** Points to the memory location to where that data is being transferred.
- r10** Points to the last address to transfer to.

The entire sequence for handling a normal transfer is four instructions. Code situated after the conditional return is used to signal that the transfer is complete.

Example 6-16 FIQ handler

```

LDR    r11, [r8, #IOData]    ; Load port data from the I/O device.
STR    r11, [r9], #4         ; Store it to memory: update the pointer.
CMP    r9, r10               ; Reached the end ?
SUBLSS pc, lr, #4            ; No, so return.
                                ; Insert transfer complete
                                ; code here.

```

Byte transfers can be made by replacing the load instructions with load byte instructions. Transfers from memory to an I/O device are made by swapping the addressing modes between the load instruction and the store instruction.

Dual-channel DMA transfer

Example 6-17 on page 6-35 is similar to Example 6-16, except that there are two channels being handled. The code is an FIQ handler. It uses the banked FIQ registers to maintain state between interrupts. It is best situated at location 0x1c.

In the example code:

r8	Points to the base address of the I/O device from which data is read.
IOWStat	Is the offset from the base address to a register indicating which of two ports caused the interrupt.
IOWPort1Active	Is a bit mask indicating if the first port caused the interrupt (otherwise it is assumed that the second port caused the interrupt).
IOWPort1, IOWPort2	Are offsets to the two data registers to be read. Reading a data register clears the interrupt for the corresponding port.
r9	Points to the memory location to which data from the first port is being transferred.
r10	Points to the memory location to which data from the second port is being transferred.
r11, r12	Point to the last address to transfer to (r11 for the first port, r12 for the second).

The entire sequence to handle a normal transfer takes nine instructions. Code situated after the conditional return is used to signal that the transfer is complete.

Example 6-17 FIQ handler

```

LDR    r13, [r8, #IOStat]      ; Load status register to find which port
                                   ; caused the interrupt.
TST     r13, #IOPort1Active
LDREQ  r13, [r8, #IOPort1]     ; Load port 1 data.
LDRNE  r13, [r8, #IOPort2]     ; Load port 2 data.
STREQ  r13, [r9], #4           ; Store to buffer 1.
STRNE  r13, [r10], #4          ; Store to buffer 2.
CMP     r9, r11                ; Reached the end?
CMPLT  r10, r12                ; On either channel?
SUBNES  pc, lr, #4             ; Return
                                   ; Insert transfer complete code here.

```

Byte transfers can be made by replacing the load instructions with load byte instructions. Transfers from memory to an I/O device are made by swapping the addressing modes between the conditional load instructions and the conditional store instructions.

Interrupt prioritization

Example 6-18 on page 6-36 dispatches up to 32 interrupt sources to their appropriate handler routines. Because it is designed for use with the normal interrupt vector (IRQ), it is branched to from location 0x18.

External hardware is used to prioritize the interrupt and present the high-priority active interrupt in an I/O register.

In the example code:

- IntBase** Holds the base address of the interrupt controller.
- IntLevel** Holds the offset of the register containing the highest-priority active interrupt.
- r13** Is assumed to point to a small full descending stack.

Interrupts are enabled after ten instructions, including the branch to this code.

The specific handler for each interrupt is entered after a further two instructions (with all registers preserved on the stack).

In addition, the last three instructions of each handler are executed with interrupts turned off again, so that the SPSR can be safely recovered from the stack.

Note

Application Note 30: *Software Prioritization of Interrupts* describes multiple-source prioritization of interrupts using software, as opposed to using hardware as described here.

Example 6-18

```

; first save the critical state
SUB    lr, lr, #4                ; Adjust the return address
                                           ; before we save it.
STMFD  sp!, {lr}                ; Stack return address
MRS    r14, SPSR                ; get the SPSR ...
STMFD  sp!, {r12, r14}          ; ... and stack that plus a
                                           ; working register too.
                                           ; Now get the priority level of the
                                           ; highest priority active interrupt.
MOV     r12, #IntBase            ; Get the interrupt controller's
                                           ; base address.
LDR     r12, [r12, #IntLevel]    ; Get the interrupt level (0 to 31).

; Now read-modify-write the CPSR to enable interrupts.

MRS     r14, CPSR                ; Read the status register.
BIC     r14, r14, #0x80          ; Clear the I bit
                                           ; (use 0x40 for the F bit).
MSR     CPSR_c, r14              ; Write it back to re-enable
                                           ; interrupts and
LDR     pc, [pc, r12, LSL #2]    ; jump to the correct handler.
                                           ; PC base address points to this
                                           ; instruction + 8
NOP                                           ; pad so the PC indexes this table.

                                           ; Table of handler start addresses
DCD     Priority0Handler
DCD     Priority1Handler
DCD     Priority2Handler
; ...
Priority0Handler
STMFD  sp!, {r0 - r11}          ; Save other working registers.
                                           ; Insert handler code here.
; ...
LDMFD  sp!, {r0 - r11}          ; Restore working registers (not r12).
```

```

; Now read-modify-write the CPSR to disable interrupts.
MRS    r12, CPSR          ; Read the status register.
ORR     r12, r12, #0x80    ; Set the I bit
                               ; (use 0x40 for the F bit).
MSR     CPSR_c, r12        ; Write it back to disable interrupts.

; Now that interrupt disabled, can safely restore SPSR then return.
LDMFD   sp!, {r12, r14}    ; Restore r12 and get SPSR.
MSR     SPSR_csfxf, r14    ; Restore status register from r14.
LDMFD   sp!, {pc}^         ; Return from handler.
Priority1Handler
; ...
```

Context switch

Example 6-19 on page 6-38 performs a context switch on the User mode process. The code is based around a list of pointers to *Process Control Blocks* (PCBs) of processes that are ready to run.

Figure 6-5 shows the layout of the PCBs that the example expects.

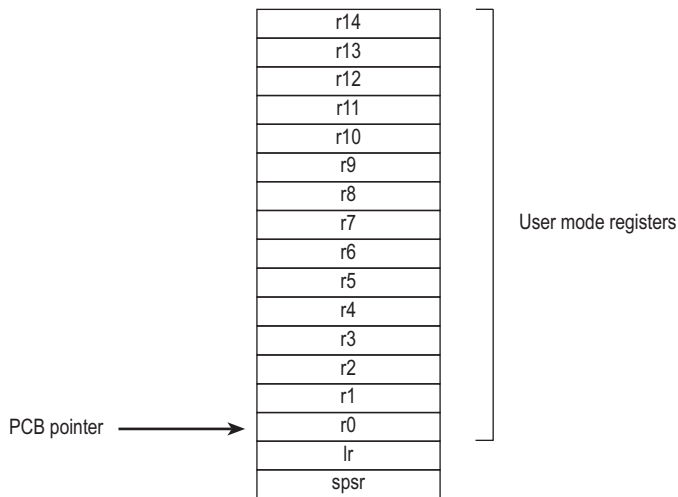


Figure 6-5 PCB layout

The pointer to the PCB of the next process to run is pointed to by r12, and the end of the list has a zero pointer. Register r13 is a pointer to the PCB, and is preserved between time slices, so that on entry it points to the PCB of the currently running process.

Example 6-19

```

STMIA    r13, {r0 - r14}^      ; Dump user registers above r13.
MRS      r0, SPSR               ; Pick up the user status
STMDB    r13, {r0, lr}         ; and dump with return address below.
LDR      r13, [r12], #4        ; Load next process info pointer.
CMP      r13, #0               ; If it is zero, it is invalid
LDMNEDB  r13, {r0, lr}         ; Pick up status and return address.
MSRNE    SPSR_cxsf, r0         ; Restore the status.
LDMNEIA  r13, {r0 - r14}^      ; Get the rest of the registers
NOP
SUBNES   pc, lr, #4             ; and return and restore CPSR.
                                   ; Insert "no next process code" here.

```

6.8 Reset handlers

The operations carried out by the Reset handler depend on the system for which the software is being developed. For example, it might:

- Set up exception vectors. See *Installing an exception handler* on page 6-14 for details.
- Initialize stacks and registers.
- Initialize the memory system, if using an MMU.
- Initialize any critical I/O devices.
- Enable interrupts.
- Change processor mode and/or state.
- Initialize variables required by C and call the main application.

See Chapter 2 *Embedded Software Development* for more information.

6.9 Undefined Instruction handlers

Instructions that are not recognized by the processor are offered to any coprocessors attached to the system. If the instruction remains unrecognized, an Undefined Instruction exception is generated. It might be the case that the instruction is intended for a coprocessor, but that the relevant coprocessor, for example a Floating-Point Accelerator (FPA), is not attached to the system. However, a software emulator for such a coprocessor might be available.

Such an emulator must:

1. Attach itself to the Undefined Instruction vector and store the old contents.
2. Examine the Undefined Instruction to see if it has to be emulated. This is similar to the way in which a SWI handler extracts the number of a SWI, but rather than extracting the bottom 24 bits, the emulator must extract bits [27:24].

These bits determine whether the instruction is a coprocessor operation in the following way:

- If bits [27:24] = b1110 or b110x, the instruction is a coprocessor instruction.
 - If bits [8:11] show that this coprocessor emulator has to handle the instruction, the emulator must process the instruction and return to the user program.
3. Otherwise the emulator must pass the exception onto the original handler (or the next emulator in the chain) using the vector stored when the emulator was installed.

When a chain of emulators is exhausted, no further processing of the instruction can take place, so the Undefined Instruction handler must report an error and quit. See *Chaining exception handlers* on page 6-44 for more information.

Note

The Thumb instruction set does not have coprocessor instructions, so there is no requirement for the Undefined Instruction handler to emulate such instructions.

6.10 Prefetch Abort handler

If the system has no MMU, the Prefetch Abort handler can report the error and quit. Otherwise the address that caused the abort must be restored into physical memory. `lr_ABT` points to the instruction at the address following the one that caused the abort, so the address to be restored is at `lr_ABT-4`. The virtual memory fault for that address can be dealt with and the instruction fetch retried. The handler therefore returns to the same instruction rather than the following one, for example:

```
SUBS    pc, lr, #4
```

6.11 Data Abort handler

If there is no MMU, the Data Abort handler must report the error and quit. If there is an MMU, the handler must deal with the virtual memory fault.

The instruction that caused the abort is at `1r_ABT-8` because `1r_ABT` points two instructions beyond the instruction that caused the abort.

The following types of instruction can cause this abort:

Single Register Load or Store (LDR or STR)

The response depends on the processor type:

- If the abort takes place on an ARM6 processor:
 - If the processor is in early abort mode and writeback was requested, the address register has not been updated.
 - If the processor is in late abort mode and writeback was requested, the address register has been updated. The change must be undone.
- If the abort takes place on an ARM7 processor, including the ARM7TDMI®, the address register has been updated and the change must be undone.
- If the abort takes place on an ARM9, ARM10, or StrongARM processor, the address is restored by the processor to the value it had before the instruction started. No further action is required to undo the change.

Swap (SWP) There is no address register update involved with this instruction.

Load Multiple or Store Multiple (LDM or STM)

The response depends on the processor type:

- If the abort takes place on an ARM6 processor or ARM7 processor, and writeback is enabled, the base register is updated as if the whole transfer had taken place.
In the case of an LDM with the base register in the register list, the processor replaces the overwritten value with the modified base value so that recovery is possible. The original base address can then be recalculated using the number of registers involved.
- If the abort takes place on an ARM9, ARM10, or StrongARM processor and writeback is enabled, the base register is restored to the value it had before the instruction started.

In each of the three cases the MMU can load the required virtual memory into physical memory. The MMU *Fault Address Register* (FAR) contains the address that caused the abort. When this is done, the handler can return and try to execute the instruction again.

You can find example Data Abort handler code in the main examples directory, in `...\databort`.

6.12 Chaining exception handlers

In some situations there can be several different sources of a particular exception. For example:

- Angel uses an Undefined Instruction to implement breakpoints. However, Undefined Instruction exceptions also occur when a coprocessor instruction is executed, and no coprocessor is present.
- Angel uses a SWI for various purposes, such as entering Supervisor mode from User mode and supporting semihosting requests during development. However, a *Real Time Operating System* (RTOS) or an application might also implement some SWIs.

In such situations the following approaches can be taken to extend the exception handling code:

- *A single extended handler.*
- *Several chained handlers.*

6.12.1 A single extended handler

In some circumstances it is possible to extend the code in the exception handler to work out what the source of the exception was, and then directly call the appropriate code. In this case, you are modifying the source code for the exception handler.

Angel has been written to make this approach simple. Angel decodes SWIs and Undefined Instructions, and the Angel exception handlers can be extended to deal with non-Angel SWIs and Undefined Instructions.

However, this approach is only useful if all the sources of an exception are known when the single exception handler is written.

6.12.2 Several chained handlers

Some circumstances require more than a single handler. Consider the situation in which a standard Angel debugger is executing, and a standalone user application (or RTOS) is then downloaded that wants to support some additional SWIs. The newly loaded application might have its own entirely-independent exception handler that it wants to install, but which cannot replace the Angel handler.

In this case the address of the old handler must be noted so that the new handler is able to call the old handler if it discovers that the source of the exception is not a source it can deal with. For example, an RTOS SWI handler would call the Angel SWI handler on discovering that the SWI was not an RTOS SWI, so that the Angel SWI handler gets a chance to process it.

This approach can be extended to any number of levels to build a chain of handlers. Although code that takes this approach enables each handler to be entirely independent, it is less efficient than code that uses a single handler, or at least it becomes less efficient the further down the chain of handlers it has to go.

Both routines given in *Installing the handlers from C* on page 6-16 return the old contents of the vector. This value can be decoded to give:

The offset for a branch instruction

This can be used to calculate the location of the original handler and enable a new branch instruction to be constructed and stored at a suitable place in memory. If the replacement handler fails to handle the exception, it can branch to the constructed branch instruction, which in turn branches to the original handler.

The location used to store the address of the original handler

If the application handler fails to handle the exception, it has to load the PC from that location.

In most cases, such calculations are not necessary because information on the debug monitor or RTOS handlers is available to you. If so, the instructions required to chain in the next handler can be hard-coded into the application. The last section of the handler must check that the cause of the exception has been handled. If it has, the handler can return to the application. If not, it must call the next handler in the chain.

Note

When chaining in a handler before a debug monitor handler, you must remove the chain when the monitor is removed from the system, then directly install the application handler.

6.13 System mode

The ARM Architecture defines a User mode that has 15 general purpose registers, a PC, and a CPSR. In addition to this mode there are other privileged processor modes, each of which has an SPSR and a number of registers that replace some of the 15 User mode general purpose registers.

Note

This section only applies to processors that implement architectures ARMv4, ARMv4T, and later.

When a processor exception occurs, the current PC is copied into the link register for the exception mode, and the CPSR is copied into the SPSR for the exception mode. The CPSR is then altered in an exception-dependent way, and the PC is set to an exception-defined address to start the exception handler.

The ARM subroutine call instruction (BL) copies the return address into r14 before changing the PC, so the subroutine return instruction moves r14 to PC (MOV pc,lr).

Together these actions imply that ARM modes that handle exceptions must ensure that another exception of the same type cannot occur if they call subroutines, because the subroutine return address is overwritten with the exception return address.

In earlier versions of the ARM architecture, this problem has been solved by either carefully avoiding subroutine calls in exception code, or changing from the privileged mode to User mode. The first solution is often too restrictive, and the second means the task might not have the privileged access it requires to run correctly.

ARMv4 and later provide a processor mode called *System* mode, to overcome this problem. System mode is a privileged processor mode that shares the User mode registers. Privileged mode tasks can run in this mode, and exceptions no longer overwrite the link register.

Note

System mode cannot be entered by an exception. The exception handlers modify the CPSR to enter System mode. See *Reentrant interrupt handlers* on page 6-31 for an example.

Chapter 7

Debug Communications Channel

This chapter explains how to use the *Debug Communications Channel* (DCC). It contains the following sections:

- *About the Debug Communications Channel* on page 7-2
- *Target transfer of data* on page 7-3
- *Polled debug communications* on page 7-4
- *Interrupt-driven debug communications* on page 7-8
- *Access from Thumb state* on page 7-9.

7.1 About the Debug Communications Channel

The EmbeddedICE® logic in ARM® cores such as ARM7TDMI® and ARM9TDMI® contains a debug communications channel. This enables data to be passed between the target and the host debugger using the JTAG port and a protocol converter such as Multi-ICE®, without stopping the program flow or entering debug state. This chapter describes how the DCC can be accessed by a program running on the target, and by the host debugger.

7.1.1 Semihosting

You can use the debug communications channel for semihosting if you are using Multi-ICE with `$semihosting_enabled=2`. See the *Multi-ICE User Guide* for more information.

7.2 Target transfer of data

The target accesses the DCC as coprocessor 14 on the core using the ARM instructions MCR and MRC.

Two registers are provided to transfer data:

Comms data read register

A 32-bit wide register used to receive data from the debugger. The following instruction returns the read register value in Rd:

MRC p14, 0, Rd, c1, c0

Comms data write register

A 32-bit wide register used to send data to the debugger. The following instruction writes the value in Rn to the write register:

MCR p14, 0, Rn, c1, c0

Note

See the appropriate Technical Reference Manual for information on accessing DCC registers for the ARM10 and ARM11 cores. The instructions used, positions of the status bits, and interpretation of the status bits are different for processors later than ARM9.

7.3 Polled debug communications

In addition to the comms data read and write registers, a comms data control register is provided by the DCC. This section includes:

- *Comms data control register*
- *Target to debugger communication on page 7-5*
- *Debugger to target communication on page 7-6.*

7.3.1 Comms data control register

The following instruction returns the control register value in Rd:

```
MRC p14, 0, Rd, c0, c0
```

Two bits in this control register provide synchronized handshaking between the target and the host debugger:

Bit 1 (W bit)	Denotes whether the comms data write register is free (from the target point of view):
W = 0	New data can be written by the target application.
W = 1	The host debugger can scan new data out of the write register.
Bit 0 (R bit)	Denotes whether there is new data in the comms data read register (from the target point of view):
R = 1	New data is available to be read by the target application.
R = 0	The host debugger can scan new data into the read register.

———— Note ————

The debugger cannot use coprocessor 14 to access the debug communications channel directly, because this has no meaning to the debugger. Instead, the debugger can read from and write to the DCC registers using the scan chain. The DCC data and control registers are mapped into addresses in the EmbeddedICE logic. To view the EmbeddedICE logic registers, see the documentation for your debugger and debug target.

7.3.2 Target to debugger communication

This is the sequence of events for an application running on the ARM core to communicate with a debugger running on the host:

1. The target application verifies that the DCC write register is free for use. It does this using the MRC instruction to read the debug communications channel control register to check that the W bit is clear.
2. If the W bit is clear, the comms data write register is clear and the application writes a word to it using an MCR instruction to coprocessor 14. The action of writing to the register automatically sets the W bit. If the W bit is set, the debugger has not emptied the comms data write register. If the application has to send another word, it must poll the W bit until it is clear.
3. The debugger polls the comms data control register through scan chain 2. If the debugger sees that the W bit is set, it can read the DCC data register to read the message sent by the application. The process of reading the data automatically clears the W bit in the comms data control register.

Example 7-1 shows how this works. The example code is available in the main examples directory, in ...\\dcc\\outchan.s.

Example 7-1

```

        AREA OutChannel, CODE, READONLY
        ENTRY
        MOV    r1,#3          ; Number of words to send
        ADR    r2, outdata    ; Address of data to send
pollout
        MRC    p14,0,r0,c0,c0 ; Read control register
        TST    r0, #2
        BNE    pollout        ; if W set, register still full
write
        LDR    r3,[r2],#4     ; Read word from outdata
                                ; into r3 and update the pointer
        MCR    p14,0,r3,c1,c0 ; Write word from r3
        SUBS   r1,r1,#1       ; Update counter
        BNE    pollout        ; Loop if more words to be written
        MOV    r0, #0x18      ; Angel_SWIreason_ReportException
        LDR    r1, =0x20026    ; ADP_Stopped_ApplicationExit
        SWI    0x123456        ; ARM semihosting SWI
outdata
        DCB    "Hello there!"
        END

```

To execute the example:

1. Assemble outchan.s:
`armasm -g outchan.s`
2. Link the output object:
`armlink outchan.o -o outchan.axf`
The link step creates the executable file outchan.axf
3. Load and execute the image. See your debugger documentation for details.

7.3.3 Debugger to target communication

This is the sequence of events for message transfer from a debugger running on the host to the application running on the core:

1. The debugger polls the comms data control register R bit. If the R bit is clear, the comms data read register is clear and data can be written there for the target application to read.
2. The debugger scans the data into the comms data read register through scan chain 2. The R bit in the comms data control register is automatically set by this.
3. The target application polls the R bit in the comms data control register. If it is set, there is data in the comms data read register that can be read by the application, using an MRC instruction to read from coprocessor 14. The R bit is cleared as part of the read instruction.

The target application code shown in Example 7-2 on page 7-7 shows this in action. The example code is available in the main examples directory, in `...\dcc\inchan.s`.

To execute the example:

1. Create an input file on the host containing, for example, `And goodbye!`.
2. Assemble inchan.s:
`armasm -g inchan.s`
3. Link the output object:
`armlink inchan.o -o inchan.axf`
The link step creates the executable file inchan.axf
4. Load the and execute the image. See your debugger documentation for details.

Example 7-2

```

        AREA InChannel, CODE, READONLY
        ENTRY
        MOV    r1,#3          ; Number of words to read
        LDR    r2, =indata    ; Address to store data read
pollin
        MRC    p14,0,r0,c0,c0 ; Read control register
        TST    r0, #1
        BEQ    pollin         ; If R bit clear then loop
read
        MRC    p14,0,r3,c1,c0 ; read word into r3
        STR    r3,[r2],#4     ; Store to memory and
                                ; update pointer
        SUBS   r1,r1,#1       ; Update counter
        BNE    pollin         ; Loop if more words to read
        MOV    r0, #0x18      ; Angel_SWIreason_ReportException
        LDR    r1, =0x20026   ; ADP_Stopped_ApplicationExit
        SWI    0x123456       ; ARM semihosting SWI

        AREA Storage, DATA, READWRITE
indata
        DCB    "Duffmessage#"
        END

```

7.4 Interrupt-driven debug communications

The examples given in *Polled debug communications* on page 7-4 demonstrate polling the DCC. You can convert these to interrupt-driven examples by connecting up COMMRX and COMMTX signals from the Embedded ICE logic to your interrupt controller.

The read and write code in Example 7-1 on page 7-5 and Example 7-2 on page 7-7 can then be moved into an interrupt handler.

See *Interrupt handlers* on page 6-29 for information on writing interrupt handlers.

7.5 Access from Thumb state

Because the Thumb® instruction set does not contain coprocessor instructions, you cannot use the debug communications channel while the core is in Thumb state.

There are three possible ways around this:

- You can write each polling routine in a SWI handler, that can then be executed while in either ARM or Thumb state. Entering the SWI handler immediately puts the core into ARM state where the coprocessor instructions are available. See Chapter 6 *Handling Processor Exceptions* for more information on SWIs.
- Thumb code can make interworking calls to ARM subroutines that implement the polling. See Chapter 4 *Interworking ARM and Thumb* for more information on mixing ARM and Thumb code.
- Use interrupt-driven communication rather than polled communication. The interrupt handler would be written in ARM instructions, so the coprocessor instructions can be accessed directly.

Appendix A

Using the Procedure Call Standard

This appendix describes how to use the *Procedure Call Standard for the ARM® Architecture* (AAPCS) with *RealView® Compilation Tools* (RVCT). It contains the following sections:

- *About the AAPCS* on page A-2
- *Using AAPCS options* on page A-4.

A.1 About the AAPCS

The AAPCS forms part of the *Application Binary Interface (ABI) for the ARM Architecture (base standard)* [BSABI] specification. It specifies the base standard for a family of *Procedure Call Standard (PCS)* variants.

For more detailed information see:

- the *ABI for the ARM Architecture (base standard)* [BSABI], `bsabi.pdf`, in `install_directory\Documentation\Specifications\...`
- the *Procedure Call Standard for the ARM Architecture* specification [AAPCS], `aapcs.pdf`, in `install_directory\Documentation\Specifications\...`
- the ARM website <http://www.arm.com>.

By writing code that adheres to the AAPCS, you can ensure that separately compiled and assembled modules can work together.

This section describes:

- *PCS variants*
- *ARM C libraries* on page A-3
- *ARM assembler* on page A-3
- *Conformance to the AAPCS* on page A-3.

A.1.1 PCS variants

The PCS variants comprise a base standard modified by options that you can select independently. Code conforming to the base standard runs faster than, and occupies less memory than, code conforming to other variants. However, code conforming to the base standard does not provide for:

- position independent data or code
- re-entry to routines with independent data for each invocation
- stack checking.

Many details of the standard are the same, whichever variant you use. The compiler or assembler sets *attributes* in the ELF object file that record the variant you have chosen. In general, you must choose one variant and then use it for all subroutines that have to work together.

A.1.2 ARM C libraries

There are several variants of the ARM C libraries (see *RealView Compilation Tools v2.2 Compiler and Libraries Guide*). The linker selects a variant to link with your object files. It selects the best variant compatible with the AAPCS options recorded in your object files. See the linker chapter in *RealView Compilation Tools v2.2 Linker and Utilities Guide*.

A.1.3 ARM assembler

The AAPCS options you use do not affect the code produced by the assembler. They are simply an assertion that the code in the input files uses a particular PCS variant. The AAPCS options cause attributes to be set in the object file produced by the assembler. The linker then uses these attributes to check compatibility of files and to select appropriate library variants.

A.1.4 Conformance to the AAPCS

extern routines compiled using the ARM compiler conform to the selected variant of the AAPCS. You are responsible for ensuring that routines written in assembly language conform to the selected variant of the AAPCS.

To conform to the AAPCS, an assembly language routine must:

- follow all details of the standard at publicly visible interfaces
- follow the AAPCS rules of stack usage at all times
- be assembled with the appropriate `--apcs` options (see *RealView Compilation Tools v2.2 Assembler Guide*).

Conformance of ARM Architectures

Architectures ARMv5T and later conform to the AAPCS, because they provide direct interworking support. However, if you are compiling or assembling code for interworking on earlier architectures, you must also use the `--apcs /interwork` option.

A.2 Using AAPCS options

Use the `--apcs [qualifiers]` command-line option to specify that you are using the AAPCS. The qualifiers that you use enable you to specify the PCS variant used by the compiler.

Here there must be:

- at least one qualifier present
- a space between `--apcs` and the list of qualifiers
- no space between the qualifiers.

Use `--apcs /none` to specify that you are not using the AAPCS. This means that AAPCS registers are not set up. Qualifiers are not allowed.

Note

For full details on all these command-line options see *RealView Compilation Tools v2.2 Compiler and Libraries Guide*.

The main AAPCS variants are described in:

- *Position independence*
- *Stack checking*
- *Interworking* on page A-5
- *Floating-point* on page A-5.

A.2.1 Position independence

See *RealView Compilation Tools v2.2 Compiler and Libraries Guide* for information on the `--apcs` position independence qualifiers.

A.2.2 Stack checking

It is possible to write assembly code in such a way that stack limit checking is irrelevant. The code in a file might not require stack limit checking, but can still be compatible with other code assembled with either `--apcs /swst` or the default `--apcs /noswst`. In this case, use the option `--apcs /swstna` to specify that software stack limit checking is not applicable.

Rules for stack limit checked code

In the stack limit checked variants of the AAPCS:

- The *stack limit pointer* (`$1`) must point at least 256 bytes above the lowest usable address in the stack.

Note

If an interrupt handler can use the User mode stack, you must allocate sufficient space for it, between `s1` and the lowest usable address in the stack, in addition to the 256 bytes.

- `s1` must not be changed by code compiled or assembled with stack limit checking selected. `s1` is altered by runtime support code.
- The value held in `sp` must always be greater than or equal to the value in `s1`.

Register usage with stack limit checking

You must not change `r10`, or restore it, in routines assembled or compiled with the stack limit checking option selected. Register `r10` is the stack limit pointer, `s1`.

In all other respects the usage of registers is the same with or without stack limit checking.

Stack checking in assembly language

If you select the software stack checking option (`/swst`) for your assembly code, it is your responsibility to write code that performs stack checking.

A.2.3 Interworking

Interworking is part of the base standard AAPCS. Therefore, ARMv5T architectures, and later, conform to the AAPCS, because they provide direct interworking support.

For earlier architectures, you must use the `--apcs /interwork` option when compiling or assembling code for interworking.

The use of registers is the same with or without interworking.

See Chapter 4 *Interworking ARM and Thumb* for more information on ARM/Thumb interworking, and the chapter on using the basic linker functionality in *RealView Compilation Tools v2.2 Linker and Utilities Guide* for information on the automatically generated interworking veneers.

A.2.4 Floating-point

The AAPCS supports the VFP architecture as an AAPCS variant. See *RealView Compilation Tools v2.2 Compiler and Libraries Guide* for more information.

Glossary

The items in this glossary are listed in alphabetical order, with any symbols and numerics appearing at the end.

AAPCS *See* Procedure Call Standard for the ARM Architecture.

ABI for the ARM Architecture (base standard) (BSABI)

The ABI for the ARM Architecture is a collection of specifications, some open and some specific to ARM architecture, that regulate the inter-operation of binary code in a range of ARM architecture-based execution environments. The base standard specifies those aspects of code generation that must be standardized to support inter-operation and is aimed at authors and vendors of C and C++ compilers, linkers, and runtime libraries.

ADS *See* ARM Developer Suite.

American National Standards Institute (ANSI)

An organization that specifies standards for, among other things, computer software.

Angel A debug monitor that enables you to develop and debug applications running on ARM architecture-based hardware. Angel can debug applications running in either ARM state or Thumb state.

ANSI *See* American National Standards Institute.

Application Specific Integrated Circuit (ASIC)

An integrated circuit designed or adapted for a specific application. Traditionally called a custom circuit.

Application Specific Standard Product (ASSP)

An integrated circuit (IC) dedicated to one specific application (like a custom IC) but with several possible customers (unlike a custom IC).

ARM Developer Suite (ADS)

A suite of software development applications, together with supporting documentation and examples, that enable you to write and debug applications for the ARM family of RISC processors. ADS is superseded by RealView Developer Suite (RVDS).

See also RealView Developer Suite.

ARM state

A processor that is executing ARM instructions is operating in ARM state. The processor switches to Thumb state (and to recognizing Thumb instructions) when directed to do so by a state-changing instruction such as BX, BLX.

See also Thumb state.

ASIC

See Application Specific Integrated Circuit.

ASSP

See Application Specific Standard Product.

Big-Endian

Memory organization where the least significant byte of a word is at a higher address than the most significant byte.

Canonical Frame Address (CFA)

In DWARF, this is an address on the stack specifying where the call frame of an interrupted function is located.

CFA

See Canonical Frame Address.

CISC

See Complex Instruction Set Computing.

Complex Instruction Set Computing (CISC)

A highly flexible but not very efficient device containing a number of instructions for specific applications in a microprocessor.

Coprocessor

An additional processor which is used for certain operations. Usually used for floating-point math calculations, signal processing, or memory management.

Debug With Arbitrary Record Format (DWARF)

ARM code generation tools generate debug information in DWARF2 format by default. From RVCT v2.2, you can optionally generate DWARF3 format (Draft Standard 9).

Deprecated

A deprecated option or feature is one that you are strongly discouraged from using. Deprecated options and features will not be supported in future versions of the product.

Doubleword	A 64-bit unit of information. Contents are taken as being an unsigned integer unless otherwise stated.
DWARF	<i>See</i> Debug With Arbitrary Record Format.
EC++	A variant of C++ designed to be used for embedded applications.
ELF	<i>See</i> Executable and Linking Format.
Embedded	Applications that are developed as firmware. Assembler functions placed out-of-line in a C or C++ program. <i>See also</i> Inline.
Executable and Linking Format (ELF)	The industry standard binary file format used by RealView Compilation Tools. ELF object format is produced by the ARM object producing tools such as armcc and armasm. The ARM linker accepts ELF object files and can output either an ELF executable file, or a partially linked ELF object.
Execution view	The address of regions and sections after the image has been loaded into memory and started execution.
Flash memory	Nonvolatile memory that is often used to hold application code.
Halfword	A 16-bit unit of information. Contents are taken as being an unsigned integer unless otherwise stated.
Heap	The portion of computer memory that can be used for creating new variables.
Image	An executable file which has been loaded onto a processor for execution. A binary execution file loaded onto a processor and given a thread of execution. An image might have multiple threads. An image is related to the processor on which its default thread runs.
Inline	Functions that are repeated in code each time they are used rather than having a common subroutine. Assembler code placed within a C or C++ program. <i>See also</i> Output section <i>and</i> Embedded.
Input section	Contains code or initialized data or describes a fragment of memory that must be set to zero before the application starts. <i>See also</i> Output section.
International Standards Organization (ISO)	An organization that specifies standards for, among other things, computer software.
Interworking	Producing an application that uses both ARM and Thumb code.

ISO	<i>See</i> International Standards Organization.
Library	A collection of assembler or compiler output objects grouped together into a single repository.
Linker	Software which produces a single image from one or more source assembler or compiler output objects.
Little-endian	Memory organization where the least significant byte of a word is at a lower address than the most significant byte.
Load view	The address of regions and sections when the image has been loaded into memory but has not yet started execution.
Local	An object that is only accessible to the subroutine that created it.
Memory Management Unit (MMU)	Hardware that controls caches and access permissions to blocks of memory, and translates virtual to physical addresses.
MMU	<i>See</i> Memory Management Unit.
Multi-ICE	A JTAG-based tool for debugging embedded systems.
Output section	Is a contiguous sequence of input sections that have the same RO, RW, or ZI attributes. The sections are grouped together in larger fragments called regions. The regions are grouped together into the final executable image. <i>See also</i> Region.
PIC	Position Independent Code. <i>See also</i> ROPI.
PID	Position Independent Data <i>or</i> the ARM Platform-Independent Development (PID) board. <i>See also</i> RWPI
Procedure Call Standard for the ARM Architecture (AAPCS)	<i>Procedure Call Standard for the ARM Architecture</i> defines how registers and the stack will be used for subroutine calls.
Profiling	Accumulation of statistics during execution of a program being debugged, to measure performance or to determine critical areas of code.
Program image	<i>See</i> Image.
Read-Only Position Independent (ROPI)	Code or read-only data that can be placed at any address.

Read Write Position Independent (RWPI)

Read/write data addresses that can be changed at runtime.

RealView ARMulator ISS (RVISS)

The latest version of the ARM simulator, RealView ARMulator ISS is supplied with RealView Developer Suite.

RVISS is a collection of programs that simulate the instruction sets and architecture of various ARM processors. This provides instruction-accurate simulation and enables ARM and Thumb executable programs to be run on non-native hardware.

RVISS provides modules that model:

- the ARM processor core
- the memory used by the processor.

There are alternative predefined models for each of these parts. However, you can create your own models if a supplied model does not meet your requirements.

RealView Compilation Tools (RVCT)

RealView Compilation Tools is a suite of tools, together with supporting documentation and examples, that enables you to write and build applications for the ARM family of *RISC* processors.

RealView Developer Suite (RVDS)

The latest suite of software development applications, together with supporting documentation and examples, that enable you to write and debug applications for the ARM family of *RISC* processors.

RealView ICE (RVI) A JTAG-based debug solution to debug software running on ARM processors.

Reduced Instruction Set Computing (RISC)

Device where the number of instructions a microprocessor runs for a specific application are reduced from a general purpose Complex Instruction Set Computing (CISC) device to create a more efficient operating system.

Reentrancy The ability of a subroutine to have more than one instance of the code active. Each instance of the subroutine call has its own copy of any required static data.

Region In an image, a region is a contiguous sequence of one to three output sections (RO, RW, and ZI). A region typically maps onto a physical memory device, such as ROM, RAM, or peripheral.

Remapping Changing the address of physical memory or devices after the application has started executing. This is typically done to enable RAM to replace ROM when the initialization has been done.

Retargeting The process of moving code designed for one execution environment to a new execution environment.

RISC	<i>See</i> Reduced Instruction Set Computing.
ROPI	<i>See</i> Read-Only Position Independent.
RTOS	Real Time Operating System.
RVCT	<i>See</i> RealView Compilation Tools.
RVDS	<i>See</i> RealView Developer Suite.
RVISS	<i>See</i> RealView ARMulator ISS.
RWPI	<i>See</i> Read Write Position Independent.
Scatter-loading	Assigning the address and grouping of code and data sections individually rather than using single large blocks.
Section	A block of software code or data for an Image. <i>See also</i> Input section.
Semihosting	A mechanism whereby the target communicates I/O requests made in the application code to the host system, rather than attempting to support the I/O itself.
Software Interrupt (SWI)	An instruction that causes the processor to call a programmer-specified subroutine. Used by ARM code to handle semihosting.
SWI	<i>See</i> Software Interrupt.
Target	The actual target processor, (real or simulated), on which the application is running.
TCM	<i>See</i> Tightly Coupled Memory.
Thread	A context of execution on a processor. A thread is always related to a processor and might or might not be associated with an image.
Thumb state	A processor that is executing Thumb instructions is operating in Thumb state. The processor switches to ARM state (and to recognizing ARM instructions) when directed to do so by a state-changing instruction such as BX, BLX. <i>See also</i> ARM state.
Tightly Coupled Memory (TCM)	Tightly Coupled Memory replaces an area of off-chip memory when enabled.
Veneer	A small block of code used with subroutine calls when there is a requirement to change processor state or branch to an address that cannot be reached in the current processor state.

Word

A 32-bit unit of information. Contents are taken as being an unsigned integer unless otherwise stated.

Index

A

AAPCS

- about 1-12, 3-2, A-2
 - assembler A-3
 - conformance A-3
 - interworking ARM and Thumb 4-2, A-5
 - libraries A-3
 - position independence 3-2
 - processes 3-6
 - read-only position independence 3-3
 - read-write position independence 3-6
 - reentrant routines 3-6
 - relocatable code 3-2
 - ROPI 3-3
 - RWPI 3-6
 - static base register 3-2, 3-6
 - threads 3-6
 - /fpic qualifier 3-4
- AAPCS options
- apcs A-5
 - fpu A-5
 - /interwork 4-4, A-5
 - /swst A-5
 - /swstna A-4
- ABI 3-2, A-2
- Accessing
- debug comms channel 7-2
- ANSI C
- see* ISO C
- ARM architecture v5T
- interworking ARM and Thumb 4-11
- ARM architectures
- ARMv6 support 1-13
- ARMv6
- alignment support 1-14
 - Byte Invariant Addressing mode 1-14
 - endian support 1-14
 - instruction generation 1-13
 - support 1-13
- ARM/Thumb synonyms 4-17

ASIC 1-7

Assembler

- embedded, *see* Embedded assembler
- inline, *see* Inline assembler
- mode changing 4-8

Assembly language

- calling C 5-11
- calling C++ 5-15
- calling from C 5-10
- calling from C++ 5-13
- embedded assembler 5-2
- interrupt handlers 6-33
- interworking ARM and Thumb 4-7, 4-18
- interworking using veneers 4-18

ASSP 1-7

B

Banked registers 6-3

Base classes

- in mixed languages 5-9

- BE32 1-14
 - BE8 1-14
 - Big-endian code 1-14
 - Big-endian data 1-14
 - Big-endian mode 1-14
 - Bit 0, use in BX instruction 4-8
 - BSABI 3-2, A-2
 - BX instruction 4-7
 - bit 0 usage 4-8
 - Byte Addressing mode 1-14
 - BE32 1-14
 - BE8 1-14
- ## C
- calling assembly language 5-10
 - calling C++ 5-14
 - calling from assembler 5-7
 - calling from assembly language 5-11
 - calling from C++ 5-7, 5-12
 - global variables from assembly language 5-4
 - Interworking ARM and Thumb 4-13
 - linkage 5-8
 - using header files from C++ 5-5
- ## Calling
- assembler from C++ 5-7
 - C from assembly language 5-7
 - C from C++ 5-7, 5-9
 - C++ from assembly language 5-7
 - interworking examples
 - interworking veneers
 - language conventions 5-7
- ## Chaining exception handlers
- 6-44
- ## CISC
- 1-3
- ## Code
- density and interworking 4-3
 - PIC 3-2
 - position-independent 3-2
- ## Comms channel, debug
- 7-1
- ## Comms data control register
- 7-4
- ## Comms data registers
- 7-3
- ## Compiler
- embedded assembler 5-2
- ## Compiler options
- apcs A-4
- ## Context switch
- 6-37
- ## Coprocessor
- 14 7-3
- ## Coprocessors
- Undefined Instruction handlers 6-40
- ## CPSR
- 6-8
- ## C++
- calling assembly language 5-13
 - calling C 5-12
 - calling conventions 5-8
 - calling from assembler 5-7
 - calling from assembly language 5-15, 5-17
 - calling from C 5-7, 5-14, 5-17
 - data types in mixed languages 5-9
- ## D
- ## Data
- PID 3-2
 - position-independent 3-2
- ## Data Abort
- exception 6-3
 - handler 6-42
 - LDM 6-42
 - LDR 6-42
 - returning from 6-12
 - STM 6-42
 - STR 6-42
 - SWP 6-42
- ## DCC
- 7-1
- ## Debug
- interrupt-driven comms 7-8
 - polled communications 7-4
- ## Debug Communications Channel
- 7-1
- ## Debuggers
- communicating with target 7-6
- ## Directives
- IMPORT 5-4
- ## Directives, assembly language
- IMPORT 5-4
- ## E
- ## Embedded assembler
- about 5-2
 - differences from inline assembly 5-3
- ## Embedded software
- application startup 2-9
 - avoiding semihosting 2-12
 - C library structure 2-6
 - default memory map 2-7
 - developing 2-1
 - developing with RVCT 2-2
 - execution mode 2-33
 - execution regions 2-15
 - hardware initialization 2-32
 - initialization sequence 2-26
 - linker placement rules 2-8
 - load regions 2-15
 - loading stack and heap 2-36
 - local memory setup 2-30
 - locating target peripherals 2-35
 - one-region model 2-21
 - overview 2-2
 - placing objects in scatter-loading file 2-18
 - placing stack and heap 2-20
 - placing symbols 2-36
 - retargeting C library 2-11
 - ROM/RAM remapping 2-28
 - root regions 2-19
 - runtime memory model 2-21
 - RVCT behavior 2-4
 - scatter-loading 2-15
 - scatter-loading and memory setup 2-30
 - scatter-loading file syntax 2-16
 - semihosting support 2-5
 - stack pointer initialization 2-31
 - tailoring image memory map 2-15
 - two-region model 2-22
 - utilizing linker generated symbols 2-38
 - vector table 2-27
- ## EMPTY attribute
- 2-39
- ## Examples
- main directory 1-2, 3-5, 3-7
- ## Exception handlers
- chaining 6-44
 - Data Abort 6-42
 - extending 6-44
 - installing 6-14
 - installing from C 6-16

- installing on reset 6-14
- Prefetch Abort 6-12, 6-41
- Reset 6-39
- returning from 6-9
- subroutines in 6-46
- SWI 6-19, 6-20, 6-21, 6-23
- Thumb 6-8
- Undefined Instruction 6-40
- writing 6-13
- Exceptions 6-2
 - Data Abort 6-12
 - entering 6-8
 - FIQ 6-11
 - handler code 6-13
 - illegal address 6-2, 6-3
 - installing handlers 6-14
 - IRQ 6-3, 6-11
 - leaving 6-8
 - Prefetch Abort 6-2, 6-11
 - priorities 6-4
 - reset 6-2
 - response by processors 6-8
 - returning from 6-11
 - SWI 6-2, 6-10
 - SWI handlers 6-19, 6-20, 6-21, 6-23
 - Undefined Instruction 6-2, 6-10
 - use of modes 6-3
 - use of registers 6-3
 - vector table 6-3, 6-14
- Execution
 - speed 4-3, 6-29
- Extending exception handlers 6-44
- extern "C" 5-5, 5-8, 5-9

F

- Fault address register 6-43
- FIQ 6-3, 6-29
 - handler 6-11, 6-29
 - registers 6-29
- Floating-point
 - AAPCS options A-5
- FPA
 - Undefined Instruction handlers 6-40
- Function pointers
 - Thumb state 4-16

I

- Illegal address 6-2
- implicit **this** 5-7
- IMPORT directive 5-4
- Inline assembler
 - about 5-2
 - differences from embedded assembly 5-3
- Installation directory
 - default location ix
- Instructions, assembly language
 - BX 4-7
 - SWI 6-19
 - SWI (Thumb) 6-6
- Interrupt handlers
 - about 6-29
 - calling subroutines 6-30
 - external 6-29
 - in assembly language 6-33
 - prioritization 6-35
 - reentrant 6-31
 - simple 6-29
- Interrupt-driven debug comms 7-8
- Interworking ARM and Thumb
 - AAPCS 4-20
 - about 4-2
 - about veneers 4-2
 - ARM architecture v5T 4-11
 - assembly language 4-7, 4-18
 - BX instruction 4-7
 - C 4-14
 - C and C++ 4-13
 - C and C++ libraries 4-16
 - compiler command-line options 4-16
 - compiling code 4-13
 - detecting calls 4-4
 - duplicate functions 4-17
 - examples 4-10, 4-14, 4-19
 - exceptions 4-3
 - function pointers 4-16
 - in AAPCS 4-2
 - leaf functions 4-13
 - mixed languages 4-18, 4-20
 - non-Thumb processors 4-14
 - rules 4-16
 - veneers 4-13, 4-14, 4-18
- IRQ 6-29

- handler 6-11
- IRQ exception 6-3
- ISO C, header files 5-9

J

- JTAG 7-2
- Jump table 6-6, 6-20

L

- Leaf functions 4-13
- Legacy objects
 - in RVCT 1-12
- Link register 6-3
- Linking
 - and interworking 4-4, 4-13

M

- Main examples directory 1-2, 3-5, 3-7
- Mangling symbol names 5-7, 5-9
- Mixed language programming
 - interworking ARM and Thumb 4-18, 4-20
- MMU 1-7

N

- Nested SWIs 6-24

O

- One-region memory model 2-21

P

- `__packed` qualifier
 - fields in structures 1-4
 - pointers 1-4
- PCS
 - variants A-2
- PIC 3-3

- about 3-2
- PID 3-6
 - about 3-2
- Pointers
 - unaligned 1-3
- Polled debug communications 7-4
- Position independence
 - code 3-2
 - data 3-2
 - linking 3-4, 3-7
- Power-up 6-2
 - reset 1-11
 - soft reset 1-11
- Prefetch Abort
 - handler 6-12, 6-41
 - returning from 6-11
- Prefetch Abort exception 6-2
- Process control blocks 6-37
- Processors
 - responding to exceptions 6-8
- Program Status Register 5-2
- Protocol converter 7-2
- PSR 5-2

R

- Reentrant routines 3-6
- References 5-9
- Registers
 - debug comms channel 7-3
- Relocatable
 - code 3-2
 - data 3-2
- Reset
 - signaling power-up 1-11
- Reset exception 6-2
- Reset handlers 6-39
- Return address 6-10
- Return instruction 6-10
- RISC 1-2, 1-3
- ROM/RAM remapping 2-28
- ROPI 3-3
 - linker options 3-4
- Runtime memory models
 - one-region model 2-21
 - two-region model 2-22
- RVCT
 - components 1-2
- legacy objects 1-12
- RWPI 3-6
 - linker options 3-7

S

- Scatter-loading
 - description file syntax 2-16
 - embedded software 2-15
 - EMPTY attribute 2-39
 - execution regions 2-15
 - load regions 2-15
 - placing objects 2-18
 - placing stack and heap 2-20
 - root region 2-19
- Semihosting
 - avoiding in embedded software 2-12
 - for embedded software 2-5
 - with DCC 7-2
- Soft reset 6-2
 - signaling power-up 1-11
- Software FPA emulator
 - Undefined Instruction handlers 6-40
- Software interrupt
 - see* SWIs
- SPSR 6-3
- SPSR 6-8
 - T bit 6-6
- Stacks 6-4
 - stack pointer 6-3
 - supervisor 6-22
- Static base register
 - sb 3-2, 3-6
- Supervisor mode 6-22
- Supervisor stack 6-22
- SWI exception 6-2
- SWI instruction 6-19
 - Thumb 6-6
- SWIs
 - calling from assembly language 6-24
 - calling from C/C++ 6-24
 - handlers 6-19, 6-20, 6-21, 6-23
 - indirect 6-27
 - jump table 6-20
 - nested 6-24

- returning from 6-10
- Supervisor mode 6-22
- Thumb state 6-6
- Symbol names, mangling 5-7, 5-9
- Symbols
 - ARM/Thumb synonyms 4-17
 - multiple definitions of 4-17
- System mode 6-46
- System V
 - /fpic qualifier 3-4

T

- Target
 - communicating with debugger 7-5
- this**, implicit 5-7
- Threads 3-6
- Thumb
 - access to DCC 7-9
 - and **__irq** 6-30
 - BX instruction 4-8
 - changing to Thumb state, example 4-9
 - exception handler 6-8
 - handling exceptions 6-8
 - interworking with ARM 4-3
 - using duplicate function names 4-17
- Thumb state
 - function pointers 4-16
- Two-region memory model 2-22

U

- Unaligned data
 - fields in structures 1-4
 - LDR accesses to halfwords 1-6
 - pointers 1-3
- Undefined Instruction exception 6-2
- Undefined Instruction handler 6-10, 6-40
- User mode 6-3

V

- Vector table 2-27, 6-3, 6-8, 6-14, 6-29

Vector table and caches 6-18

Veneers

see Interworking

Symbols

__use_two_region_memory 2-22

